

Don't Talk All at Once!

Elevating Privileges on macOS by Audit Token Spoofing

Framework

XPC

Access a low-level interprocess communication mechanism.

macOS 10.10+

Mac Catalyst 13.0+

Overview

XPC provides a lightweight mechanism for basic interprocess communication. It allows you to create lightweight helper tools, called XPC services, that perform work on behalf of your app. The `launchd` system daemon manages these services, launching them on demand, shutting them down when idle, and restarting them if they crash. Benefits of XPC services include:

- Centralize work from multiple processes or mediate access to a shared resource.
- Delegate work so it continues beyond a client's life cycle.
- Privilege isolation to narrow the scope of access for different functionality.



Hello!

I'm Thijs Alkemade

Security Researcher at Computest

Computest
always on.



About me

- > Thijs Alkemade (<https://infosec.exchange/@xnyhps>)
- > Security researcher at Computest Sector 7 (The Netherlands)
- > Other work:
 - > 0-click RCE in Zoom, Pwn2Own Vancouver 2021
 - > Winning Pwn2Own Miami 2022 with 5 ICS vulnerabilities
 - > Last year's OBTS talk: "Process injection: breaking all macOS security layers with a single vulnerability"



XPC

- > XPC is often used between different privilege boundaries
 - > For example: privileged daemon performing requests from lower privileged processes
- > These often need to do an authorisation check: **“should this process be allowed to do this?”**
 - > Does it have a specific entitlement?
 - > Is the client signed by a specific Team ID?
 - > Does it have a specific TCC permission?
 - > Is the client sandboxed?

XPC

> `xpc_connection_get_pid` is **not safe** to use here!

Don't Trust the PID!

Stories of a simple logic bug and where to find it

Samuel Groß (@5aelo)


Issue 1223: MacOS/iOS userspace entitlement checking is racy

Reported by ianbeer@google.com on Thu, Mar 23, 2017 at 9:37 PM GMT+1

Project Member

Learn XPC exploitation - Part 2: Say no to the PID!

@WOJCIECH REGUŁA · APR 23, 2020 · 4 MIN READ

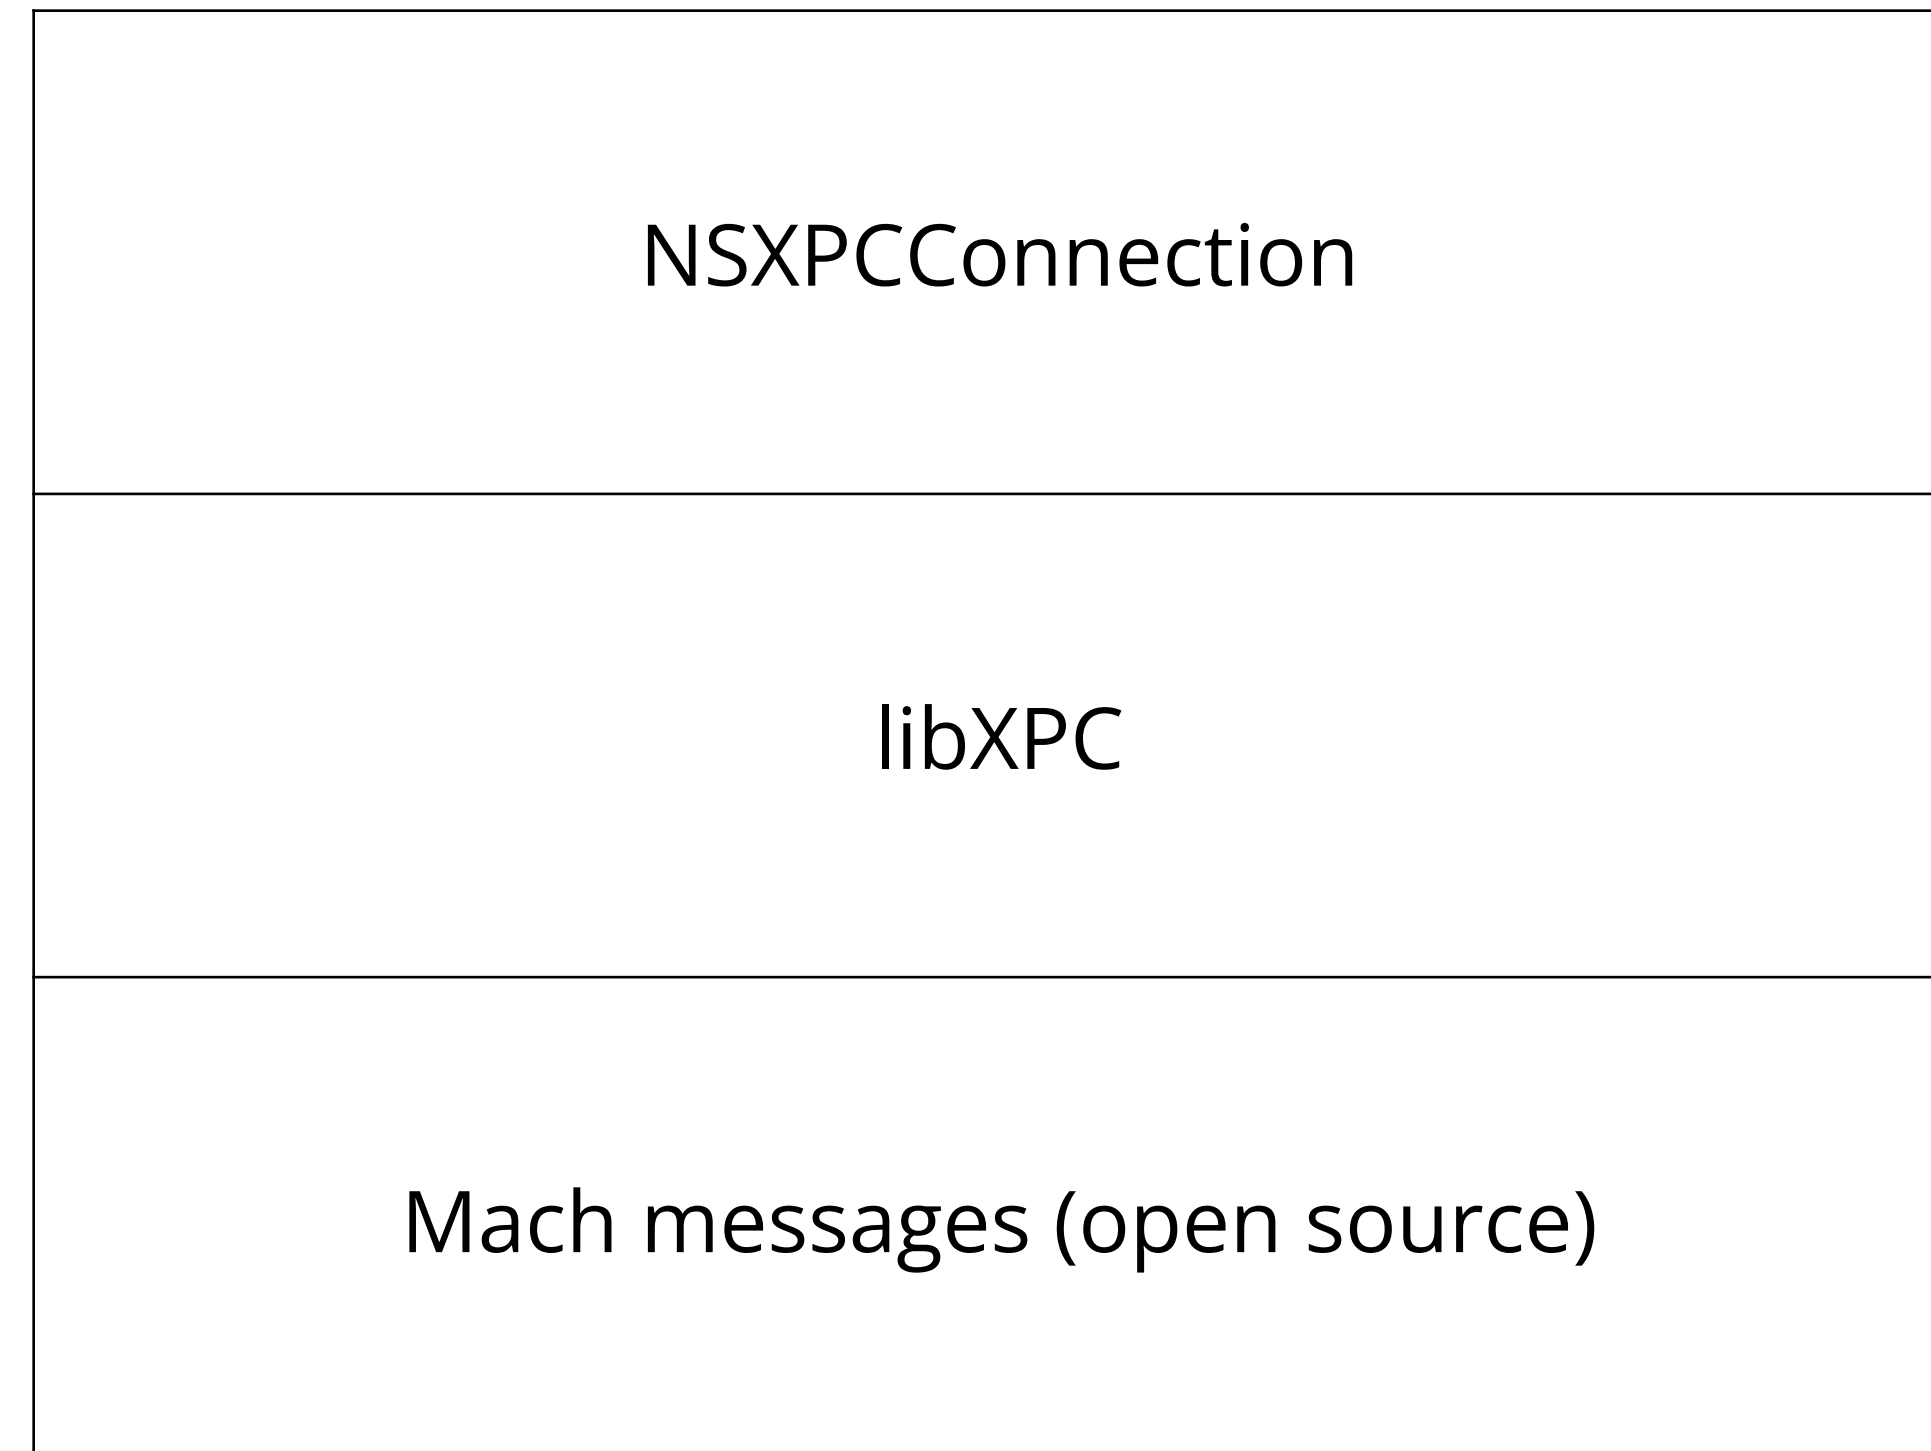


OBTS v3.0: "Abusing & Security XPC in macOS apps" - Wojciech Regula

XPC

- > It is well known that `xpc_connection_get_pid` is unsafe for authorising XPC clients
 - > It is vulnerable to a race condition:
 - > Send message
 - > Exec an authorised process (without changing PID)
 - > Hope the exec is done before the authorisation check
- > `xpc_connection_get_audit_token` is better. Audit tokens are safe, because they contain a **PID version**
- > But as I will show today, there are situations where even this function can be vulnerable to a race condition
- > Better: `xpc_dictionary_get_audit_token`.


XPC architecture



Mach messages 101

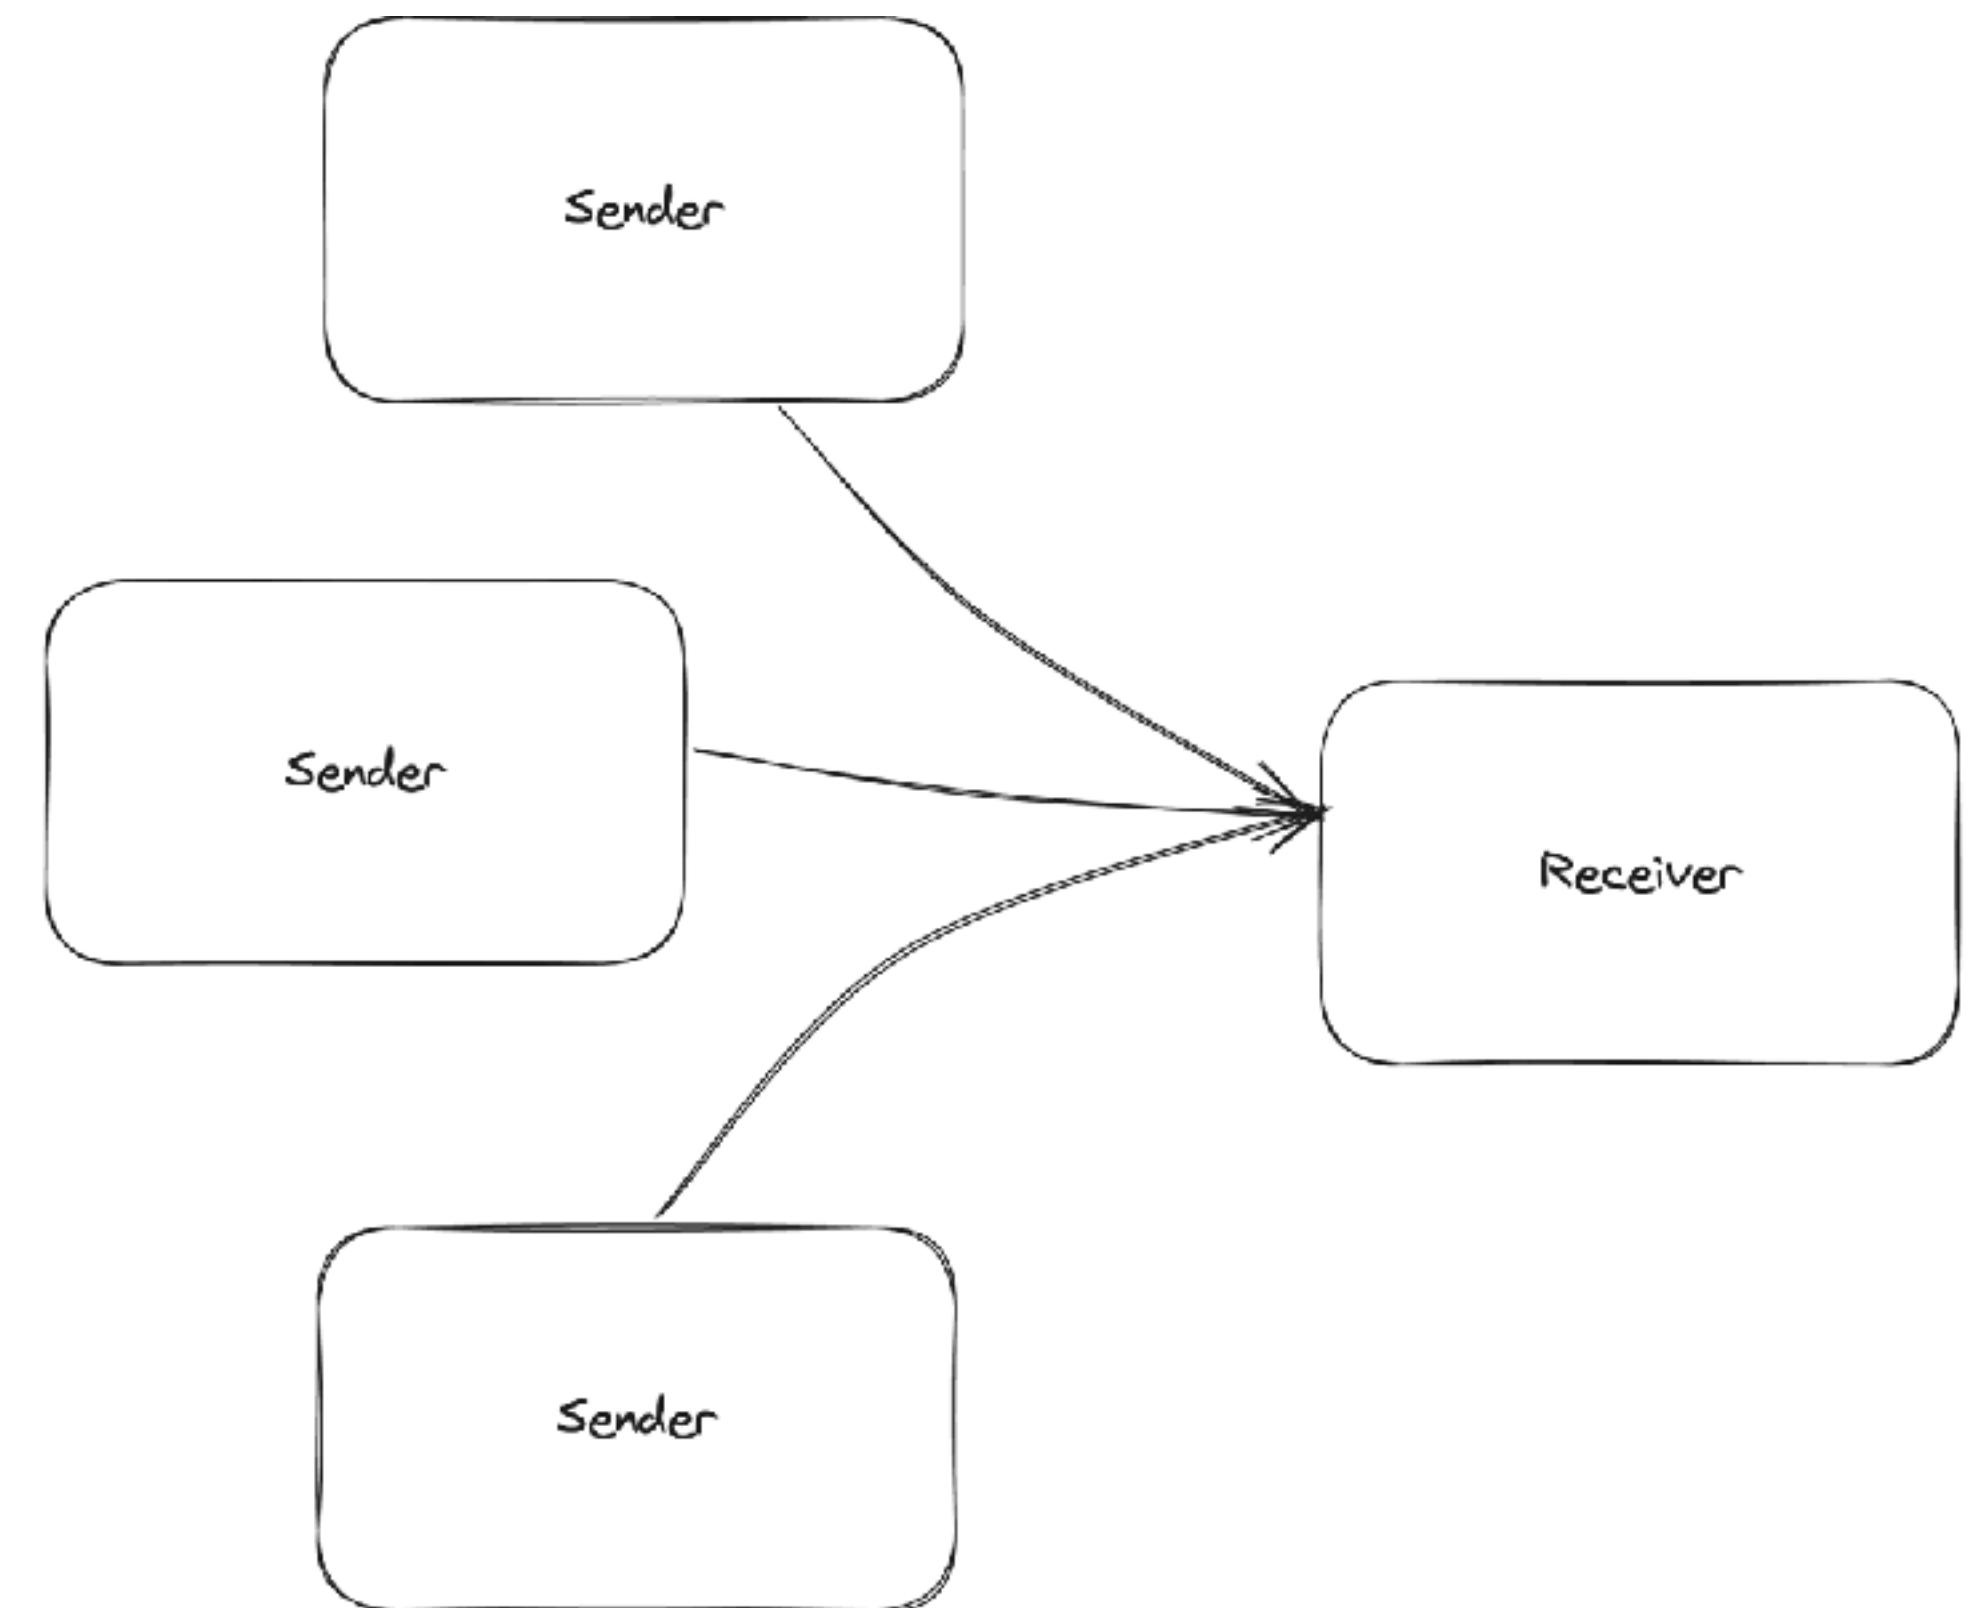
Mach messages 101

- > Mach messages are sent over “mach ports”
- > A mach port is a **single receiver, multiple sender** communications channel managed by the kernel
 - > The actual object exists in the kernel, process only has a reference to it (the **port name**)
- > Used extensively throughout the system
 - > Some kernel API's are just special mach ports



Mach messages 101

- > Three types of “rights” exist for a mach port:
 - > Receive right (only one process can have this)
 - > Send right
 - > Send-once right



Mach messages 101

- > Send rights can be **duplicated** (or converted into a send-once)
- > Port rights can be **transferred** with a mach message
- > For example, the “local_port” field of a mach message can be used to transfer a send-once right for a reply to a message

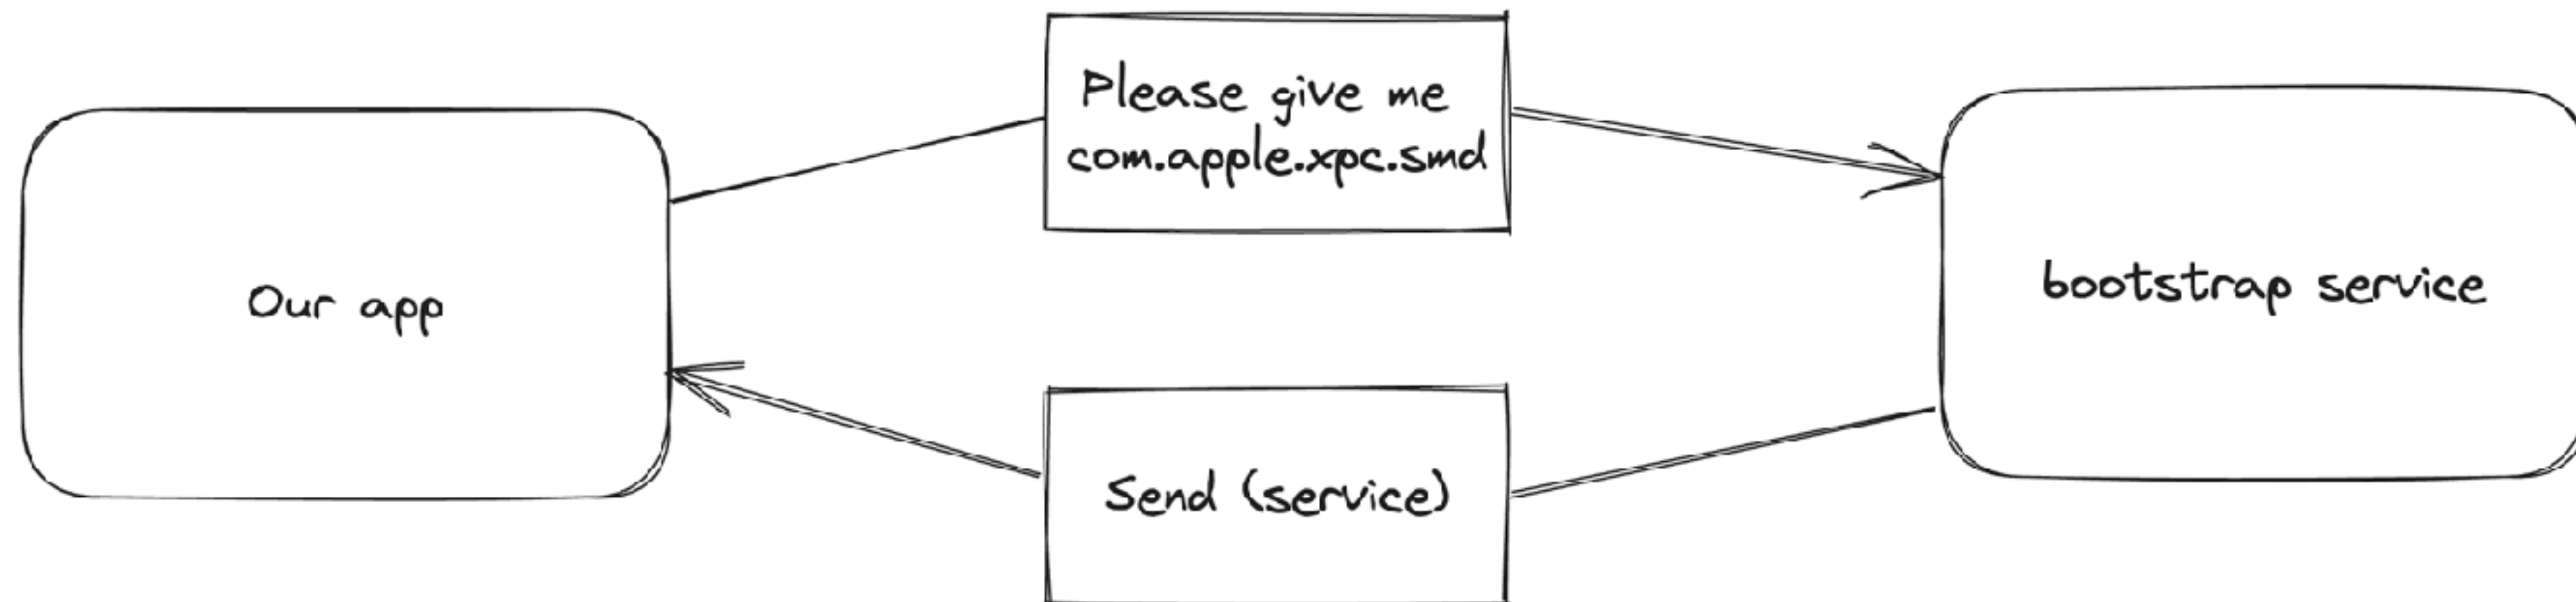
```
typedef struct {  
    mach_msg_bits_t      msg_bits;  
    mach_msg_size_t      msg_size;  
    mach_port_t          msg_remote_port;  
    mach_port_t          msg_local_port;  
    mach_port_name_t      msg_voucher_port;  
    mach_msg_id_t        msg_id;  
} mach_msg_header_t;
```

Mach messages 101

- > Some terminology help:
 - > A mach **port** is the entire communication channel (unlike “ports” in TCP/UDP)
 - > “Holds a send/receive **right**” just means “this is a sender/the receiver”
 - > Talking about rights is easier when they get transferred

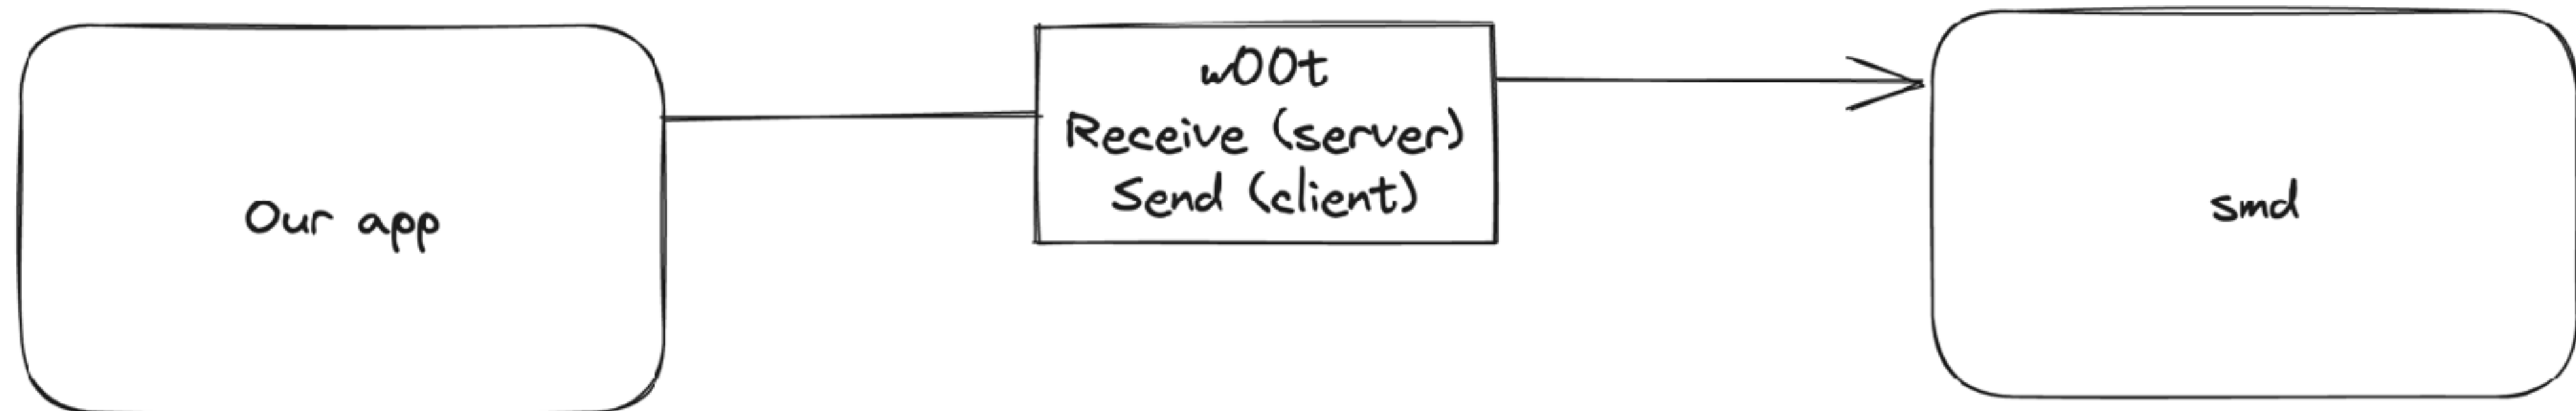
Establishing an XPC connection

- > To establish an XPC connection, a client first needs (a send right for) the **service port**
- > For a mach service, this is retrieved from the **bootstrap service** (in launchd)
- > To get the service port:
 1. Client asks for service port for a name
 2. launchd looks up the name and duplicates the associated send right
 3. launchd sends the duplicated send right back




Establishing an XPC connection

- > Then, the client generates two new mach ports: a **server port** and a **client port**
- > It sends a mach message (id: w00t) to the service port. This message:
 - > Transfers the **receive right** of the **server** port
 - > Transfers the **send right** to the **client** port
- > If the server accepts the connection, it starts listening on the server port




Establishing an XPC connection

- > Server can send messages to the client port, client can send messages to the server port
- > Note that this doesn't use the reply port mentioned earlier
 - > It is used for XPC, but only for `xpc_connection_send_message_with_reply{_sync}`



Audit tokens

- > How do those audit tokens come in here?
- > For the mach_msg syscall, it is possible to ask the kernel to append a “trailer” with the audit token using the flag MACH_RCV_TRAILER_AUDIT
- > Every time a message or reply is received, _xpc_connection_set_creds copies the audit token from the trailer to the connection object



The vulnerability

But wait...

> We have seen the following:

1. XPC connections use authorisation checks that assume it's a **one-to-one** communication channel
2. Mach ports are single receiver, **multiple sender** communication channels
3. `xpc_connection_get_audit_token` returns the audit token saved for the **most recently received message**



Research question

Can we set up an XPC connection where **multiple processes** can send a message?

And if so, can we use that to **spoof the audit token**?

Research question

Answer: **yes**, the client generates the mach ports for the connection, so can specifically construct them.

However, audit token spoofing only works under **limited circumstances**.

An example: smd

Variant 1

An example

- > smd is the **Service Manager Daemon**
- > For example: SMJobBless, which you can use to install a privileged helper tool.
 - > See also “Job(s) Bless Us! Privileged Operations on macOS” by Julia Vashchenko at OBTSv3
- > Our goal: install a privileged helper tool, **without** the user entering their password

Job(s) Bless Us!
Privileged Operations on macOS

An example

- > SMJobBless calls smd to perform the installation
- > smd has multiple actions, specified by the “routine” in the request
- > Routine 1004 is used by SMJobBless
 - > This routine is handled on a different dispatch queue using `dispatch_async`

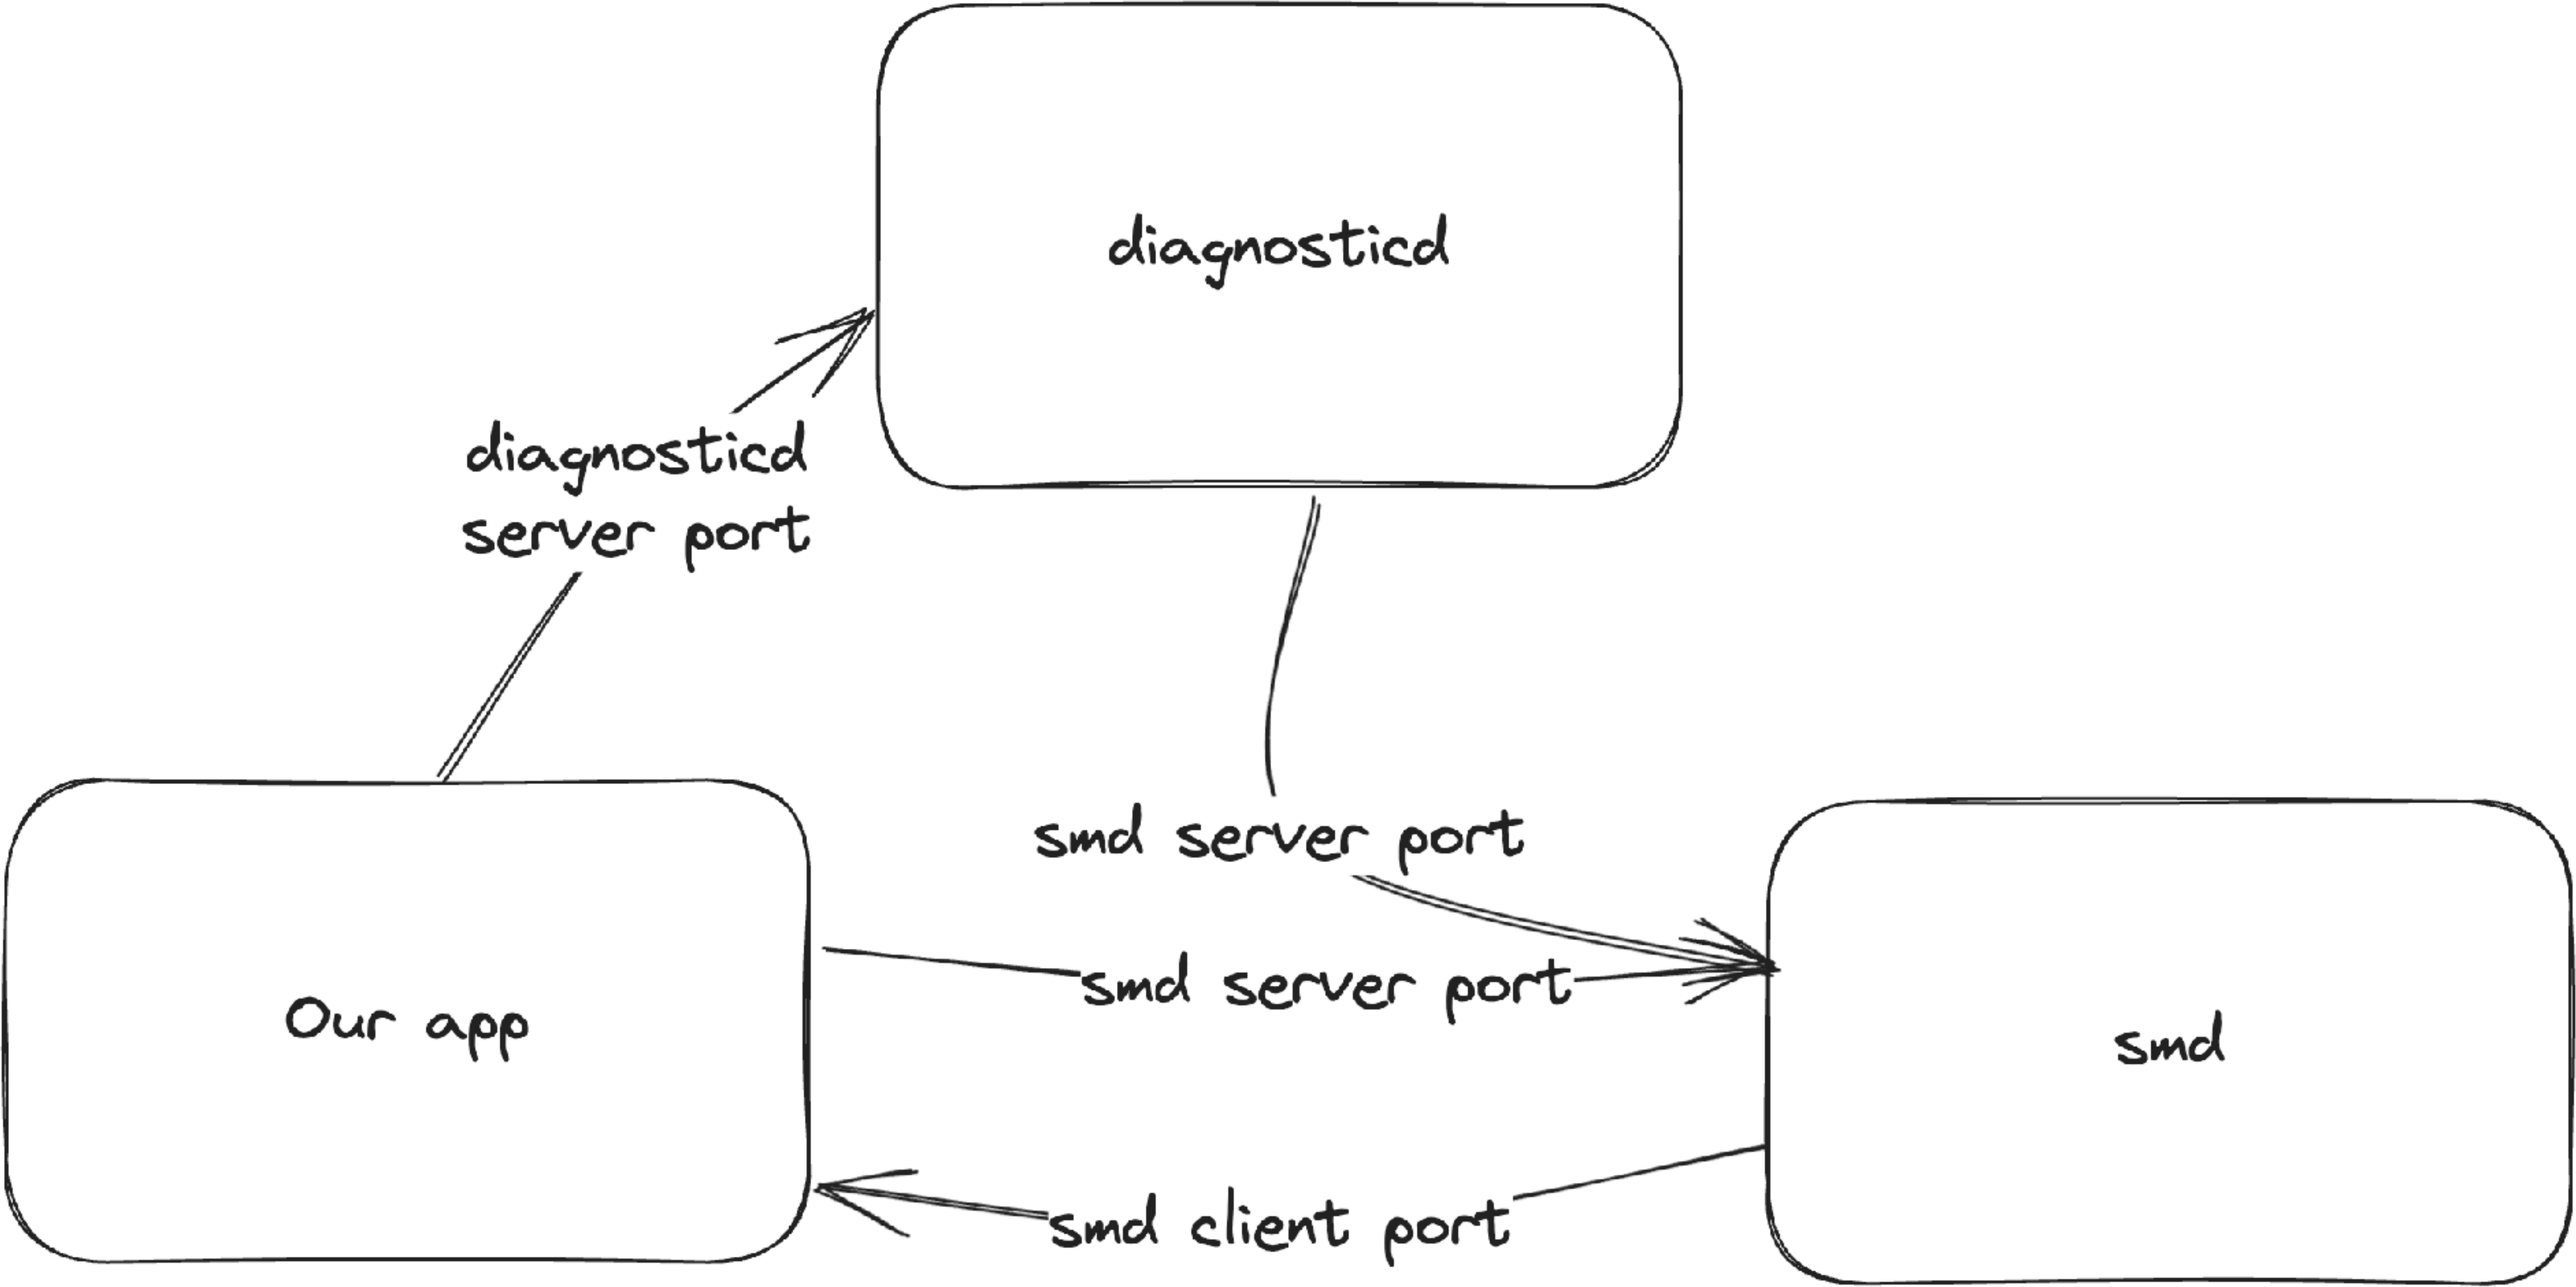
```
case 1004LL:
    buf.i64[0] = (__int64)_NSConcreteStackBlock;
    buf.i64[1] = 3254779904LL;
    buf.i64[2] = (__int64)handle_bless;
    buf.i64[3] = (__int64)&unk_100008180;
    v49 = objc_retain(v4);
    v50 = objc_retain(v5);
    dispatch_async((dispatch_queue_t)queue, &buf);
    goto LABEL_25;
```

An example

- > A process is allowed to install a new privileged helper tool if it passes **one** of the following checks:
 1. It is running as root
 2. It has the entitlement `com.apple.private.xpc.unauthenticated-bless`.
 3. It has an authorisation reference for `com.apple.ServiceManagement.bless-helper`
- > The third one is what is used if you call `SMJobBless`
 - > Our goal is to pass the first

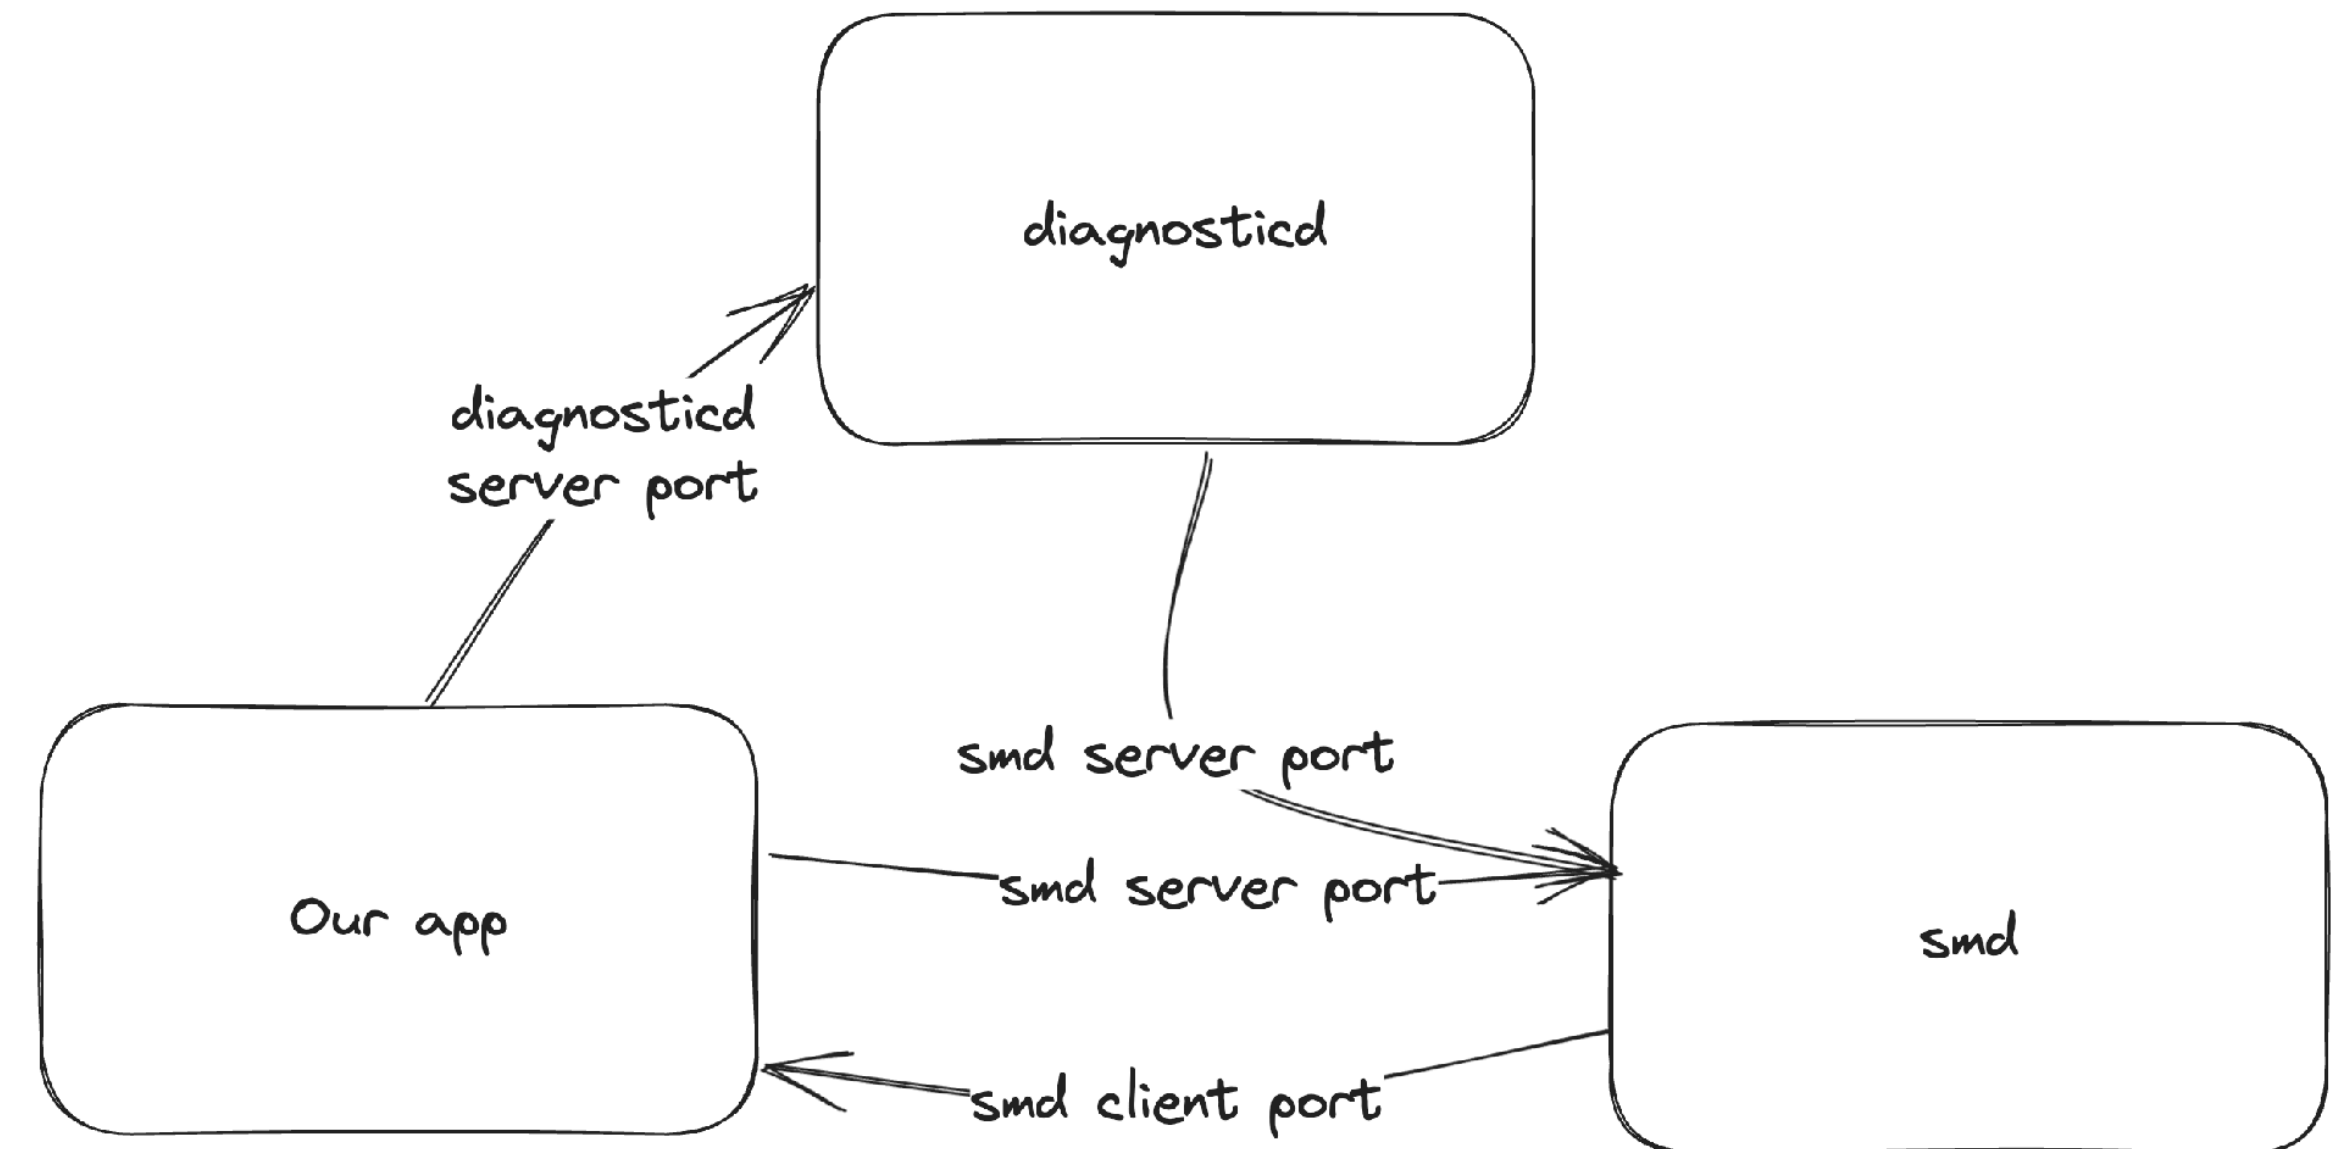
Second service: diagnosticsd

- > `diagnosticd` is another daemon, which can be used to get live diagnostics of a process.
- > It runs as root, so it can pass the first check in `smd`
- > While monitoring a process, this daemon sends multiple messages per second about what the process is doing.




Exploiting smd

1. We establish a **normal** XPC connection to smd
2. Then, we set up another XPC connection to `diagnostcd`, but instead of a send right to our client port, we **duplicate our send right** to smd's server port and send that.
3. Now `diagnostcd` will start streaming messages to smd instead of us



Exploiting smd

- > Now we spam routine 1004 message to smd. If we are lucky, a message from `diagnostcd` overwrites the audit token at the right moment.
- > If the authorisation check uses `diagnostcd`'s audit token, smd thinks we are running as root, so our install is allowed and we achieve privesc



Exploiting smd

- > This is a race, but we can retry quickly. smd doesn't care if the authorisation check fails, and doesn't mind receiving many unrecognised messages from `diagnosticd`
- > Within a few seconds, we install our privileged helper tool and elevate privileges!
- > Since Ventura, the user gets a notification of this. But they are common, hard to identify and our payload has already executed
- > Also works from a sandboxed app (kinda)

Variant 1

- > To summarise the requirements:
 - > Two mach services: the **exploited** service and the **impersonated** service
 - > Our process must be allowed to communicate with **both**
 - > The exploited service must have an authorisation check we cannot pass, but the impersonated service can
 - > The impersonated service can be asked to send normal messages (not replies)
 - > The authorisation check must be performed **asynchronously** from the XPC event handler

Reply port forwarding

Variant 2

Other instances?

- > Second variant: **reply** overwriting the audit token of a **normal message**
- > XPC reply messages (typically) get handled on a different dispatch queue than normal messages
- > Needs the mach service to send a message to my app which expects a reply
- > Concurrently receiving a normal message and a reply can overwrite the audit token of the message with that of the reply, **even if the message is handled synchronously**

Impersonated

Our app

Exploited

Message
(reply port p)

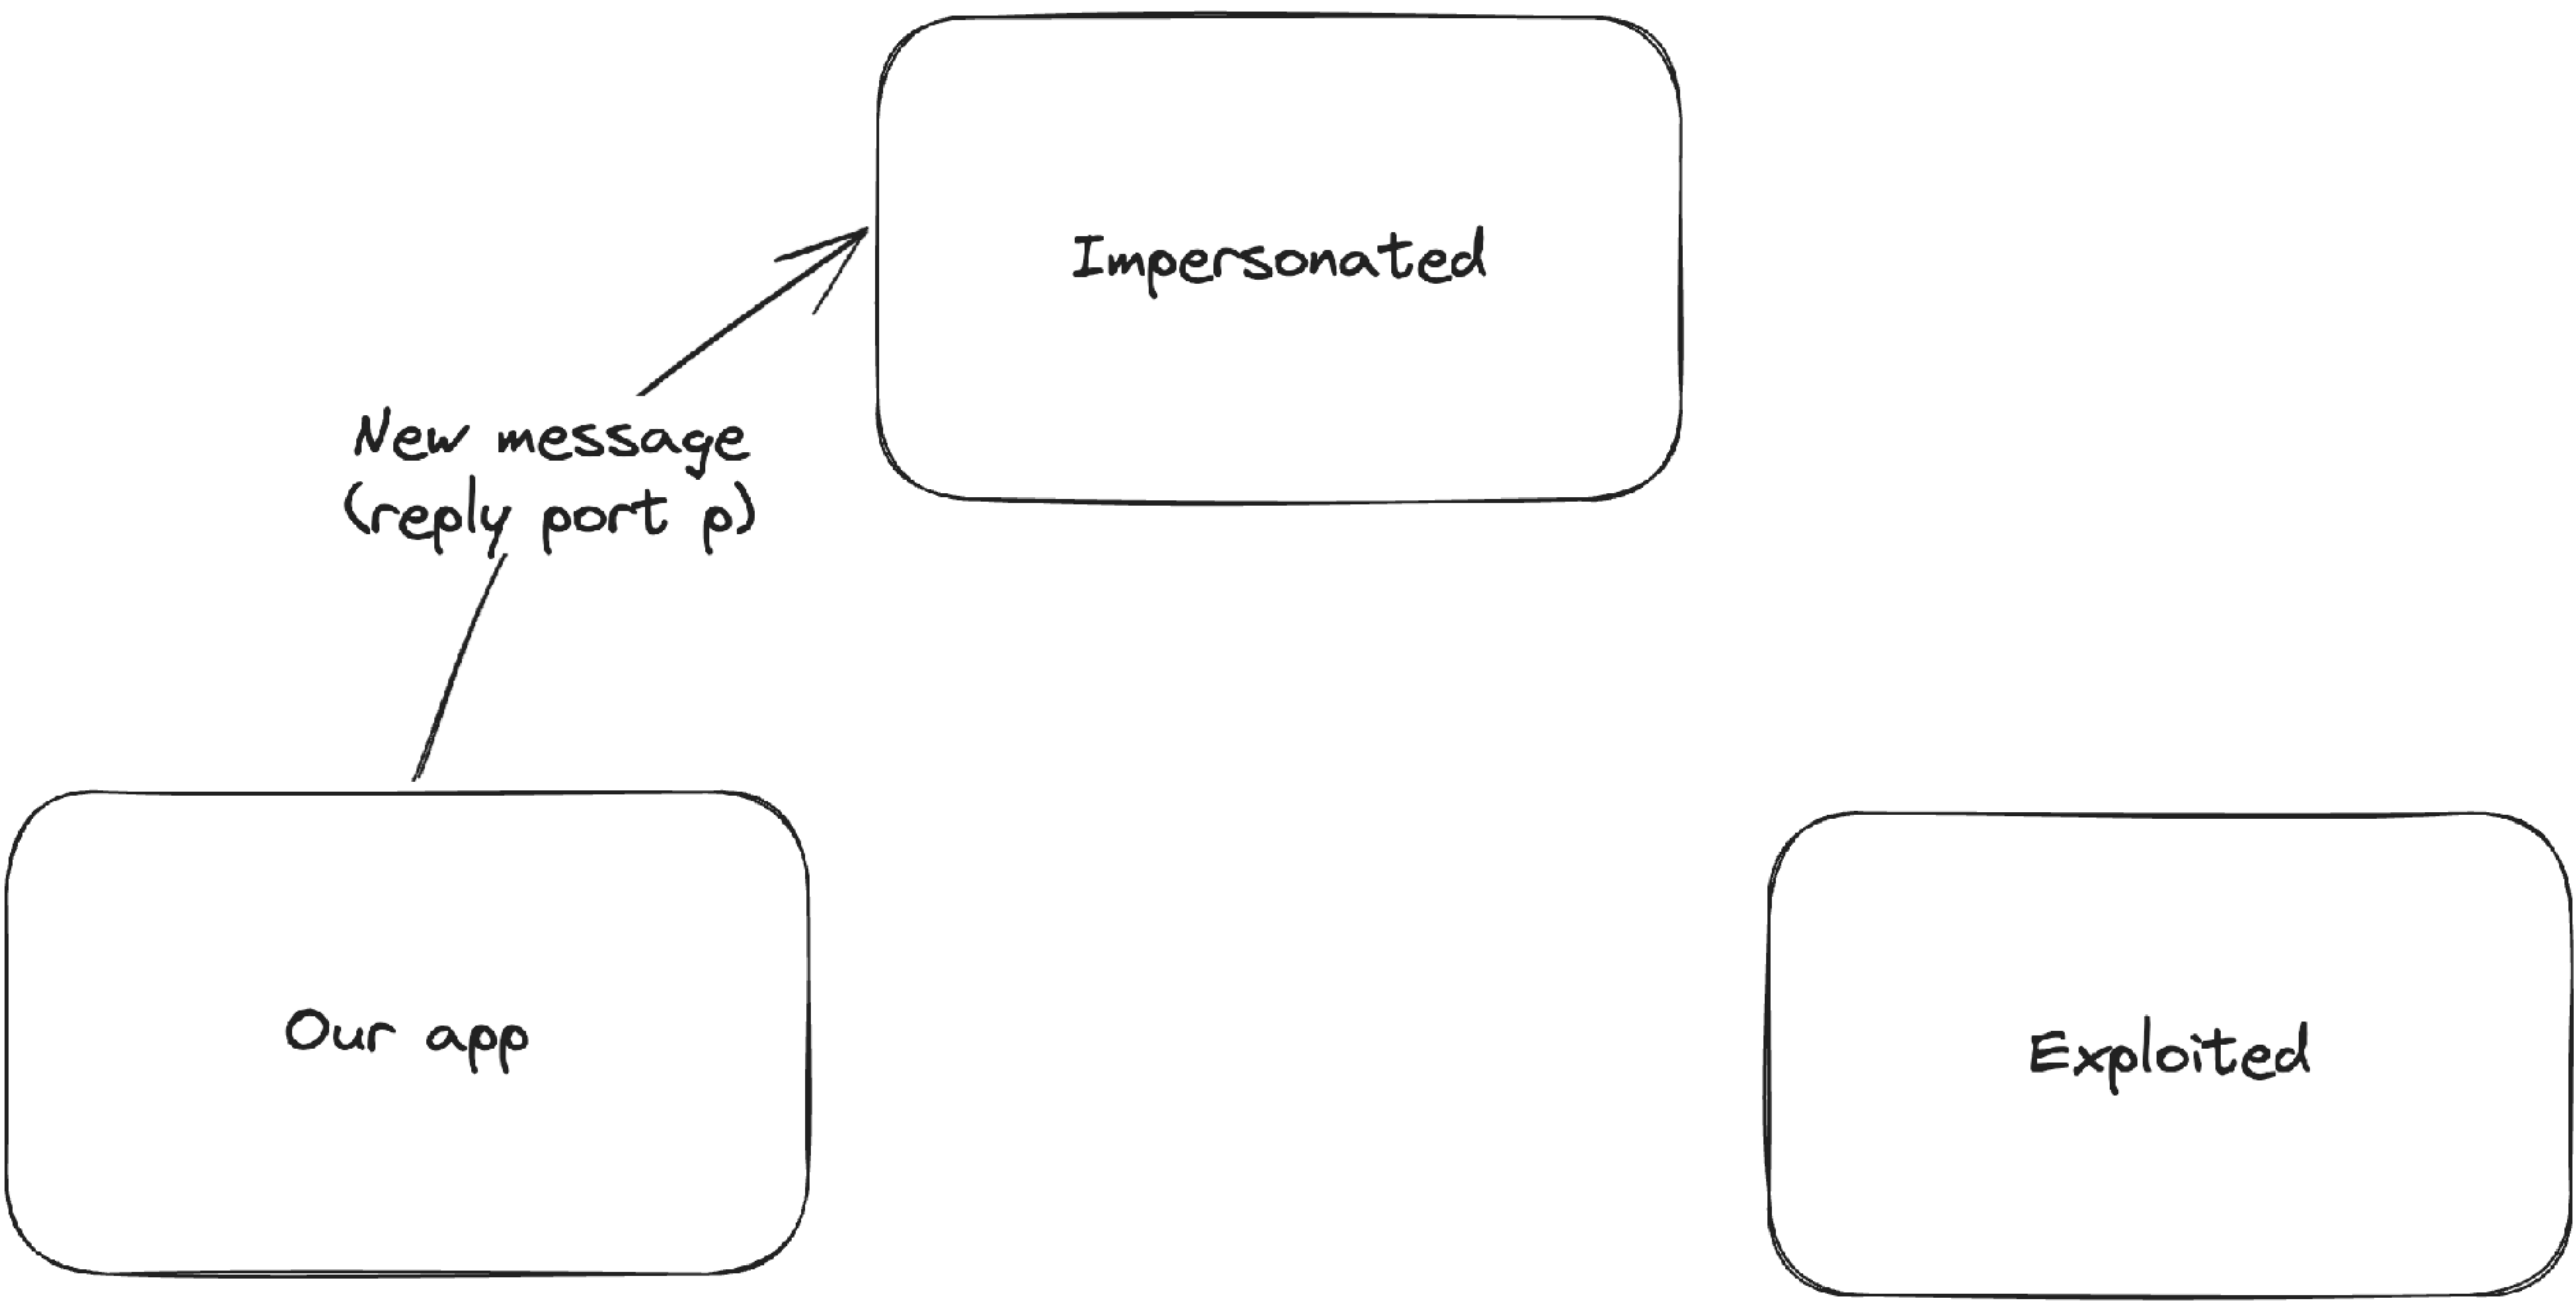



New message
(reply port p)

Impersonated


Our app

Exploited





Reply



Our app

Exploit message



Exploited



Other instances?

- > No instances found, but replicated with two custom launch agents

| | |
|---------------------|---|
| test_mach_service | Reply received: <OS_xpc_dictionary: dicti |
| test_mach_service | *** AUDIT TOKEN CHANGED!!! |
| test_mach_service | Got message: <OS_xpc_dictionary: dictiona |
| second_mach_service | Message received: <OS_xpc_dictionary: dic |
| test_mach_service | Reply received: <OS_xpc_dictionary: dicti |
| test_mach_service | Got message: <OS_xpc_dictionary: dictiona |
| second_mach_service | Message received: <OS_xpc_dictionary: dic |
| test_mach_service | Reply received: <OS_xpc_dictionary: dicti |
| test_mach_service | *** AUDIT TOKEN CHANGED!!! |
| test_mach_service | Got message: <OS_xpc_dictionary: dictiona |
| second_mach_service | Message received: <OS_xpc_dictionary: dic |

Summary

- > Verifying a connection **before accepting** is not vulnerable (e.g. `-listener: shouldAcceptNewConnection:`)
- > XPC event handlers for a connection are never called concurrently
 1. But, if `xpc_connection_get_audit_token` is called **asynchronously** (i.e. not from an event handler), it may get the wrong audit token
 2. Receiving a **reply** concurrently with a **message** can also trigger the audit token to be wrong

More instances

Maybe in iOS?

Other instances?

- > Impact on macOS was demonstrated
- > But one issue doesn't warrant a structural fix, does it also affect iOS?
- > Spent many days (spread out over more than a year) trying to find an instance where it is exploitable on iOS

Other instances?

- > Used **Frida** to look for `xpc_connection_get_audit_token`, used the backtrace to check if not in an XPC event handler
 - > (Using a **Security Research iPhone**, if you're interested, sign-ups are open!)
 - > Only finds functionality that is used and in the monitored daemon
- > Also decompiled a lot of iOS mach services, very time intensive
- > Tried writing a **Ghidra script** to search for paths from `xpc_connection_set_event_handler` to `dispatch_async` and then from that block to `xpc_connection_get_audit_token`...
 - > Parsing blocks and the shared dyld cache made this very challenging
- > Gave up, submitted what we had

The fix

- > Suggested structural fix: **drop a message if the audit token is not equal to the saved audit token**
 - > UID may change, but PID and PID version really should **never** change
- > Apple's fix: changed the vulnerable call to `xpc_connection_get_audit_token` to `xpc_dictionary_get_audit_token` in `smd`
 - > `xpc_dictionary_get_audit_token` uses the audit token from the mach message's trailer, which is safe
- > Issue may still affect other services, so... good luck anyone else who wants to look at this!

Hardening advice for XPC services

- > Drop the connection after receiving an unrecognised message
- > Perform authorisation checks **before accepting** a connection wherever possible
 - > If not, use `xpc_dictionary_get_audit_token`
 - > Alternatively, save the audit token during the accept handler and use that (works for `NSXPCConnection` too)
- > There are new APIs to automatically perform a codesigning check before accepting:
 - > – `[NSXPCConnection setCodeSigningRequirement:]` (since macOS 13.0)
 - > `xpc_connection_set_peer_code_signing_requirement` (since macOS 12.0)

Summary

- > XPC is a one-to-one communication channel, built on mach messages, which are a multiple sender, single receiver communication channel
- > We can create an XPC connection with multiple senders. The audit token for such a connection becomes unreliable
- > We have shown how we exploited this to achieve privilege escalation on macOS using `smd`
- > Apple has not addressed this structurally, so this may still affect other services and platforms

Want to read our full write-up?



SECT<7>R
powered by Computest

sector7.computest.nl

Bonus: sandbox escape

- > The smd exploit still works if sandbox is enabled! 🥳
 - > But it does require including the to be installed helper tool in the app's bundle and Info.plist...
- > Exploiting existing software won't work
- > Mac App Store review will almost certainly reject this (easy static check)
- > Would only help to convince users to download and run the app: "it's safe, just look at the sandbox entitlement!"

Bonus: sandbox escape

- > More generally: the sandbox prevents mach service **lookups**, so if you transfer the mach ports of an XPC connection to a sandboxed process, it can use it fine!
- > Suppose you have:
 - > Unsandboxed code execution as a normal user
 - > Sandboxed code execution as root
- > Then opening a connection as the normal user and moving it to the root process could offer certain opportunities
 - > In smd routine 1004 doesn't work anymore, but maybe 1000 does something interesting?