

IOKit Fundamentals

Contents

Introduction to I/O Kit Fundamentals 9

[Who Should Read This Document?](#) 9

[Organization of This Document](#) 9

[See Also](#) 11

What Is the I/O Kit? 13

[Before You Begin](#) 13

[I/O Kit Features](#) 14

[Design Principles of the I/O Kit](#) 15

[Limitations of the I/O Kit](#) 16

[Language Choice](#) 16

[Using Namespaces in an I/O Kit Driver](#) 17

[Using Static Constructors in an I/O Kit Driver](#) 17

[The Parts of the I/O Kit](#) 18

[Frameworks and Libraries](#) 18

[Applications and Tools](#) 19

[Other I/O Kit Resources](#) 20

[Should You Program in the Kernel?](#) 21

[When Code Should Reside in the Kernel](#) 21

[Alternatives to Kernel-Resident Code](#) 22

Architectural Overview 23

[Driver Layering](#) 23

[Families and Drivers](#) 24

[Drivers and Nubs](#) 24

[The Anatomy of an I/O Connection](#) 25

[The Runtime Environment of Device Drivers](#) 27

[Runtime Features](#) 27

[Kernel Programming Constraints](#) 28

[The I/O Registry and the I/O Catalog](#) 29

[Driver Matching](#) 30

[The I/O Kit Class Hierarchy](#) 31

[The OS Classes](#) 32

[The General I/O Kit Classes](#) 33

The I/O Kit Family Classes	34
Controlling Devices From Outside the Kernel	35
The Device-Interface Mechanism	36
POSIX Device Files	38
The I/O Registry	39
I/O Registry Architecture and Construction	39
The I/O Registry Explorer	41
Driver and Device Matching	43
Driver Personalities and Matching Languages	43
Driver Matching and Loading	48
Driver Matching	48
Device Probing	49
Driver Loading	50
Device Matching	51
The Base Classes	52
The libkern Base Classes	53
Object Creation and Disposal (OSObject)	54
Runtime Type Information (OSMetaClass)	55
Defining C++ Classes in libkern	58
The I/O Kit Base Classes	61
Dynamic Driver Registration (IORegistryEntry)	62
Basic Driver Behavior (IOService)	63
I/O Kit Families	69
Drivers and Families	69
Families As Libraries	71
Library Versioning	72
Library Loading	72
The Programmatic Structure of Families	73
Typical Classes	74
Naming and Coding Conventions	74
Creating An I/O Kit Family	75
Handling Events	77
Work Loops	77
Work Loop Architecture	78
Shared and Dedicated Work Loops	79

Examples of Obtaining Work Loops	80
Event Sources	82
Handling Interrupts	83
Handling Timer Events	89
I/O Requests and Command Gates	90

Managing Data 95

Handling I/O Transfers	95
Memory Descriptors and Memory Cursors	96
Memory in an I/O Request	97
Issues With 64-Bit System Architectures	100
Relaying I/O Requests	103
More on Memory Descriptors	103
More on Memory Cursors	104
DMA and System Memory	105
Dealing With Hardware Constraints	106
IOMemoryCursor Subclasses	107

Managing Power 108

Power Events	109
The Power Plane: A Hierarchy of Power Dependencies	109
Devices and Power States	111
Deciding How to Implement Power Management in Your Driver	112
Implementing Basic Power Management	113
Implementing Advanced Power Management	117
Defining and Using Multiple Power States	117
Changing the Power State of a Device	120
Implementing Idleness Determination and Idle Power Saving	123
Receiving Notification of Power-State Changes in Other Devices	125
Receiving Shutdown and Restart Notifications	127
Keeping Power On for Future Device Attachment	128

Managing Device Removal 129

The Phases of Device Removal	129
Making Drivers Inactive	130
Clearing I/O Queues	131
Detaching and Releasing Objects	131

Base and Helper Class Hierarchy 132

Bibliography 133

System Internals 133

Websites - Online Resources 133

Glossary 135

I/O Kit Family Reference 142

ADB 142

ATA and ATAPI 143

Audio 145

FireWire 146

Graphics 148

 A Note on NDRV Compatibility 150

HID 151

Network 152

PC Card 156

PCI and AGP 157

SBP-2 158

SCSI Parallel 160

SCSI Architecture Model 161

Serial 165

Storage 166

 IOMedia Filter Schemes 169

 IOMedia Properties 170

 Accessing IOMedia From Applications 171

USB 172

Devices Without I/O Kit Families 176

 Imaging Devices 176

 Digital Video 176

 Sequential Access Devices (Tape Drives) 176

 Telephony Devices 177

 Vendor-Specific Devices 177

Document Revision History 178

Index 181

Figures, Tables, and Listings

What Is the I/O Kit? 13

- Table 1-1 Frameworks and libraries of the I/O Kit 19
- Table 1-2 Applications used in driver development 19
- Table 1-3 Command-line tools used in driver development 20

Architectural Overview 23

- Figure 2-1 Driver objects as clients and providers 25
- Figure 2-2 Driver objects in a connection for a SCSI disk driver 26
- Figure 2-3 I/O Kit extended class hierarchy 32
- Figure 2-4 An application controlling a SCSI device through a device interface. 37

The I/O Registry 39

- Figure 3-1 Two planes in the I/O Registry 40
- Figure 3-2 A sample I/O Registry Explorer window 42

Driver and Device Matching 43

- Listing 4-1 A partial listing of an XML personality for an Ethernet controller 44
- Listing 4-2 Driver personalities for the AppleUSBAudio driver 46

The Base Classes 52

- Figure 5-1 The base classes of the I/O Kit class hierarchy 52
- Figure 5-2 Driver object life-cycle functions 64
- Table 5-1 OSMetaClass type-casting and introspection APIs 58
- Listing 5-1 Implementing an init method 59
- Listing 5-2 Creating an instance and calling its init method 60
- Listing 5-3 Implementing the free function 61

I/O Kit Families 69

- Figure 6-1 A driver's relationships with I/O Kit families 70
- Figure 6-2 OSBundleLibraries and the dependency tree 73
- Table 6-1 API prefixes reserved by Apple 75
- Listing 6-1 The OSBundleLibraries property 71

Handling Events 77

Figure 7-1	Driver objects sharing a work loop	80
Figure 7-2	A work loop and its event sources	85
Listing 7-1	Creating a dedicated work loop	81
Listing 7-2	Adding an interrupt event source to a work loop	85
Listing 7-3	Disposing of an <code>IOInterruptEventSource</code> object	86
Listing 7-4	Setting up an <code>IOFilterInterruptEventSource</code>	87
Listing 7-5	Creating and registering a timer event source	89
Listing 7-6	Disposing of a timer event source	90
Listing 7-7	Creating and registering a command gate	91
Listing 7-8	Issuing an I/O request through the command gate	92
Listing 7-9	Disposing of an <code>IOCommandGate</code> object	93

Managing Data 95

Figure 8-1	The role of the user client in an I/O command	98
Figure 8-2	The principal I/O Kit objects in an I/O transfer	99
Table 8-1	Subclasses of <code>IOMemoryDescriptor</code>	104
Table 8-2	Apple-provided subclasses of <code>IOMemoryCursor</code>	107

Managing Power 108

Figure 9-1	The power plane shown in I/O Registry Explorer	110
Table 9-1	Fields and appropriate values in the <code>IOPMPowerState</code> structure	117
Table 9-2	Power flags that describe device capabilities	118
Listing 9-1	Building the power-state array and registering the driver	119
Listing 9-2	Getting notification of system shutdown or restart	128

Managing Device Removal 129

Figure 10-1	Phases of device removal	130
-------------	--------------------------	-----

I/O Kit Family Reference 142

Figure A-1	Storage family driver stack	167
Table A-1	Clients and providers of the ADB family	143
Table A-2	Clients and providers of the ATA and ATAPI family	144
Table A-3	Clients and providers of the Audio family	146
Table A-4	Clients and providers of the FireWire family	148
Table A-5	Clients and providers of the Graphics family	149
Table A-6	Clients and providers of the HID family	152
Table A-7	Clients and providers of the Network family	154
Table A-8	Clients and providers of the PCI and AGP family	158
Table A-9	Clients and providers of the SBP2 family	159
Table A-10	Clients and providers of the SCSI Parallel family	160

Table A-11	SCSI Architecture Model family—Transport Driver layer	161
Table A-12	Clients and providers of the Serial family	166
Table A-13	Storage family (IOMedia) properties	170
Table A-14	Clients and providers of the USB family	175

Introduction to I/O Kit Fundamentals

This document explains the terminology, concepts, architecture, and basic mechanisms of the I/O Kit, Apple's object-oriented framework for developing device drivers for OS X. It contains essential background information for anyone wanting to create device drivers for this platform.

Who Should Read This Document?

There are two general types of I/O Kit developers, and this document tries to be useful to both. The first type is the developer creating a device driver that is to be resident in the kernel; the second type is the application developer who is using an I/O Kit device interface to communicate with hardware. Some chapters contain information useful to both types of developers, and others contain information that is of interest only to writers of kernel-resident drivers.

Obviously there are things *I/O Kit Fundamentals* does not cover. It does not, for example, describe the use of the development tools or the use of specific driver programming interfaces. But it does help you to understand the hows and whys of the I/O Kit, enabling you to obtain the most value from the more specific documentation and examples.

Organization of This Document

I/O Kit Fundamentals gives a broad, conceptual description of the I/O Kit and device-driver development on OS X. It contains the following chapters:

- [What Is the I/O Kit?](#) (page 13)
Describes the features and benefits of the I/O Kit, and also discusses the philosophy and decisions informing its design.
- [Architectural Overview](#) (page 23)
Gives a high-level description of the I/O Kit's architecture, essential concepts, and basic mechanisms.
- [The I/O Registry](#) (page 39)
Describes the I/O Registry, a dynamic database capturing the client/provider relationships among active driver objects.
- [Driver and Device Matching](#) (page 43)

Explains the matching process by which the most appropriate client drivers are found for registered providers. It also summarizes the procedure processes in user space follow to find suitable devices and their drivers.

- [The Base Classes](#) (page 52)

Describes the base classes that each driver object directly or indirectly inherits from. It includes discussions of object construction and disposal, driver objects as I/O Registry entries, and the driver life cycle.

- [Handling Events](#) (page 77)

Explains the architecture and usage of work loops and event sources, mechanisms that the I/O Kit uses to process events such as interrupts and I/O requests in a protected single-threaded environment.

- [Managing Data](#) (page 95)

Describes how to use memory cursors, memory descriptors, and related objects to handle I/O transfers. It also discusses how drivers should deal with hardware constraints, such as those imposed by DMA engines.

- [Managing Power](#) (page 108)

Explains the concepts of OS X power management and describes different ways drivers can power-manage their devices.

- [Managing Device Removal](#) (page 129)

Explains how to respond to device removal (hot-swapping).

- [I/O Kit Family Reference](#) (page 142)

Displays a class hierarchy chart for each family and provides family-specific information that might differ from generic I/O Kit information.

- [Base and Helper Class Hierarchy](#) (page 132)

Provides a class hierarchy chart for all I/O Kit classes that are not members of a specific family.

- [Document Revision History](#) (page 178)

Lists changes to this document.

- [Bibliography](#) (page 133)

Lists additional sources for information on OS X and related topics.

- [Glossary](#) (page 135)

Defines key terms used in this document.

See Also

Once you've absorbed the information in *I/O Kit Fundamentals*, you should be able to forge ahead and actually create a device driver. Apple provides several documents and other sources of information to help you with your efforts:

- *IOKit Device Driver Design Guidelines* describes the general steps required to design, code, debug, and build a device driver that will be resident in the kernel.
- *Accessing Hardware From Applications* discusses how to use the I/O Kit's "device interface" feature; it also includes information on serial and storage I/O via BSD device files.
- *Kernel Extension Programming Topics* contains a collection of tutorials that introduce you to the development tools and take you through the steps required to create, debug, and package kernel extensions and I/O Kit drivers (a type of kernel extension). It also includes information on other aspects of kernel extensions.
- Documentation that provides in-depth information on writing drivers for specific driver families and related reference documentation is available in [Hardware & Drivers Documentation](#).
- *Kernel Programming Guide* provides an overview of the architecture and components of the OS X kernel environment (Mach, BSD, networking, file systems, I/O Kit). All developers who intend to program in the kernel (including device-driver writers) should read this document.
- *Mac Technology Overview* provides an introduction to OS X as a whole, which is useful for developers new to the platform.

Of course, you can always browse the header files shipped with the I/O Kit, which are installed in `Kernel.framework/Headers/iokit` (kernel-resident) and `IOKit.framework/Headers` (device interface).

You can also view developer documentation in Xcode. To do this, select Help from the Xcode menu and then click Show Documentation Window.

You can browse the BSD man pages for more information on BSD and POSIX APIs in two ways: You can type `manfunction_name` in a Terminal window (for example, `man gdb`) or you can view an HTML version at *OS X Man Pages*.

If you're ready to develop a universal binary version of a device driver to run in an Intel-based Macintosh, first read *Universal Binary Programming Guidelines, Second Edition*. Then, see *IOKit Device Driver Design Guidelines* for an overview of issues of particular interest to device driver developers. Related information that is specific to a particular device type is available in the documents listed at [Hardware & Drivers Documentation](#).

Apple maintains several websites where developers can go for general and technical information on OS X.

- Apple Developer Connection Reference Library (<http://developer.apple.com/referencelibrary/index.html>) contains a comprehensive collection of technical resources, including documentation, sample code, and Technical Notes.
- Apple Developer Connection: OS X (<http://developer.apple.com/devcenter/macosx>) offers SDKs, release notes, product notes and news, and other resources and information related to OS X.
- The AppleCare Support site (<http://www.apple.com/support>) provides a search feature that enables you to locate technical articles, manuals, specifications, and discussions on OS X and other areas.

What Is the I/O Kit?

The I/O Kit is a collection of system frameworks, libraries, tools, and other resources for creating device drivers in OS X. It is based on an object-oriented programming model implemented in a restricted form of C++ that omits features unsuitable for use within a multithreaded kernel. By modeling the hardware connected to an OS X system and abstracting common functionality for devices in particular categories, the I/O Kit streamlines the process of device-driver development.

This chapter talks about the inherent capabilities of the I/O Kit (and of the drivers developed with it), about the decisions informing its design, and about the I/O Kit when considered as a product. It also offers some caveats and guidelines for those considering developing kernel software such as device drivers.

Before You Begin

You might have developed device drivers for other platforms—Mac OS 9, perhaps, or BSD or another flavor of UNIX. One thing you’ll discover reading this document is how different the approach is with the I/O Kit. Although writing drivers for OS X requires new ways of thinking and different ways of programming, you are amply rewarded for shifting to this new approach. The I/O Kit simplifies driver development and supports many categories of devices. Once you get the basics of the I/O Kit down, you’ll find it a relatively easy and efficient matter to create device drivers.

Before you attempt driver development with the I/O Kit, Apple highly recommends certain prerequisites. Because the framework uses an object-oriented programming model, which is implemented in a restricted subset of C++, it helps to know C++ or object-oriented concepts in general. Also, device drivers are not the same thing as applications because, being kernel-resident, they must abide by more restrictive rules. Knowledge of kernel programming is therefore very useful.

Indeed, programming in the kernel is discouraged except when it is absolutely necessary. Many alternatives for communicating with hardware and networks exist at higher levels of the system, including the “device interface” feature of the I/O Kit described in [Controlling Devices From Outside the Kernel](#) (page 35) See [Should You Program in the Kernel?](#) (page 21) for more on alternatives to kernel programming.

I/O Kit Features

From its inception, the fundamental goal for the I/O Kit has been to accommodate and augment native features and capabilities of OS X, particularly those of the kernel environment. As the driver model for OS X, the I/O Kit supports the following features:

- Dynamic and automatic device configuration (plug-and-play)
- Many new types of devices, including graphics acceleration and multimedia devices
- Power management (for example, “sleep” mode)
- The kernel’s enforcement of protected memory—separate address spaces for kernel and user programs
- Preemptive multitasking
- Symmetric multiprocessing
- Common abstractions shared between types of devices
- Enhanced development experience—new drivers should be easy to write

The I/O Kit supports these kernel features with its new model for device drivers and adds some additional features:

- An object-oriented framework implementing common behavior shared among all drivers and types (families) of drivers
- Many families for developers to build upon
- Threading, communication, and data-management primitives for dealing with issues related to multiprocessing, task control, and I/O-transfers
- A robust, efficient match-and-load mechanism that scales well to all bus types
- The I/O Registry, a database that tracks instantiated objects (such as driver instances) and provides information about them
- The I/O Catalog, a database of all I/O Kit classes available on a system
- A set of device interfaces—plug-in mechanisms that allows applications and other software in “user space” to communicate with drivers
- Excellent overall performance
- Support for arbitrarily complex layering of client and provider objects

The I/O Kit's object-oriented programming model is implemented in a restricted subset of C++. Object-orientation just in itself is an advantage in driver development because of the code reusability it fosters. Once you are familiar with the I/O Kit, you can write device drivers much more quickly and efficiently than you can using a procedural model. In addition, code reusability decreases the memory footprint of drivers; drivers ported from Mac OS 9, for example, have been up to 75% smaller in OS X.

Design Principles of the I/O Kit

OS X is largely the product of two strains of operating-system technology: Mac OS 9 (and its predecessors) and BSD. Given this pedigree, one might have expected Apple to adopt the device-driver model of Mac OS 9 or FreeBSD. Instead, Apple chose to redesign the model. Several reasons motivated this decision.

First, neither the Mac OS 9 driver model nor the FreeBSD driver model offers a set of features rich enough to meet the needs of OS X. The OS X kernel is significantly more advanced than its Mac OS precursors; it handles memory protection, preemptive multitasking, multiprocessing, and other features not present in previous versions of Mac OS. Although FreeBSD is capable of handling these features, the BSD model does not offer other features expected in a modern operating system, including automatic configuration, driver stacking, power management, and dynamic loading of devices.

Thus the primary motivation behind the I/O Kit was the inadequacy of the currently available driver models. The redesign of the I/O architecture had to take advantage of and support the operating-system features of OS X. Toward this end, the I/O Kit's designers settled on an object-oriented programming model that abstracted the kernel capabilities and hardware of an OS X system and provided a view of this abstraction to the upper layers of the operating system. The compelling part of this abstraction is the implementation of behavior common to all device drivers (or types of device drivers) in the classes of the I/O Kit.

As an example, consider virtual memory. In Mac OS 9, virtual memory is not a fundamental part of the operating system; it is an option. Because of this, a developer must always take virtual memory into account when creating a driver, and this raises certain complications. In contrast, virtual memory is an inherent capability of OS X and cannot be turned off. Because virtual memory is a fundamental and assumed capability, knowledge of it is incorporated into system software and driver writers do not have to take it into account.

The I/O Kit functions as a kind of foundation and coordinator for device drivers. This is a departure from previous driver models. In Mac OS 9, all software development kits (SDKs) are independent of each other and duplicate common functionality. OS X delivers the I/O Kit as part of a single kernel development kit (KDK); all portions of the KDK rest on common underpinnings. OS X helps developers take advantage of hardware complexity without requiring them to encode software complexity into each new device driver. In most cases, they need only add the specific code that makes their drivers different.

Another part of the design philosophy of the I/O Kit is to make the design completely open. Rather than hiding APIs in an attempt to protect developers from themselves, all I/O Kit source code is available as part of Darwin. Developers can use the source code as an aid to designing (and debugging) new drivers.

Limitations of the I/O Kit

Although the I/O Kit supports most types of hardware on an OS X system, it does not fully support all hardware. One category of such devices are those used for imaging, among them printers, scanners, and digital cameras. The I/O Kit provides only limited support for these devices, handling communication with these devices through the FireWire and USB families. Applications or other programs in user space are responsible for controlling the particular characteristics of these devices (see [Controlling Devices From Outside the Kernel](#) (page 35) for details). If your application needs to drive an imaging device, you should use the appropriate imaging software development kit (SDK).

Although the I/O Kit attempts to represent the hierarchy and dynamic relationships among hardware devices and services in an OS X system, some things are difficult to abstract. It is in these gray areas of abstraction, such as when layering violations occur, that driver writers are more on their own. Even when the I/O Kit representation is clean and accurate, the reusability of I/O Kit family code can be limited. All hardware can have its own quirks and a driver's code must take these quirks into account.

Language Choice

Apple considered several programming languages for the I/O Kit and chose a restricted subset of C++.

C++ was chosen for several reasons. The C++ compiler is mature and the language provides support for system programming. In addition, there is already a large community of Macintosh (and BSD) developers with C++ experience.

The restricted subset disallows certain features of C++, including

- Exceptions
- Multiple inheritance
- Templates
- Runtime type information (RTTI)—the I/O Kit uses its own implementation of a runtime typing system

These features were dropped because they were deemed unsuitable for use within a multithreaded kernel. If you feel you need these features, you should reconsider your design. You should be able to write any driver you require using I/O Kit with these restrictions in place.

Using Namespaces in an I/O Kit Driver

Note that you can use namespaces in your I/O Kit driver. The use of namespaces can help you avoid name collisions and may make your code easier to read and more maintainable. Be sure to use reverse-DNS format for the namespace name (for example, `com.mycompany`) to avoid potential namespace collisions.

If you decide to use namespaces in your in-kernel I/O Kit driver, do not declare any subclass of `OSObject` in a namespace or your driver will not load. At present, the loader does not support `OSObject`-derived classes that require qualification, such as the one shown below:

```
namespace com.mycompany {  
    class com.mycompany.driver.myClass : public IOService { // This is not allowed.  
        OSDeclareDefaultStructors (com.mycompany.driver.myClass);  
    };  
};
```

Using Static Constructors in an I/O Kit Driver

In OS X v10.4, GCC 4.0 is the default compiler for all new projects, including I/O Kit drivers. This section describes a particular difference between GCC 3.3 and GCC 4.0 that may affect the compatibility of your in-kernel driver between OS X v10.3.x and OS X v10.4.x. For more information on the differences between GCC 3.3 (the default compiler in OS X v10.3) and GCC 4.0, including porting guidance, see *GCC Porting Guide*.

If you perform static construction within a function in a C++ I/O Kit driver (or other KEXT) compiled with GCC 3.3 or earlier, be aware that the same KEXT compiled with GCC 4.0 will no longer load successfully. This is because GCC 4.0 is more strict about taking and releasing locks in the kernel environment. If you perform in-function static construction in your I/O Kit driver compiled with GCC 4.0, you will probably see the following error when you try to load it:

```
kld():Undefined symbols:  
__cxa_guard_acquire  
__cxa_guard_release
```

The solution to this problem is simple: move the static constructor to a global namespace. For example, suppose that your I/O Kit driver includes an in-function static construction, such as in the code shown below:

```
class com_mycompany_driver_mystaticclass;  
void com_mycompany_driver_myclass::myfunction(void)  
{
```

```
static com_mycompany_driver_mystaticclass staticclass;
staticclass.anotherfunction();
}
```

You can avoid loading errors by changing this code to avoid in-function static construction, as in the code shown below:

```
class com_mycompany_driver_mystaticclass;
static com_mycompany_driver_mystaticclass staticclass;
void com_mycompany_driver_myclass::myfunction(void)
{
    staticclass.anotherfunction();
}
```

Note that you may be able to avoid the load errors associated with in-function static construction without changing your code if you compile your KEXT with GCC 4.0 using the `-fno-threadsafe-statics` compiler option, but this may lead to other problems. Specifically, unless you can guarantee thread safety in other ways, compiling your KEXT with this option may break your code.

The Parts of the I/O Kit

Physically and electronically, the I/O Kit is composed of many parts: frameworks and libraries, development and testing tools, and informational resources such as example projects, documentation, and header files. This section catalogs these parts and indicates where they are installed and how they can be accessed.

Frameworks and Libraries

The I/O Kit is based on three C++ libraries. All of them are packaged in frameworks, but only `IOKit.framework` is a true framework. The Kernel framework exists primarily to expose kernel header files, including those of `libkern` and `IOKit`. The code of these “libraries” is actually built into the kernel; however, drivers (when loaded) do link against the kernel as if it were a library.

Table 1-1 Frameworks and libraries of the I/O Kit

Framework or library	Description and location
Kernel/IOKit	The library used for developing kernel-resident device drivers. Headers location: <code>Kernel.framework/Headers/IOKit</code>
Kernel/libkern	The library containing classes useful for all development of kernel software. Headers location: <code>Kernel.framework/Headers/libkern</code>
IOKit	The framework used for developing device interfaces. Location: <code>IOKit.framework</code>

Applications and Tools

You use a handful of development applications to build, manage, debug, examine, and package device drivers. [Table 1-2](#) (page 19) lists the applications used in driver development; these applications are installed in `/Developer/Applications`.

Table 1-2 Applications used in driver development

Application	Description
Xcode	The primary development application for OS X. Xcode manages projects, provides a full-featured code editor, builds projects according to arbitrarily complex rules, provides a user interface for software configuration, and acts as a front end for debugging and documentation searches.
I/O Registry Explorer	Enables the graphical exploration of the contents and structure of the I/O Registry.
Package Maker	Creates an installation package for the Installer application; used for deployment of kernel extensions (including device drivers).

[Table 1-3](#) (page 20) describes the command-line tools used in developing device drivers with the I/O Kit; all tools are located in `/usr/sbin/` or `/sbin`.

Note: You can view on-line documentation of these tools (called man pages in the UNIX world) by entering a command in the shell provided by the Terminal application. The command is `man`, and the main argument to the `man` command is the name of the tool for which you want to see documentation. For example, to see the man page for the `kextload` tool, enter the following line in Terminal:

```
man kextload
```

Table 1-3 Command-line tools used in driver development

Tool	Description and location
<code>ioreg</code>	Prints the contents of the I/O Registry (a command-line version of the I/O Registry Explorer application).
<code>kextload</code>	Loads a kernel extension (such as device driver) or generates a statically linked symbol file for remote debugging.
<code>kextunload</code>	Unloads a kernel extension (if possible).
<code>kextstat</code>	Prints statistics about currently loaded drivers and other kernel extensions.
<code>iostat</code>	Displays kernel I/O statistics on terminal, disk, and CPU operations.
<code>ioclasscount</code>	Displays instance count of a specified class.
<code>ioalloccount</code>	Displays some accounting of memory allocated by I/O Kit objects in the kernel.
<code>kextcache</code>	Compresses and archives kernel extensions (including drivers) so they can be automatically loaded into the kernel at boot time.
<code>gcc</code>	Apple’s version of the GNU C++ compiler; Xcode automatically invokes it with the correct set of flags for I/O Kit projects.
<code>gdb</code>	Apple’s version of the GNU debugger; Xcode automatically invokes it with the correct set of flags for I/O Kit projects.

Other I/O Kit Resources

Several informational resources are included with the I/O Kit “product,” particularly documentation and header files. Some of these resources are described in the preceding chapter, [Introduction to I/O Kit Fundamentals](#) (page 9)

The I/O Kit is part of the Darwin Open Source project. Apple maintains a website where you can find much information related to the I/O Kit and other Open Source projects managed by Apple. The following two locations are of particular interest:

- Open Source Projects—<http://developer.apple.com/darwin/projects/>

Here you can find links to the Darwin and Darwin Streaming projects, among other projects. Also featured are links to documentation and tools.

- Mailing lists—<http://developer.apple.com/darwin/mail.html>

This page features links that will put you on the Darwin-Development and DarwinOS-Users mailing lists, among others.

Should You Program in the Kernel?

If you are thinking of writing code for the kernel environment, think carefully. Programming in the kernel can be a difficult and dangerous task. And often there is a way to accomplish what you want to do without touching the kernel.

Software that resides in the kernel tends to be expensive. Kernel code is “wired” into physical memory and thus cannot be paged out by the virtual memory system. As more code is put into the kernel, less physical memory is available to user-space processes. Consequently, paging activity will probably intensify, thereby degrading system performance.

Kernel code is also inherently destabilizing, much more so than application code. The kernel environment is a single process, and this means that there is no memory protection between your driver and anything else in the kernel. Access memory in the wrong place and the entire system can grind to a halt, a victim of a kernel panic.

Moreover, because kernel code usually provides services to numerous user-space clients, any inefficiencies in the code can be propagated to those clients, thereby affecting the system globally.

Finally, kernel software is a real pain to write. There are subtleties to grapple with that are unknown in the realm of application development. And bugs in kernel code are harder to find than in user-space software.

With all this in mind, the message is clear. It is in everyone's best interest to put as little code as possible into the kernel. And any code that ends up in the kernel should be honed and rigorously tested.

When Code Should Reside in the Kernel

A handful of situations warrant loading a driver or extension into the kernel environment:

- The software is used by the kernel environment itself.

- User-space programs will frequently use the software.
- The software needs to respond directly to primary interrupts (those delivered by the CPU's interrupt controller).

If the software you are writing does not match any of these criteria, it probably doesn't belong in the kernel. If your software is a driver for a disk, a network controller, or a keyboard, it should reside in the kernel. If it is an extension to the file system, it should live in the kernel. If, on the other hand, it is used only now and then by a single user-space program, it should be loaded by the program and reside within it. Drivers for printers and scanners fall into this latter category.

Alternatives to Kernel-Resident Code

Apple provides a number of technologies that might let you accomplish what you want to do and stay out of the kernel. First are the higher-level APIs that give you some hardware-level access. For example, Open Transport is a powerful resource for many networking capabilities, and Quartz Compositor enables you to do some fairly low-level things with the graphics subsystem.

Second, and just as important, is the device-interface technology of the I/O Kit framework. Through a plug-in architecture, this technology makes it possible for your application to interact with the kernel to access hardware. In addition, you can—with a little help from the I/O Kit—use POSIX APIs to access serial, storage, or network devices. See [Controlling Devices From Outside the Kernel](#) (page 35) for a summary of device interfaces and see the document *Accessing Hardware From Applications* for a full discussion of this technology.

Note: Objective-C does not provide device-level I/O services. However, in your Cocoa application, you can call the C APIs for device-level functionality that the I/O Kit and BSD provide. Note that you can view the man pages that document BSD and POSIX functions and tools at *OS X Man Pages*.

Architectural Overview

As you can with any complex system, you can look at the design of the I/O Kit from various angles and at different granularities. This chapter introduces you to the more important architectural elements and conceptual domains of the I/O Kit:

- Hardware modeling, the layering of driver objects, and the roles played by families, drivers, and nubs
- The runtime environment of device drivers
- The I/O Kit Registry and I/O Catalog
- Driver matching
- The I/O Kit class hierarchy
- Device interfaces

Keep in mind that this chapter is an overview, and so the discussion it devotes to each of these topics is intentionally brief. Later chapters cover most of these topics in more detail. In the case of device interfaces, the document *Accessing Hardware From Applications* describes the technology in great detail.

Driver Layering

Central to the design of the I/O Kit is a modular, layered runtime architecture that models the hardware of an OS X system by capturing the dynamic relationships among the multiple pieces—hardware and software—involved in an I/O connection. The layers of the connection, comprising driver objects and the families these objects are members of, are stacked in provider-client relationships.

The chain of interconnected services or devices starts with a computer’s logic board (and the driver that controls it) and, through a process of discovery and “matching,” extends the connection with layers of driver objects controlling the system buses (PCI, USB, and so on) and the individual devices and services attached to these buses.

You can view the layering of driver objects in a running OS X system using the I/O Registry Explorer application, included in the developer version of OS X. The developer version also includes a command-line version of the application, `ioreg`, that you can run in a Terminal window to display current I/O Registry information.

This section examines the I/O Kit’s layered architecture and describes the major constituent elements: families, drivers, and nubs.

Families and Drivers

An I/O Kit family is one or more C++ classes that implement software abstractions common to all devices of a particular type. The I/O Kit has families for bus protocols (such as SCSI Parallel, USB, and FireWire), for storage (disk) devices, for network services (including Ethernet), for human-interface devices (such as mice, keyboards, and joysticks), and for a host of other devices.

A driver becomes a member of a family through inheritance; the driver's class is almost always a subclass of some class in a family. By being a member of a family, the driver inherits the data structures (instance variables) and the behaviors that are common to all members of the family. For example, all SCSI controllers have certain things they must do, such as scanning the SCSI bus; the SCSI Parallel family defines and implements this scanning functionality. Thus, you do not need to include scanning code in your new SCSI controller driver (unless you require a different scanning implementation).

Most I/O Kit development involves writing specific driver classes, each of which inherits from the superclass in the family that provides the functionality the driver requires. A driver for an Ethernet controller, for example, inherits from the `IOEthernetController` superclass in the Network family. The bulk of a driver's interaction with its own family involves implementing member functions that the family invokes. These are typically client configuration and I/O requests. Some families also define objects and functions for your driver to use. The exact nature of these objects and functions depends on the family your driver works with.

However, a driver typically works with two families. In addition to the family a driver is a member of, the driver class must communicate with a nub published by the family for the bus or protocol the device is attached to. A nub (as [Drivers and Nubs](#) (page 24) explains in detail) is an object that defines an access point and communication channel for a given protocol. A family (usually representing a bus such as PCI or USB) acts as the driver's provider through the nub that it creates. The driver uses the nub to attach into the I/O Registry and communicate with its device. As an example, a PCI Ethernet driver would use an `IOPCIDevice` nub from the PCI family to attach to and communicate over the PCI bus. A driver's main interaction with the nub involves the issuing of requests or commands on whatever bus the nub is a client of. A SCSI device driver, for example, issues SCSI command blocks and checks results through the nub.

For more on families, particularly the nature and composition of superclasses in a family, see [The I/O Kit Family Classes](#) (page 34)

Drivers and Nubs

The I/O Kit supports two broad types of driver objects. The first is the nub, an object that defines an access point and communication channel, usually for a given protocol, such as PCI, USB, or Ethernet. The second type is the specific driver for an individual device or service. The specific driver communicates with the hardware, through a nub, to perform I/O operations. Both drivers and nubs in the I/O Kit must inherit from the `IOService` class.

A driver is an I/O Kit object that manages a specific piece of hardware. Drivers are written as kernel extensions and are usually installed in the Extensions folder (at `/System/Library/Extensions`.) See [Kernel Extension Overview](#) in *Kernel Programming Guide* for more information about creating and installing kernel extensions.

When a driver is selected for a device, but before it is loaded into the kernel (as a kernel extension), all required families—in terms of superclasses and their dependencies—are loaded to provide the common functionality for the driver and others of its type. (Of course, if these families have already been loaded, this step is not necessary.) After all requirements for the driver are met, the driver is loaded and instantiated as an object. See [The Anatomy of an I/O Connection](#) (page 25) for an illustration of this process.

A nub is an I/O Kit object that represents a communication channel for a device or logical service and mediates access to the device and service. For example, a nub could represent a bus, a disk, a disk partition, a graphics adaptor, or a keyboard. It might help to think of a nub as the software representation of a device slot or connector. Nubs also provide services such as arbitration, power management, and driver matching (see [The I/O Registry and the I/O Catalog](#) (page 29)).


Nubs act as bridges between two drivers and, by extension, between two families. A driver communicates with a nub (and the nub's family) as its client and may, through its family, publish a nub which finds (by matching) a driver for which it is a provider. Usually a driver publishes one nub for each individual device or service it controls; however, when a driver supports a specific piece of hardware it can act as its own nub.

The Anatomy of an I/O Connection

The I/O Kit's layered architecture models the chain of connections between the system's hardware buses and devices, gathering common functionality into classes your driver can interact with. Each layer is a client of the layer below it and a provider of services to the layer above it. Broad groupings of layers, defined by the I/O Kit families, define the functionality common to a general type of I/O provider, such as networking or PCI bus devices.

Consider [Figure 2-1](#) (page 25) which illustrates a typical layering of client and provider objects for a PCI-based Ethernet controller driver in the Network family.

Figure 2-1 Driver objects as clients and providers




As this diagram shows, your driver typically fits between two families, inheriting from a class in the upper-layer family and using the services of the lower-layer family. In the case of the Ethernet controller, the driver participates in a stack of C++ objects comprising instances of classes from the networking and PCI families:

IONetworkStack (interface managing object)	Connects I/O Kit objects to the BSD networking facilities.
IOEthernetInterface (nub)	Manages device-independent data transmission and reception.
Controller Driver (driver)	Operates the Ethernet controller through the IOPCIDevice object. This object inherits from a networking family class called IOEthernetController.
IOPCIDevice (nub)	Match point for the controller; provides basic PCI bus interaction to the controller.
IOPCIBridge (driver)	Manages the PCI bus. (Other objects provide services to the IOPCIBridge; their specific identities depend on the hardware configuration.)

Another way to look at a stack of driver objects in a typical I/O connection is to consider the stack from a dynamic perspective. In other words, what happens when an OS X system discovers a new device attached to it? How is the stack of driver objects constructed? For this, let's use the example of a SCSI disk drive; the general order of creation or discovery in [Figure 2-2](#) (page 26) is left to right.

Figure 2-2 Driver objects in a connection for a SCSI disk driver



This figure illustrates how a SCSI disk driver, a member of the Storage family, is connected to the PCI bus. As each individual connection is made, the newly created driver or nub is also added to the I/O Registry (described in [The I/O Registry and the I/O Catalog](#) (page 29)). The chain of connections takes place in several steps:

1. The PCI bus controller driver, a member of the PCI family, discovers a PCI device and announces its presence by creating a nub (`IOPCIDevice`).
2. The nub identifies (matches) an appropriate device driver—in this case, a SCSI controller driver—and requests that it be loaded. Loading the SCSI controller driver causes the SCSI Parallel family, and all families that it depends on, to be loaded as well. The SCSI controller driver is given a reference to the `IOPCIDevice` nub.

3. The SCSI controller driver, which is a client of the PCI family and a provider of SCSI Parallel family services, scans the SCSI bus for devices that might be clients of these services. Upon finding such a device (a disk), the driver announces the device's presence by creating a nub (`IIOSCSIDevice`).
4. The nub, by going through the matching procedure, finds a device driver (a disk driver) that is appropriate for the device and requests that this driver be loaded. Loading the disk driver causes the Storage family, and all families that it depends on, to be loaded as well. The disk driver is now a client of the SCSI Parallel family and a member of the Storage family. The disk driver is given a reference to the `IIOSCSIDevice` nub.

In many cases, applications and other “user space” programs can use the I/O Kit's device-interface technology to drive devices (including mass storage devices), obviating the need for a kernel-resident driver. See [Controlling Devices From Outside the Kernel](#) (page 35) for an overview of this plug-in technology.

The Runtime Environment of Device Drivers

The I/O Kit provides a runtime environment with several powerful features for driver writers, including:

- A dynamic, layered driver architecture that allows drivers to be loaded and unloaded at any time and delays reserving costly system resources until they're needed
- Standard facilities for managing data during common I/O operations
- A robust system for protecting access to driver resources during I/O operations, which frees driver writers from having to write their own code to enable and disable interrupts and manage locks on the driver's private resources
- Access to services in the libkern C++ library (on which the I/O Kit itself is based) to manage collections, perform atomic operations, and byte-swap values for use on different kinds of hardware

The following section summarizes each of these features.

Runtime Features

I/O Kit drivers can be loaded and unloaded or activated and deactivated at any time, through events initiated by software—as when networking stacks are brought up and down—and by hardware—as when a USB device is added to or removed from the bus. Nearly all drivers must work within the context of a dynamically changing system. The I/O Kit makes this easier by defining a standard life cycle for driver objects. By implementing a small set of functions, summarized in [The General I/O Kit Classes](#) (page 33) your driver can gracefully handle the addition and removal of devices and services, as well as changes induced by the power-management system.

Nearly all I/O operations require the same preparation in OS X:

- Paging virtual memory into physical memory
- Wiring memory down so it can't be paged out during I/O operations
- Building scatter/gather lists that describe the data buffers to read or write

The I/O Kit provides a set of utility classes to help drivers prepare memory for I/O operations and to build scatter/gather lists, including the `IOMemoryDescriptor` and `IOMemoryCursor` classes. For more information on these facilities, see the chapter [Managing Data](#) (page 95)

Drivers running in a multithreaded system must be able to protect their resources from reentrant or concurrent access. The I/O Kit includes a small set of classes for this purpose. A work loop object runs a dedicated thread and manages a gating mechanism for exclusive access to data. Other objects, called event sources, use the gating mechanism to serialize function calls that access critical resources, closing the work loop gate before invoking the function. For more information on work loops and event sources, see the chapter [Handling Events](#) (page 77)

The libkern C++ library, on which the I/O Kit itself is based, provides services commonly needed by drivers, including:

- Arithmetic and logical operations that are guaranteed to be atomic
- Byte-swapping of values between big-endian and little-endian formats
- Classes for common collections of data such as strings, arrays, and dictionaries

For more information on the libkern classes see [The OS Classes](#) (page 32) as well as the libkern reference documentation installed with the OS X Developer package.

Kernel Programming Constraints

Kernel code is always held resident in physical memory, and cannot be paged out by the virtual memory system. This makes kernel resources much more expensive than application program resources. Your driver should reside in the kernel if:

- It takes primary interrupts (in which case it must live in the kernel)
- Its primary client resides in the kernel; for example, mass storage drivers because file-system stacks reside in the kernel

Drivers for disks, network controllers, and keyboards, for example, reside in the kernel. If your driver is only occasionally used by one user-space program at a time, it should be loaded by the program and reside within it. Drivers for such devices as scanners and printers reside within user-space programs, using the I/O Kit's device-interface mechanism to communicate with devices. For more information on device interfaces, see [Controlling Devices From Outside the Kernel](#) (page 35)

Even if your driver resides in the kernel, you should minimize the amount of kernel-resident code and the amount of processing done by that code. For example, a dedicated application that controls an interrupt-driven hardware device should supply a driver that puts the minimum code in the kernel needed to service the interrupt, make the data available to its client, then return. For additional reasons to be cautious about programming in the kernel, see [Should You Program in the Kernel?](#) (page 21)

If your driver must reside in the kernel, you should be aware of the following issues:

- Most importantly, the kernel is a single program—there is no memory protection between your driver and the rest of the kernel. A kernel-resident driver that behaves badly can crash or hang the operating system.
- A more subtle issue is that function call stacks within the kernel are limited to 16 KB. Be careful not to declare large local variables in functions. Whenever possible, you should preallocate buffers and reuse them.

Kernel-resident drivers have full access to kernel programming interfaces. However, because of their low level of operation, drivers should use only Mach calls and not BSD calls. Many parts of the BSD kernel code aren't currently safe for multithreading or multiprocessing. In any case, drivers rarely need to perform Mach calls directly, as the I/O Kit provides interfaces to most of the kernel-level functionality needed by a driver.

The I/O Registry and the I/O Catalog

The I/O Registry is a dynamic database that records the network of driver objects participating in hardware connections on an OS X system and tracks the provider-client relationships among those objects. A device driver must be recorded in the I/O Registry to participate in most I/O Kit services.

The I/O Registry is a critical part of the I/O Kit because it supports the dynamic features of the operating system, which allows users to add or remove devices (particularly FireWire or USB devices) to and from a running system and have them immediately available, without the need for a reboot. As hardware is added, the system automatically finds and loads the necessary drivers and updates the I/O Registry to reflect the new device configuration; when hardware is removed, the appropriate drivers are unloaded and the Registry is updated again. The Registry always resides in system memory and is not stored on disk or archived between boots.

The I/O Registry structures its data as an inverted tree. Each object in the tree descends from a parent node and can have one or more child nodes; if it is a “leaf” node, it has no children. Almost every node in the tree represents a driver object: a nub or an actual driver. These objects must inherit from the `IORegistryEntry` class (which is the superclass of `IOService`, the superclass of all driver classes). The central characteristic of `IORegistryEntry` objects is a list of associated properties. These properties reflect the personality used in driver matching (see [Driver Matching](#) (page 30)) and otherwise add information about a driver. The properties captured in the Registry derive from each driver’s information property list, a file in the driver KEXT containing key-value pairs describing the driver’s characteristics, settings, and requirements.

Another dynamic database, the I/O Catalog, works closely with the I/O Registry. The I/O Catalog maintains entries for all available drivers on a system. When a nub discovers a device, it requests a list of all drivers of the device’s family from the I/O Catalog.

You can examine the I/O Registry using the I/O Registry Explorer application and the `ioreg` command-line tool, both included in the OS X Developer package. You can also programmatically explore and manipulate the properties of Registry entries using the member functions of the `IORegistryEntry` class. From applications and other programs in user space, you can search for and access driver information in the I/O Registry using APIs in the I/O Kit framework.

For more about the I/O Registry and the I/O Catalog see the chapter [The I/O Registry](#) (page 39) For further information about the `IORegistryEntry` class see [Dynamic Driver Registration \(IORegistryEntry\)](#) (page 62) in the chapter “[The Base Classes](#) (page 52)”

Driver Matching

A primary function of nubs is to provide matching services, matching drivers to devices. Unlike in Mac OS 8 and 9, drivers are not loaded automatically simply because they are installed. In OS X, a driver must first be matched to an existing device before that driver can be loaded.

Driver matching is an I/O Kit process in which a nub, after discovering a specific hardware device, searches for the driver or drivers most suited to that device. To support driver matching, each device driver defines one or more **personalities** that specify the kinds of devices it can support. This information is stored in XML dictionaries defined in the information property list in the driver’s bundle. The dictionary values specify whether a driver is a candidate for a particular device.

When a nub detects a device, the I/O Kit finds and loads a driver for the nub in three distinct phases, using a subtractive process until a successful candidate is found. The phases of matching are:

1. **Class matching**—eliminates drivers of the wrong device class.

2. **Passive matching**—examines each remaining driver’s personalities for properties specific to the device, eliminating those drivers that don’t match.
3. **Active matching**—the remaining driver candidates probe the device to verify that they can drive it.

When a matching driver is found, its code is loaded and an instance of the principal class listed in the personality is created. At this point the driver’s life cycle begins. See [Driver Object Life Cycle](#) (page 64) in the chapter “[The Base Classes](#) (page 52)” for details.

For a detailed discussion of driver personalities and the matching process, see the chapter [Driver and Device Matching](#) (page 43)

The I/O Kit Class Hierarchy

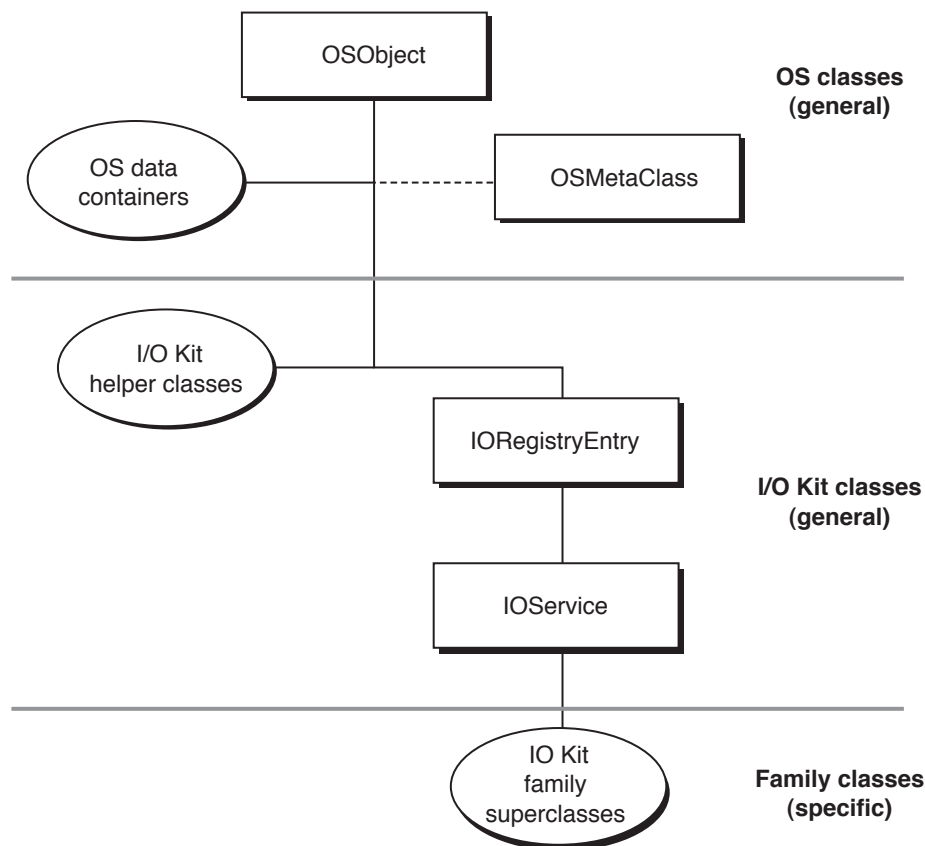
The I/O Kit encompasses dozens of C++ classes and is itself an extension of the libkern C++ library, the foundation for loadable kernel modules. Taken together, the I/O Kit and libkern would seem to form a forbiddingly large and complex hierarchy of classes. Yet the essential structure of that hierarchy is fairly simple, as [Figure 2-3](#) (page 32) illustrates.

You can break down the extended I/O Kit class hierarchy into three broad groupings:

- The classes of libkern (sometimes called the OS classes because of their “OS” prefix)
- The I/O Kit base classes and helper classes
- The classes of the I/O Kit families

For information on where the binaries and header files of libkern and the I/O Kit library are installed, see [Frameworks and Libraries](#) (page 18). For detailed information on the features and interfaces of the base classes in this hierarchy—OSObject, OSMetaClass, IORegistryEntry, and IOService—see [The Base Classes](#) (page 52).

Figure 2-3 I/O Kit extended class hierarchy



Note: The appendix [Base and Helper Class Hierarchy](#) (page 132) contains a hierarchy chart for the I/O Kit base and helper classes; the appendix [I/O Kit Family Reference](#) (page 142) includes class hierarchy charts for most I/O Kit families.

The OS Classes

The I/O Kit is built on top of the libkern C++ library; that is to say, the root superclass for the I/O Kit–specific classes is IORegistryEntry, which inherits from libkern’s OSObject. As is the I/O Kit, libkern is written in a subset of C++ suitable for use in loadable kernel modules. Specifically, the libkern C++ environment excludes the C++ exception-handling and the Runtime Type Information (RTTI) facilities. Instead the OS base classes implement a suitable equivalent of the RTTI feature, among other things.

At the root of the extended hierarchy is the `OSObject` class and closely related to this class is the `OSMetaClass` class. All other OS classes are “helper” classes for such things as collections and other data containers. The following summarizes the roles these classes play:

- `OSMetaClass` implements a runtime type information (RTTI) mechanism, enables some degree of object introspection, and supports runtime allocation of objects derived from `OSObject` by class name.
- `OSObject` features APIs for reference counting (`retain` and `release`), memory management of retained objects, and the automatic disposal of objects when they are no longer needed. `OSObject` also provides a dynamic default implementation of the `init` and `free` methods.
- The OS data containers are subclasses of `OSObject` whose instances encapsulate various types of data values (such as booleans, numbers, strings) and implement and iterate over collections such as arrays and dictionaries.

The OS data containers approximately match their user-space counterparts, the Core Foundation containers, in both name and behavior. Because the characteristics of the OS and Core Foundation classes are so similar, the system can easily convert a Core Foundation type to an OS type and vice versa. For example, a `CFArray` object is transformed into an `OSArray` when crossing the user-kernel boundary.

The OS classes are generally useful for all code written for the kernel, not just device drivers. For example, kernel extensions implementing networking services or file systems can also take advantage of these classes. `OSObject` in particular is an essential common superclass for kernel code. For one kernel module (KMOD) to reference objects created by another KMOD, the objects must ultimately derive from `OSObject`. Most I/O Kit classes assume that the objects being passed around are derived from `OSObject`.

If you are porting existing C++ code to the I/O Kit, you are not required to use the OS classes. But if you decide to forgo the features that these classes provide, such as reference counting or data containers, you'll probably need to implement them yourself.

For more on the `OSObject` and `OSMetaClass` classes, see [The libkern Base Classes](#) (page 53) in the chapter “[The Base Classes](#) (page 52)”

The General I/O Kit Classes

The middle group of the extended class hierarchy comprises the IO Kit base classes— `IORegistryEntry` and `IOService`—and a set of helper classes for resource management, data management, and thread and input control. This group of I/O Kit classes is designated “general” because all device-driver classes can potentially make use of them.

The root class of the I/O Kit hierarchy is `IORegistryEntry`; by virtue of inheritance from `IORegistryEntry`, an I/O Kit object can be a node in the I/O Registry and have one or more property tables (driver personalities) associated with it. `IORegistryEntry` implements a number of features:

- It manages connection into the Registry through a driver's `attach` and `detach` entry points
- It manages the property tables defining driver personalities using `OSDictionary` objects
- It implements locking in the Registry, allowing updates to the Registry to be made atomically

`IOService` is the sole direct subclass of `IORegistryEntry`. Almost all I/O Kit family superclasses inherit, directly or indirectly, from `IOService`. Most importantly, `IOService` specifies the life cycle of device drivers within a dynamic runtime environment. Through matching pairs of virtual functions—such as `init/free`, `start/stop`, and `open/close`—`IOService` defines how driver objects initialize themselves, attach themselves into the I/O Registry, perform all necessary allocations, and then reverse the effects of these actions in the proper order. To support its management of driver life cycles, `IOService` provides matching services (assisting with probing, for instance) and instantiates drivers based on the existence of a provider. In addition, `IOService` includes member functions that are useful for various purposes, including:

- Notification and messaging
- Power management
- Device memory (mapping and accessing)
- Device interrupts (registering, unregistering, enabling, causing, and so on)

For more on the `IORegistryEntry` and `IOService` classes, see [The I/O Kit Base Classes](#) (page 61) in the chapter “[The Base Classes](#) (page 52)”

Most I/O Kit helper classes have several functions related to the runtime environment of device drivers:

- Implementing work loops and event sources (interrupts, timers and commands) along with associated locks and queues
- Implementing memory cursors and memory descriptors for managing the data involved in I/O transfers

For more on the I/O Kit helper classes, see the chapters [Handling Events](#) (page 77) and [Managing Data](#) (page 95)

The I/O Kit Family Classes

Most drivers are instances of a subclass of a class in an I/O Kit family; the family classes, in turn, tend to be subclasses of `IOService`. Your driver class should be a direct subclass of the most appropriate family class for the driver you're trying to write. For example, if you are writing an Ethernet controller driver, your driver class should be a subclass of `IOEthernetController`, not `IONetworkController` (the superclass of `IOEthernetController`).

The I/O Kit has over a dozen families, each with its own set of classes; these families include the following:

- ADB

- ATA and ATAPI
- Audio
- FireWire
- Graphics
- HID (Human Interface Devices)
- Network
- PC Card
- PCI and AGP
- SBP-2
- SCSI Architectural Model
- SCSI Parallel
- Serial
- Storage
- USB

Apple will add additional families as they are developed. If you require a family for a device and it's currently not supported, you can try writing your own family classes. However, don't assume that a new family is required if one does not currently exist. In many cases, the `IOService` class provides all the services a driver requires; such "family-less" drivers can support many devices that are specific to certain vendors. For more information on families, see the chapter [I/O Kit Families](#) (page 69) and the appendix [I/O Kit Family Reference](#) (page 142)

Controlling Devices From Outside the Kernel

Perhaps one of the more compelling features of OS X is the inviolable separation it enforces between the virtual address spaces of processes. Unless laborious arrangements are made for shared memory, one process cannot directly touch data mapped to another process's address space. This separation enhances the stability and reliability of the system by preventing memory trashers and similar annoyances from bringing processes down.

Even more important is the separation between the address spaces of the kernel and of all other processes, which are sometimes said (from the perspective of the kernel) to inhabit "user space." If an application or other program in user space somehow violates the address space of the kernel, the whole system can come crashing down. To make this separation between kernel and user space even more airtight, programs in user space cannot even directly call kernel APIs. They must make system calls to (indirectly) access kernel APIs.

Sometimes, however, a program in user space needs to control or configure a device, and thus needs access to I/O Kit services in the kernel. For example, a game might need to set monitor depth or sound volume, or a disk-backup program might need to act as the driver for a tape drive. Other examples of applications that must somehow interact with the kernel to drive hardware include those running or interpreting data from scanners, joysticks, and digital cameras.

To answer this requirement, the I/O Kit includes two mechanisms: device interfaces and POSIX device nodes. Through a plug-in architecture and well-defined interfaces, the device-interface mechanism enables a program in user space to communicate with a nub in the kernel that is appropriate to the type of device it wishes to control. Through the nub the program gains access to I/O Kit services and to the device itself. For storage, serial, and networking devices, applications can obtain the information they need from the I/O Kit to access and control these devices using POSIX APIs.

Keep in mind that there are some family services that the I/O Kit refuses to export to user space as device interfaces; these services are available only inside the kernel. An example is the PCI family. For reasons of stability and security, external access to PCI resources is forbidden. The appendix [I/O Kit Family Reference](#) (page 142) identifies the families that export device interfaces.

This section summarizes information in *Accessing Hardware From Applications*. Refer to this document for a complete description of device interfaces and how to use them.

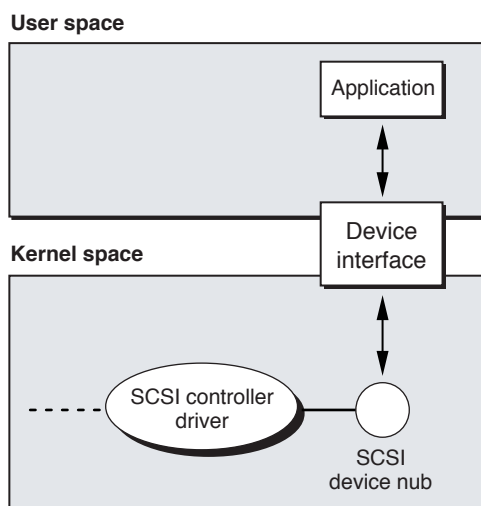
The Device-Interface Mechanism

A device interface is a plug-in interface between the kernel and a process in user space. The interface conforms to the plug-in architecture defined by Core Foundation Plug-in Services (CFPlugIn), which, in turn, is compatible with the basics of Microsoft's Component Object Model (COM). In the CFPlugIn model, the kernel acts as the plug-in host with its own set of well-defined I/O Kit interfaces, and the I/O Kit framework provides a set of plug-ins (device interfaces) for applications to use.

Conceptually, a device interface straddles the boundary between user space and the kernel. It handles negotiation, authentication, and similar tasks as if it were a kernel-resident driver. On the user-space side, it enables communication with the application (or other program) through its exported programmatic interfaces. On the kernel side it enables communication with an appropriate I/O Kit family through a nub created by a driver object of that family. From the kernel's perspective, a device interface appears to be a driver and is known as a "user client." From the application's perspective, the device interface appears as a set of functions that it can call and through which it can pass data to the kernel and receive data back from it. That's because, at an elemental level, a device interface is a pointer to a table of function pointers (although it can also include data fields). Applications, once they obtain an instance of a device interface, can call any of the functions of the interface.

[Figure 2-4](#) (page 37) illustrates the architecture of a device interface, showing an application that has acquired access to a SCSI hard disk through a device interface. It is best to view this diagram as a variation of [Figure 2-2](#) (page 26) which shows the series of driver-object connections made for a kernel-resident SCSI disk driver.

Figure 2-4 An application controlling a SCSI device through a device interface.



At the start, the same series of actions—device discovery, nub creation, matching, driver loading—occurs from the PCI bus driver to the SCSI device nub. But then the SCSI device nub matches and loads the device interface as its driver instead of a kernel-resident driver.

Before an application can use the device-interface mechanism to access a device, it must find the device. It accomplishes this through a process called device matching. In device matching, an application creates a “matching dictionary” that specifies the properties of the target device, then calls an I/O Kit function, passing in the dictionary. The function searches the I/O Registry and returns one or more matching driver objects that the application can then use to load an appropriate device interface. For more on this topic, see [Device Matching](#) (page 51)

If you develop a custom driver that is not a subclass of a class in an I/O Kit family, and you want applications to be able to access the driver, you have to write your own device interface. Any code that communicates between user space and the kernel must use one or more of the following facilities:

- BSD system calls
- Mach IPC
- Mach shared memory

The I/O Kit uses primarily Mach IPC and Mach shared memory. In contrast, the networking and file-system components of OS X use primarily BSD system calls.

POSIX Device Files

BSD, a central component of the OS X kernel environment, exports a number of programmatic interfaces that are consistent with the POSIX standard. These interfaces enable communication with serial, storage, and network devices through device files. In any UNIX-based system such as BSD, a device file is a special file located in `/dev` that represents a block or character device such as a terminal, disk drive, or printer. If you know the name of a device file (for example, `disk0s2` or `mt0`) your application can use POSIX functions such as `open`, `read`, `write`, and `close` to access and control the associated device.

The I/O Kit dynamically creates the device files in `/dev` as it discovers devices. Consequently, the set of device files is constantly changing; different devices might be attached to the device files in `/dev` at any one time, and the same devices might have different device-file names at different times. Because of this, your application cannot hard-code device file names. For a particular device, you must obtain from the I/O Kit the path to its device file through a procedure involving device matching. Once you have the path, you can use POSIX APIs to access the device.

Note that you can access networking services from user space using the BSD socket APIs. However, you should generally use sockets only if the higher-level networking APIs in the Carbon and Cocoa environments do not provide you with the features you require.

The I/O Registry

The I/O Registry is a dynamic database that describes a collection of “live” objects (nubs or drivers) and tracks the provider-client relationships between them. When hardware is added or removed from the system, the Registry is immediately updated to reflect the new configuration of devices. A dynamic part of the I/O Kit, the Registry is not stored on disk or archived between boots. Instead, it is built at each system boot and resides in memory.

The I/O Registry is made accessible from user space by APIs in the I/O Kit framework. These APIs include powerful search mechanisms that allow you to search the Registry for an object with particular characteristics. You can also view the current state of the Registry on your computer using applications provided with the developer version of OS X.

This chapter describes the I/O Registry architecture and the planes the Registry uses to represent relationships between objects. It also provides an overview of device matching and introduces applications that allow you to browse the Registry.

I/O Registry Architecture and Construction

It is most useful to think of the I/O Registry as a tree: Each object is a node that descends from a parent node and has zero or more child nodes. The Registry follows the definition of a tree in nearly all respects, with the exception of a small minority of nodes that have more than one parent. The primary example of this situation is a RAID disk controller where several disks are harnessed together to appear as a single volume. Exceptional cases aside, however, viewing the Registry as a tree will help you visualize how it is constructed and updated.


At boot time, the I/O Kit registers a nub for the Platform Expert, a driver object for a particular motherboard that knows the type of platform the system is running on. This nub serves as the root of the I/O Registry tree. The Platform Expert nub then loads the correct driver for that platform, which becomes the child node of the root. The Platform driver discovers the buses that are on the system and it registers a nub for each one. The tree continues to grow as the I/O Kit matches each nub to its appropriate bus driver, and as each bus driver discovers the devices connected to it and matches drivers to them.

When a device is discovered, the I/O Kit requests a list of all drivers of the device’s class type from another dynamic database, the I/O Catalog. Whereas the I/O Registry maintains the collection of objects active in the currently running system, the I/O Catalog maintains the collection of available drivers. This is the first step in a three-step process known as driver matching that is described in [Driver and Device Matching](#) (page 43)

Information such as class type is kept in the driver's information property list, a file containing XML-structured property information. The property list describes a driver's contents, settings, and requirements in the form of a dictionary of key-value pairs. When read into the system, this information is converted into OS containers such as dictionaries, arrays, and other types. The I/O Kit uses this list in driver matching; a user application can search the I/O Registry for objects with specific properties in a process known as device matching. You can also view the property lists of your computer's currently loaded drivers using I/O Registry Explorer, an application that displays the Registry.

Keeping the tree-like structure of the I/O Registry in mind, now visualize each node extending into the third dimension like a column. The two-dimensional Registry tree, with the Platform Expert nub at its root, is now visible on a plane that cuts perpendicularly through these columns. The I/O Kit defines a number of such planes (you can think of them as a set of parallel planes cutting through the columns at different levels). See [Figure 3-1](#) (page 40) for an illustration of this structure.

Figure 3-1 Two planes in the I/O Registry



There are six planes defined in the I/O Registry:

- Service
- Audio
- Power
- Device Tree
- FireWire
- USB

Each plane expresses a different provider-client relationship between objects in the I/O Registry by showing only those connections that exist in that relationship. The most general is the Service plane which displays the objects in the same hierarchy in which they are attached during Registry construction. Every object in the Registry is a client of the services provided by its parent, so every object's connection to its ancestor in the Registry tree is visible on the Service plane.

The other planes show more specific relationships:

- The Audio plane provides a representation of the audio signal chain that Core Audio framework and its plug-ins use to discover information about the audio signal paths between the system's audio devices.
- The Power plane shows the power interdependencies between I/O Registry objects, allowing you to trace the flow of power from provider to client and discover which objects might be affected if a particular device is powered down.
- The Device Tree plane represents the Open Firmware device hierarchy.
- The FireWire and USB planes each represent the internal hierarchies defined by those standards.

It is important to remember the following points about planes in the I/O Registry:

- All I/O Registry objects exist on all planes, but on any individual plane, only those objects connected by the relationship defined by that plane are visible.
- A driver does not get attached to the Registry on any one particular plane. Instead it may participate in a plane's connections if its provider-client relationships with other objects fit that plane's definition.

The I/O Registry Explorer

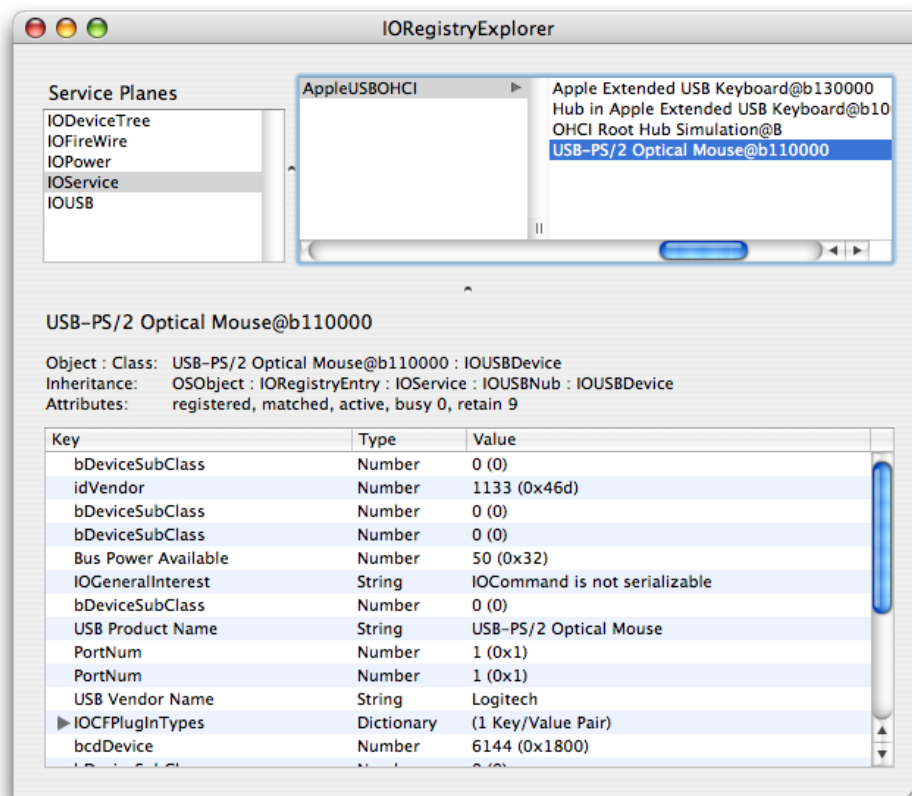
The developer version of OS X provides an application called the I/O Registry Explorer that you can use to examine the configuration of devices on your computer. I/O Registry Explorer provides a graphical representation of the I/O Registry tree. By default, it displays the Service plane, but you can choose to examine any plane. The command-line equivalent, `ioreg`, displays the tree in a Terminal window. This tool has the advantage of allowing you to cut and paste sections of the tree if, for example, you want to send that information in an email message. You can get a complete description of the usage of `ioreg` by typing `man ioreg` at the shell prompt in the Terminal application.

When you open I/O Registry Explorer, a divided window appears with I/O Registry objects in the upper right, the six planes in the upper left, and the property list of the selected object in the lower half of the window. An object followed by a disclosure triangle indicates that it is a parent node. You can traverse the I/O Registry tree by clicking a parent node and dragging the scroller to the right to display its children. [Figure 3-2](#) (page 42) shows an example of a property list in the I/O Registry Explorer window.

Commands in the Tools menu help you search the I/O Registry and examine its contents:

- **Dump Registry Dictionary to Output** places the I/O Registry contents into the console log (viewable through the Console application in `/Applications/Utilities`) if the I/O Registry Explorer was opened from the Finder.
- **Inspector** displays the property list of the currently selected object in ASCII form. Selecting a particular property in the main window causes its value to be displayed in the Inspector window.
- **Force Registry Update** updates I/O Registry Explorer's picture of the I/O Registry to reflect any changes that may have occurred since you first opened the application.
- **Find** performs a case insensitive search on your input string and, if successful, displays the path to the occurrence of the string with object names separated by colons.

Figure 3-2 A sample I/O Registry Explorer window



Driver and Device Matching

Before a device—or any service provider—can be used, a driver for it must be found and loaded into the kernel. The I/O Kit defines a flexible, three-phase matching process that narrows a pool of candidate drivers down to one or more drivers. The final candidate (or, if multiple candidates, the most eligible one) is then loaded and given the first opportunity to manage the device or service provider.

The matching process makes use of the matching dictionaries defined as XML key-value pairs in a driver's information property list. Each matching dictionary specifies a personality of the driver, which declares its suitability for a device or service of a particular type.

This chapter discusses driver personalities and the matching language that describes them. It then describes the matching process, which uses the information in the driver personalities to identify the most appropriate driver for a detected device. The chapter also briefly discusses the device-matching procedure that applications use for loading device interfaces. See *Accessing Hardware From Applications* for the complete details.

Driver Personalities and Matching Languages

Each device driver, considered as a loadable kernel extension (KEXT), must define one or more personalities that specify the kinds of devices it can support. This information is stored in XML matching dictionaries defined in the information property list (`Info.plist`) in the driver's KEXT bundle. A dictionary in this sense is a collection of key-value pairs where the XML tags `<key>` and `</key>` enclose the key. Immediately following the key are the tags enclosing the value; these tags indicate the data type of the value; for example,

```
<integer>74562</integer>
```

would define an integer value.

Each matching dictionary is itself contained within the information property list's `IOPKitPersonalities` dictionary.

The dictionary values of a personality specify whether a driver is a candidate for a particular device. All values in the personality must match for the driver to be selected for the device; in other words, a logical AND is performed on the values. Some of the keys may take a list of space-delimited values, which are generally examined in an OR fashion. Thus you might have a “model” key for a certain PCI card driver personality that takes a list of model numbers, each identifying a supported model from a specific card vendor.

The specific keys that are required depend on the family. A driver for a PCI card, for example, can define a value that is checked against the PCI vendor and device ID registers. Some families, such as the PCI family, provide fairly elaborate matching strategies. For instance, consider this key-value pair:

```
<key>IOPCIMatch</key>
<string>0x00789004&0x00ffffff 0x78009004&0xff00ffff</string>
```

This expression, which is used to match various Adaptec SCSI cards, consists of two compound values, each of which can be a valid match. To evaluate these values, the driver family reads the 32-bit vendor and device ID from the PCI card and masks it with the value to the right of each ampersand. The result of that operation is then compared with the value to the left of the ampersand to determine if there is a match.

[Listing 4-1](#) (page 44) shows a partial listing of a driver personality from the XML file for an Ethernet controller driver.

Listing 4-1 A partial listing of an XML personality for an Ethernet controller

```
<key>IOKitPersonalities</key>
  <dict>
    <dict>
      <!-- Each personality has a different name. -->
      <key>Name</key>      <string>PCI Matching</string>

      <!-- ... some keys not shown ... -->

      <!-- The name of the class IOKit will instantiate when probing. -->
      <key>IOClass</key>   <string>ExampleIntel82558</string>

      <!-- IOKit matching properties
      -- All drivers must include the IOProviderClass key, giving
      -- the name of the nub class that they attach to. The provider
      -- class then determines the remaining match keys. A personality
      -- matches if all match keys do; it is possible for a driver
      -- with multiple personalities to be instantiated more than once
      -- if several personalities match.
      -->
      <key>IOProviderClass</key>
```

```
        <string>IOPCIDevice</string>

        <!-- IOPCIDevice matching uses any of four possible PCI match
        -- criteria. This personality just uses IOPCIMatch to check the
        -- device/vendor ID.
        -->
        <key>IOPCIMatch</key>
            <string>0x12298086</string>

        <!-- The initial match score for this personality.-->
        <key>IOProbeScore</key>    <integer>400</integer>
    </dict>

    <dict>
        <!-- Can have additional personalities. -->
        <!-- ... (not shown) -->
    </dict>
</dict>
```

As mentioned in Listing 4-1 every driver must include the `IOProviderClass` key with a value that identifies the nub to which the driver attaches. In very rare cases, a driver might declare `IOResources` as the value of its `IOProviderClass` key. `IOResources` is a special nub attached to the root of the I/O Registry that makes resources, such as the BSD kernel, available throughout the system. Traditionally, drivers of virtual devices match on `IOResources` because virtual devices do not publish nubs of their own. Another example of such a driver is the HelloIOKit KEXT (described in Creating a Device Driver with Xcode) which matches on `IOResources` because it does not control any hardware.

Important: Any driver that declares `IOResources` as the value of its `IOProviderClass` key must also include in its personality the `IOMatchCategory` key and a private match category value. This prevents the driver from matching exclusively on the `IOResources` nub and thereby preventing other drivers from matching on it. It also prevents the driver from having to compete with all other drivers that need to match on `IOResources`. The value of the `IOMatchCategory` property should be identical to the value of the driver's `IOClass` property, which is the driver's class name in reverse-DNS notation with underbars instead of dots, such as `com_MyCompany_driver_MyDriver`.

Because a driver can contain multiple matching dictionaries, each one defining a different personality for the driver, the same driver code can be loaded for different devices. For purposes of competition, the I/O Kit treats each personality as if it were a driver. If, in any single personality, all of the properties required by the family match, the driver's code is loaded and given a chance to run for that device.

Your driver can have more than one personality for a variety of reasons. It could be that the driver (as packaged in the KEXT) supports more than one type of device, or more commonly, multiple versions of the same type of device. Another reason might be that the driver supports similar devices, each of which is attached to the system on different buses; for example, Zip drives can be attached to USB, FireWire, SCSI, ATAPI, and other buses. Because each of these attaches to a different nub class, it has different matching values. The personalities of a driver can also range from device-generic to device-specific. The personalities of the `AppleUSBAudio` driver ([Listing 4-2](#) (page 46)) illustrate this.

Listing 4-2 Driver personalities for the `AppleUSBAudio` driver

```
<key>IOKitPersonalities</key>
<dict>
  <key>AppleUSBAudioControl</key>
  <dict>
    <key>CFBundleIdentifier</key>
    <string>com.apple.driver.AppleUSBAudio</string>
    <key>IOClass</key>
    <string>AppleUSBAudioDevice</string>
    <key>IOProviderClass</key>
    <string>IOUSBInterface</string>
    <key>bInterfaceClass</key>
    <integer>1</integer>
    <key>bInterfaceSubClass</key>
    <integer>1</integer>
  </dict>
</dict>
```

```
<key>AppleUSBAudioStream</key>
<dict>
  <key>CFBundleIdentifier</key>
  <string>com.apple.driver.AppleUSBAudio</string>
  <key>IOClass</key>
  <string>AppleUSBAudioDMAEngine</string>
  <key>IOProviderClass</key>
  <string>IOUSBInterface</string>
  <key>bInterfaceClass</key>
  <integer>1</integer>
  <key>bInterfaceSubClass</key>
  <integer>2</integer>
</dict>
<key>AppleUSBTrinityAudioControl</key>
<dict>
  <key>CFBundleIdentifier</key>
  <string>com.apple.driver.AppleUSBAudio</string>
  <key>IOClass</key>
  <string>AppleUSBTrinityAudioDevice</string>
  <key>IOProviderClass</key>
  <string>IOUSBInterface</string>
  <key>bConfigurationValue</key>
  <integer>1</integer>
  <key>bInterfaceNumber</key>
  <integer>0</integer>
  <key>idProduct</key>
  <integer>4353</integer>
  <key>idVendor</key>
  <integer>1452</integer>
</dict>
</dict>
```

This matching dictionary defines three personalities: `AppleUSBAudioControl`, `AppleUSBAudioStream`, and `AppleUSBTrinityAudioControl`. In matching for a detected USB Trinity audio-control device, the `AppleUSBTrinityAudioControl` would be chosen; for any other audio-control device, the generic personality (`AppleUSBAudioControl`) would match.

One common property of personalities is the probe score. A probe score is an integer that reflects how well-suited a driver is to drive a particular device. A driver may have an initial probe-score value in its personality and it may implement a `probe` function that allows it to modify this default value, based on its suitability to drive a device. As with other matching values, probe scores are specific to each family. That's because once matching proceeds past the class-matching stage, only personalities from the same family compete. For more information on probe scores and what a driver does in the `probe` function, see [Device Probing](#) (page 49)

Driver Matching and Loading

At boot time and at any time devices are added or removed, the process of driver matching occurs for each detected device (or other service provider). The process dynamically locates the most suitable driver in `/System/Library/Extensions` for the device or service.

As described in [Driver Matching](#) (page 30) in the chapter [Architectural Overview](#) (page 23) the matching process is triggered when a bus controller driver scans its bus and detects a new device attached to it. For each detected device the controller driver creates a nub. The I/O Kit then initiates the matching process and obtains the values from the device to use in matching (for example, examining the PCI registers). Once a suitable driver is found for the nub, the driver is registered and loaded. That driver, in turn, may create its own nub (possibly through behavior inherited from its family), which initiates the matching process to find a suitable driver.

Driver Matching

When a nub detects a device, the I/O Kit finds and loads a driver for the nub in three distinct phases, using a subtractive process. In each phase, drivers that are *not* considered to be likely candidates for a match are subtracted from the total pool of possible candidates until a successful candidate is found.

The matching process proceeds as follows:

1. In the **class matching** step, the I/O Kit narrows the list of potential drivers by eliminating any drivers of the wrong class for the *provider* service (that is, the nub). For example, all driver objects that descend from a SCSI class can be ruled out when the search is for a USB driver.
2. In the **passive matching** step, the driver's personality (specified in a driver's XML information property list) is examined for properties specific to the provider's family. For example, the personality might specify a particular vendor name.

3. In the **active matching** step, the driver's probe function is called with reference to the nub it is being matched against. This function allows the driver to communicate with the device and verify that it can in fact drive it. The driver returns a probe score that reflects its ability to drive the device. See [Device Probing](#) (page 49) for more information. During active matching, the I/O Kit loads and probes all candidate drivers, then sorts them in order of highest to lowest probe score.

The I/O Kit then chooses the remaining driver with the highest probe score and starts it. If the driver successfully starts, it is added to the I/O Registry and any remaining driver candidates are discarded. If it does not start successfully, the driver with the next highest probe score is started, and so on. If more than one driver is in the pool of possible candidates, the more generic driver typically loses out to the more specific driver if both claim to be able to drive the device.

Device Probing

During the active matching phase, the I/O Kit requests each driver in the pool of remaining candidates to probe the device to determine if they can drive it. The I/O Kit calls a series of member functions defined in the `IOService` class and overridden in some cases by the driver's class. These functions, and the order in which they are called, are

```
init()
attach()
probe()
detach()
free() /* if probe fails */
```

These functions comprise the first part of a driver's life cycle (see [Driver Object Life Cycle](#) (page 64) in the chapter [The Base Classes](#) (page 52) for the full story). Note that four of these functions form complementary pairs, one nested inside the other: `init` and `free` are one pair, and `attach` and `detach` are the other.

During active matching, the code of a candidate driver is loaded and an instance of the principal class listed in the personality is created. The first function invoked is `init`, which is the `libkern` equivalent of the constructor function for the class. For I/O Kit drivers, this function takes as its sole argument an `OSDictionary` object containing the matching properties from the selected driver personality. The driver can use this to identify what specific personality it's been loaded for, determine what level of diagnostic output to produce, or otherwise establish basic operating parameters. I/O Kit drivers typically don't override the `init` function, performing their initialization in later stages.

However, if you do override the `init` function—or almost any other function of the driver life cycle—you must take care to do two things. The first is to invoke your superclass's implementation of the function. When you do this depends on the function; for example, in implementing `init` you should invoke the superclass's

implementation as the first thing, and in `free` you should invoke it as the last statement of the function. The second general rule is that you should undo in the second function of a pair what you've done in the first function; thus, if you allocate memory for any reason in `init`, you should free that memory in `free`.

Next, the `attach` function (which is bracketed with the `detach` function) is called. The default implementation of `attach` attaches the driver to the nub through registration in the I/O Registry; the default implementation of `detach` detaches the driver from the nub. A driver can override the default implementations, but rarely needs to do so.

After `attach` the `probe` function is invoked. The I/O Kit always calls a driver's `probe` function if the driver's matching dictionary passively matches the provider (the nub). A driver may choose not to implement `probe`, in which case `IOService`'s default implementation is invoked, which simply returns `this`.

The `probe` function takes as arguments the driver's provider and a pointer to a probe score. The probe score is a signed 32-bit integer initialized to a value specified in the driver's personality (or to zero if not explicitly initialized). The driver with the highest initial probe score is given the first chance to start operating the device. The purpose of the `probe` function is to offer drivers an opportunity to check the hardware and to modify their default probe scores as assigned in their personalities. A driver can check device-specific registers or attempt certain operations, adjusting its probe score up or down based on how well suited it is for the device it is examining. Whatever it finds, each driver must leave the hardware in the same state it was in when `probe` was invoked so the next driver can probe the hardware in its original state.

A driver, in its `probe` function, returns a driver object (`IOService *`) if the probe was successful and returns zero otherwise. The returned object is usually the driver itself, but the driver can return another driver that is more suited to the provider. The probe score is an in-out parameter, which `probe` can modify based on what it discovers about the device.

Driver Loading

After all drivers have probed the device, the one with the highest probe score is attached and its `start` function, which must be implemented by all drivers, is invoked. The `start` function initializes the device hardware and prepares it for operation. If the driver succeeds in starting, it returns `true`; the remaining candidate driver instances are discarded and the driver that started successfully continues operating. If the driver cannot initialize the hardware it must leave the hardware in the state it was in when `start` was invoked and return `false`. The failing driver is then detached and discarded, and the candidate driver with the next highest probe score is given a chance to start.

Some time after this occurs, all loaded drivers that are not currently in use are unloaded.

Device Matching

A user application that requires access to a device must first search for that device and then acquire the appropriate device interface to communicate with it. This process is known as device matching. Unlike driver matching, device matching searches the I/O Registry for a driver that is already loaded.

To perform device matching, follow these basic steps:

1. Establish a connection with the I/O Kit by obtaining a Mach port.
2. Define a dictionary that specifies the type of device to search for in the I/O Registry. The search can be refined by setting additional values in the dictionary. For example, a search for IOMedia objects can be narrowed down to find all ejectable media. You can find the values to match in the device header files (such as `IOSCSIDevice.h` or `IOATADevice.h`), by referring to the family-specific documentation, or by looking at the information property lists displayed in output from the I/O Registry Explorer application.
3. Obtain a list of all objects in the Registry that match your dictionary and choose the appropriate device.
4. Access the device you have chosen by obtaining a device interface for it. This step is explained more fully in [Controlling Devices From Outside the Kernel](#) (page 35)

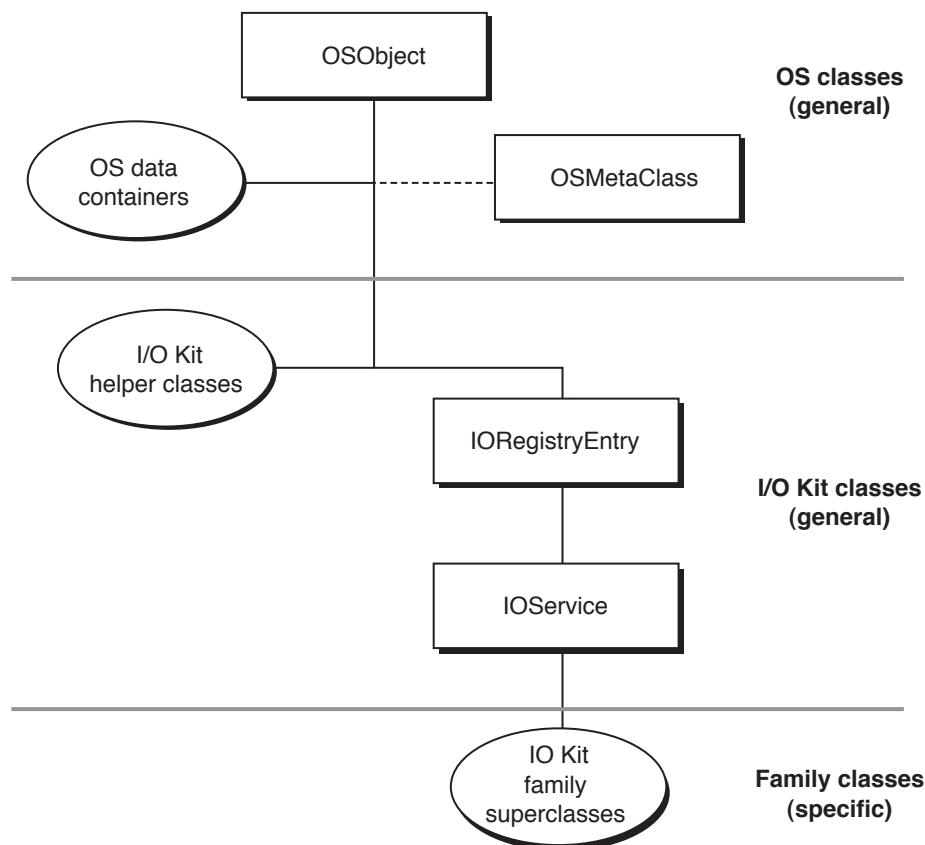
See the document *Accessing Hardware From Applications* for a full description of device matching.

The Base Classes

The I/O Kit is an object-oriented framework consisting primarily of dozens, if not hundreds, of C++ classes. These classes can be organized by virtue of their inheritance relationships in a class hierarchy. As with all class hierarchies, the I/O Kit's can be depicted as an inverted tree, with childless nodes—classes without any subclasses—as the leaves of the tree. Carrying the analogy further, the classes at the trunk and, especially, the root of the tree are those that most classes of the hierarchy inherit from. These are the base classes.

Figure 5-1 (page 52) shows the general outline of the I/O Kit's class hierarchy and the positions of the base classes within this hierarchy.

Figure 5-1 The base classes of the I/O Kit class hierarchy



As the diagram illustrates, the base classes specific to the I/O Kit are **IOService** and **IORegistryEntry**; also included as base classes—through inheritance—are the libkern library's **OSObject** and (in a special sense) **OSMetaClass**.

Given the centrality of these classes, it is apparent how important it is to understand them. They provide not just the behavior and data structures that all other classes of the I/O Kit inherit. They define the *structure* of behavior for kernel and driver objects: how objects are created and disposed of, how metaclass information is captured and revealed, how driver objects should behave within a dynamic runtime environment, and how the client/provider relationships among driver objects are dynamically established. If you're writing device drivers using the I/O Kit, you're going to have to deal with the base classes in your code early on and frequently thereafter, so it's a good idea to become familiar with them.

This chapter also gives an overview of some generally useful functions and data types. Even though these functions and types do not properly belong in a discussion of base classes (since they are not affiliated with any class), their utility in a variety of circumstances makes them almost as central as any of the base classes.

The libkern Base Classes

The I/O Kit is built on top of the libkern C++ library, which is written in a subset of C++ suitable for use in loadable kernel modules. Specifically, the libkern C++ environment does not support multiple inheritance, templates, the C++ exception-handling facility, and runtime type information (RTTI). The C++ RTTI facility is omitted because it doesn't support dynamic allocation of classes by name, a feature required for loading kernel extensions. RTTI also makes considerable use of exceptions. However, the libkern C++ environment defines its own runtime typing system, which does support dynamic loading.

Exceptions are forbidden in the kernel for reasons of both cost and stability. They increase the size of the code, thereby consuming precious kernel memory, and introduce unpredictable latencies. Further, because I/O Kit code may be invoked by many client threads, there's no way to guarantee that an exception will be caught. Using `try`, `throw`, or `catch` in any kernel extension is not supported and will result in a compilation error. Although you can't use exceptions in an I/O Kit driver, your driver should always check return codes where appropriate.

Apple highly recommends that you base all kernel C++ code, including that for device drivers, on the libkern base classes, `OSObject` and `OSMetaClass`, and observe the conventions prescribed by those classes (see [Type Casting, Object Introspection, and Class Information](#) (page 57)). Classes that are completely private to your driver need not be based on `OSObject` and need not follow these conventions. Such classes, however, will be limited in their interaction with libkern classes. For example, all libkern collection classes store objects that inherit from `OSObject`. Custom classes that don't inherit from `OSObject` can't be stored in libkern collections such as `OSDictionary` or `OSArray` objects.

Important: At present, the loader does not allow the use of any OSObject subclass that requires qualification, such as a nested class or a class declared within a namespace (for an example of a namespace declaration, see [Language Choice](#) (page 16)). For example, the following nested class declaration in an I/O Kit driver would prevent the driver from loading:

```
class com.mycompany.driver.myClass {  
    class myNestedClass : public IOService {}; // This is not allowed.  
};
```

Object Creation and Disposal (OSObject)

OSObject is at the root of the extended I/O Kit hierarchy. It inherits from no (public) superclass, and all other libkern and I/O Kit classes (except for OSMetaClass) inherit from it. OSObject implements the dynamic typing and allocation features needed to support loadable kernel modules. Its virtual functions and overridden operators define how objects are created, retained, and disposed of in the kernel. OSObject is an abstract base class, and therefore cannot itself be instantiated or copied.

Object Construction

The standard C++ constructors cannot be used in libkern because these constructors use exceptions to report failures; as you may recall, the restricted form of C++ chosen for libkern excludes exceptions. So the main purpose of the OSObject class (and also of the OSMetaClass class) is to reimplement object construction.

For constructing objects, OSObject defines the `init` function and overrides the `new` operator. The `new` operator allocates memory for an object and sets the object's reference count to one. After it uses the `new` operator, the client must call the `init` function on the new object to perform all initializations required to make it a usable object. If the `init` call fails, then the client must immediately release the object.

In support of OSObject's `init` and `new`, the OSMetaClass class implements macros related to object construction. These macros bind a class into the kernel's runtime typing facility and automatically define functions that act as the constructor and destructor for the class. See [Runtime Type Information \(OSMetaClass\)](#) (page 55) for more information on these macros and OSMetaClass's implementation of RTTI.

Subclasses of OSObject do not explicitly implement their constructors and destructors since these are essentially created through the OSMetaClass macros. Moreover, you typically invoke neither constructor and destructor functions, nor the C++ `new` and `delete` operators. These functions and operators are reserved for use by the dynamic typing and allocation facilities, which implicitly define them for a class. In their place, OSObject defines a convention for creating and initializing objects. Subclasses do, however, typically override the `init` function to perform initializations specific to the class.

Most libkern and I/O Kit classes define one or more static functions for creating instances. The naming convention varies from class to class, but the name is usually either the base name of the class itself (with a lowercase first letter), or some form of `with...` where the name describes the initialization arguments. For example, `OSArray` defines the static creation functions `withCapacity`, `withObjects`, and `withArray`; `IOTimerEventSource` defines `timerEventSource`; and `IOMemoryCursor` defines `withSpecification`. If a class doesn't have static creation functions, you must use `new` and then invoke the initialization method that takes the place of the C++ constructor, as shown in [Listing 5-2](#) (page 60)

For an overview of the boilerplate code you need to specify your class's constructor and destructor functions, see [Type Casting, Object Introspection, and Class Information](#) (page 57)

Object Retention and Disposal

`OSObject` defines a reference-counting and automatic-deallocation mechanism to support the safe unloading of kernel extensions. For this mechanism it uses three virtual member functions—`retain`, `release`, and `free`—and overrides the `delete` operator. Of these, the only functions you should call in your code are `retain` and `release`, and you should follow certain conventions that dictate when to call them.

Newly created objects and copied objects have a reference count of one. If you have created or copied a libkern object and have no need to keep it beyond the current context, you should call `release` on it. This decrements the object's reference count. If that count is zero, the object is deallocated; specifically, the `release` method invokes the alternative destructor, named `free`, and finally invokes the `delete` operator. If you don't own an object—that is, you did not create or copy it—and you want to keep it past the current context, call `retain` on it to increment its reference count. If you did not create, copy, or call `retain` on an object, you should never call `release` on it.

In addition, some functions that return objects pass ownership to the caller, meaning the caller must release the object when it is finished with it, while others don't. See the reference documentation for a given function to find out if your code needs to retain or release an object it receives.

Never invoke the `delete` operator explicitly to free an object. Also, never call `free` directly to free an object; however, you may (and should, in most circumstances) override the `free` function to deallocate memory allocated in your `init` function.

Runtime Type Information (OSMetaClass)

Although libkern's restricted form of C++ excludes the native runtime type information (RTTI) facility, `OSMetaClass` implements an alternative runtime typing facility that does support dynamic allocation of classes by name. `OSMetaClass` is not a base class in the true sense; no public libkern or I/O Kit class inherits from it. However, `OSMetaClass` provides APIs and functionality that are essential for object construction and destruction. `OSMetaClass` itself is an abstract class and cannot be directly constructed.

The functionality that `OSMetaClass` offers all libkern-based code includes the following:

- A mechanism for tracking the class hierarchy dynamically
- Safe loading and unloading of kernel modules

The runtime typing facility enables the system to track how many instances of each libkern (and I/O Kit) class are currently extant and to assign each of these instances to a kernel module (KMOD).

- Automatic construction and deconstruction of class instances
- Macros and functions for dynamic type casting, type discovery, membership evaluation, and similar introspective behavior
- Dynamic allocation of libkern class instances based on some indication of their class type, including C-string names

In libkern's runtime typing facility, one static metaclass instance (derivative of `OSMetaClass`) is created for every class in a kernel module (KMOD) loaded into the kernel. The instance encapsulates information on the class's name, size, superclass, kernel module, and the current count of instances of that class. The process of loading a kernel module takes place in two phases, the first initiated by the `preModLoad` member function and the second by the `postModLoad` function. During the `preModLoad` phase, `OSMetaClass` statically constructs, within the context of a single, lock-protected thread, a metaclass instance for each class in the module. In the `postModLoad` phase, `OSMetaClass` links together the inheritance hierarchy of constructed metaclass objects, inserts the metaclass instances into the global register of classes, and records for each instance the kernel module it derived from. See the `OSMetaClass` reference documentation for more on `preModLoad`, `postModLoad`, and related functions.

The created store of metaclass information forms the basis for the capabilities of `OSMetaClass` listed above. The following sections explore the more important of these capabilities in some detail.

Object Construction and Dynamic Allocation

One of the features of `OSMetaClass` is its ability to allocate libkern objects based upon some indication of class type. Subclasses of `OSMetaClass` can do this dynamically by implementing the `alloc` function; the class type is supplied by the `OSMetaClass` subclass itself. You can also allocate an instance of any libkern class by calling one of the `allocClassWithName` functions, supplying an appropriate identification of class type (`OSSymbol`, `OSString`, or C string).

Freshly allocated objects have a retain count of 1 as their sole instance variable and are otherwise uninitialized. After allocation, the client should immediately invoke the object's initialization function (which is `init` or some variant of `init`).

OSMetaClass defines a number of runtime type-declaration macros and object-construction macros based on the `alloc` function. Based on the type of class (virtual or otherwise), you must insert one of these macros as the first statement in class declarations and implementations:

`OSDeclareDefaultStructors`

Declares the data and interfaces of a class, which are needed as runtime type information. By convention this macro should immediately follow the opening brace in a class declaration.

`OSDeclareAbstractStructors`

Declares the data and interfaces of a virtual class, which are needed as runtime type information. By convention this macro should immediately follow the opening brace in a class declaration. Use this macro when the class has one or more pure virtual methods.

`OSDefineMetaClassAndStructors`

Defines an OSMetaClass subclass and the primary constructors and destructors for a non-abstract subclass of OSObject. This macro should appear at the top of the implementation file just before the first function is implemented for a particular class.

`OSDefineMetaClassAndAbstractStructors`

Defines an OSMetaClass subclass and the primary constructors and destructors for a subclass of OSObject that is an abstract class. This macro should appear at the top of the implementation file just before the first function is implemented for a particular class.

`OSDefineMetaClassAndStructorsWithInit`

Defines an OSMetaClass subclass and the primary constructors and destructors for a non-abstract subclass of OSObject. This macro should appear at the top of the implementation file just before the first function is implemented for a particular class. The specified initialization routine is called once the OSMetaClass instance has been constructed at load time.

`OSDefineMetaClassAndAbstractStructorsWithInit`

Defines an OSMetaClass subclass and the primary constructors and destructors for a subclass of OSObject that is an abstract class. This macro should appear at the top of the implementation file just before the first function is implemented for a particular class. The specified initialization routine is called once the OSMetaClass instance has been constructed at load time.

See [Type Casting, Object Introspection, and Class Information](#) (page 57) for more information on using these macros, including examples of usage.

Type Casting, Object Introspection, and Class Information

OSMetaClass defines many macros and functions that you can use in almost any situation. They help you safely cast from one type to another, discover an arbitrary object's class, determine if an object inherits from a given superclass, find out how many instances of a given class are still allocated, and yield other useful information. [Table 5-1](#) (page 58) summarizes these macros and functions.

Table 5-1 OSMetaClass type-casting and introspection APIs

Function or macro	Description
OSTypeID	This macro returns the type ID of a class based on its name.
OSTypeIDInst	This macro returns the type ID of the class a given instance is constructed from.
OSCheckTypeInst	This macro checks if one instance is of the same class type as another instance.
OSDynamicCast	This macro dynamically casts the class type of an instance to a suitable class. It is basically equivalent to RTTI's <code>dynamic_cast</code> .
isEqualTo	This function verifies if the invoking OSMetaClass instance (which represents a class) is the same as another OSMetaClass instance. The default implementation performs a shallow pointer comparison.
metaCast (multiple)	This set of functions determines if an OSMetaClass instance (which represents a class) is, or inherits from, a given class type. The type can be specified as OSMetaClass, OSSymbol, OSString, or a C string.
modHasInstance	Returns whether a kernel module has any outstanding instances. This function is usually called to determine if a module can be unloaded.
getInstanceCount	This function returns the number of instances of the class represented by the receiver.
getSuperClass	This function returns the receiver's superclass.
getClassName	This function returns the name (as a C string) of the receiver.
getClassSize	This function returns the size (in bytes) of the class represented by the receiver.

Defining C++ Classes in libkern

When implementing a C++ class based on `OSObject`, you invoke a pair of macros based upon the `OSMetaClass` class. These macros tie your class into the libkern runtime typing facility by defining a metaclass and by defining the constructor and destructor for your class that perform RTTI bookkeeping tasks through the metaclass.

The first macro, `OSDeclareDefaultStructors` declares the C++ constructors; by convention you insert this macro as the first element of the class declaration in the header file. For example:

```
class MyDriver : public IOEthernetController
{
```

```
OSDeclareDefaultStructors(MyDriver);  
/* ... */  
};
```

Your class implementation then uses the companion macro, `OSDefineMetaClassAndStructors`, to define the constructor and destructor, as well as the metaclass that provides the runtime typing information. `OSDefineMetaClassAndStructors` takes as arguments the name of your driver and the name of its superclass. It uses these to generate code that allows your driver class to be loaded and instantiated while the kernel is running. For example, `MyDriver.cpp` might begin like this:

```
#include "MyDriver.h"  
  
// This convention makes it easy to invoke superclass methods.  
#define super    IOEthernetController  
  
// You cannot use the "super" macro here, however, with the  
// OSDefineMetaClassAndStructors macro.  
OSDefineMetaClassAndStructors(MyDriver, IOEthernetController);
```

The definition of the `super` macro allows convenient access to superclass methods without having to type the whole name of the superclass every time. This is a common idiom of libkern and I/O Kit class implementations.

In place of the C++ constructor and destructor, your class implements an initialization method and a `free` method. For non-I/O Kit classes, the initialization method takes whatever arguments are needed, can have any name (although it usually begins with `init`), and returns a C++ `bool` value. The `free` method always takes no arguments and returns `void`.

The initialization method for your driver class should invoke the appropriate superclass initialization method before doing anything else, as shown in [Listing 5-1](#) (page 59). If the superclass returns `false`, your class's initialization method should abort, release any allocated resources, and return `false`. Otherwise your class can perform its initialization and return `true`. When the libkern C++ runtime system creates an instance of a class, it zero-fills all of the member variables, so you don't need to explicitly initialize anything to zero, false, or null values.

Listing 5-1 Implementing an `init` method

```
bool MyDriver::init(IOPhysicalAddress * paddr)  
{
```

```
    if (!super::init()) {  
        // Perform any required clean-up, then return.  
        return false;  
    }  
    physAddress = paddr; // Set an instance variable.  
    return true;  
}
```

To create an instance using the initialization method, you write code such as this:

Listing 5-2 Creating an instance and calling its init method

```
MyDriver * pDrv = new MyDriver; // This invokes the predefined constructor  
                                // of MyDriver itself  
  
if (!pDrv) {  
    // Deal with error.  
}  
  
if (!pDrv->init(memAddress)) {  
    // Deal with error.  
    pDrv->release(); // Dispose of the driver object.  
}
```

Because this makes creating instances more cumbersome, you may want to write a convenience method in the manner of many of the kernel C++ classes, as for example:

```
MyDriver * MyDriver::withAddress(IOPhysicalAddress *paddr)  
{  
    MyDriver * pDrv = new MyDriver;  
  
    if (pDrv && !pDrv->init(paddr)) {  
        pDrv->release();  
        return 0;  
    }  
}
```

```
    }  
    return pDrv;  
}
```

Using this convenience method, you can create an instance of your driver with code like the following:

```
MyDriver * pDrv = MyDriver::withAddress(paddr);  
  
if (!pDrv) {  
    // Deal with error of not being able to create driver object.  
}  
else {  
    // Go on after successful creation of driver object.  
}
```

A class's `free` method should release any resources held by the instance and then invoke the superclass's `free` method, as in [Listing 5-3](#) (page 61):

Listing 5-3 Implementing the free function

```
void MyDriver::free(void)  
{  
    deviceRegisterMap->release();  
    super::free();  
    return;  
}
```

Again, note that your code should never invoke `free` or the `delete` operator directly with objects based on the `OSObject` class. Always call `release` on such objects to dispose of them.

The I/O Kit Base Classes

All driver objects based on the I/O Kit inherit from the two base classes `IORegistryEntry` and `IOService`. The second of these classes, `IOService`, directly inherits from `IORegistryEntry` and all driver objects ultimately inherit from `IOService`. The `IORegistryEntry` class defines a driver object as a node in the I/O Registry, and `IOService` defines the life cycle of a driver object as well as implementing other behavior common to drivers.

The close inheritance relationship between `IORegistryEntry` and `IOService` might invite speculation as to why these classes weren't designed as one class. The reason is performance. Having `IORegistryEntry` as a superclass of `IOService` is an optimization because, in terms of memory footprint, the `IORegistryEntry` object is much more lightweight.

Dynamic Driver Registration (`IORegistryEntry`)

An `IORegistryEntry` object defines a node (or entry) in the I/O Registry. As the chapter [The I/O Registry](#) (page 39) explains in detail, the I/O Registry is a dynamic database that captures the current graph of “live” driver objects, tracking the client/provider relationships among these objects and recording the properties that describe their personalities. The I/O Registry plays an essential role in the dynamic features of OS X; when users add or remove hardware, the system uses the Registry in the driver-matching process and immediately updates it to reflect the new configuration of devices.

Each `IORegistryEntry` object has two dictionaries (that is, instances of `OSDictionary`) associated with it. One is the property table for the object, which is typically also a driver object. This property table is the matching dictionary that specifies one of the driver's personalities. (See [Driver Personalities and Matching Languages](#) (page 43) for information on personalities.) The other dictionary of an `IORegistryEntry` object is the plane dictionary, which specifies how the object is connected to other objects in the registry.

In addition to reflecting all client/provider relationships among driver objects, the I/O Registry identifies subsets of these relationships. Both the totality of the Registry tree and the subsets of it are called **planes**. Each plane expresses a different provider/client relationship between objects in the I/O Registry by showing only those connections that exist in that relationship. Often the plane relationship is one of a dependency chain. The most general plane is the Service plane which displays the total hierarchy of registry entries. Every object in the Registry is a client of the services provided by its parent, so every object's connection to its ancestor in the Registry tree is visible on the Service plane. In addition to the Service plane, there are the Power, Audio, Device, FireWire, and USB planes. For more information on planes, see [I/O Registry Architecture and Construction](#) (page 39)

It is possible to have an `IORegistryEntry` object that is not also an `IOService` object. Such an object could be used purely for holding information associated with that node in the Registry. However, there is little actual need for such objects.

The `IORegistryEntry` class includes many member functions that driver objects might find useful; these functions fall into several categories:

- Property-table functions allow you to set, get, and remove properties of an `IORegistryEntry` object's property table as well as serializing property tables. Some `getProperty` functions perform a synchronized, recursive search through the Registry for the property of a given key.

- Positional functions let an `IORegistryEntry` object manipulate its position in the Registry tree. It can locate, identify, and attach to or detach from another `IORegistryEntry` object.
- Iteration functions enable your code to traverse the entire Registry tree, or a portion of it, and optionally invoke an “applier” callback function on `IORegistryEntry` objects encountered.

See the reference documentation for `IORegistryEntry` for details.

Basic Driver Behavior (`IOService`)

Every driver object in the I/O Kit is an instance of a class that ultimately inherits from the `IOService` class. `IOService` most importantly defines, through complementary pairs of virtual functions, a driver’s life cycle within a dynamic runtime environment. It manages the matching and probing process, implements default matching behavior, and registers drivers and other services. But the `IOService` class also provides a wealth of functionality for many other purposes, including:

- Accessing a driver’s provider, clients, state, and work loop
- Posting notifications and sending messages to other driver objects or services
- Managing power in devices
- Implementing user clients (device interfaces)
- Accessing device memory
- Registering and controlling interrupt handlers

This section first describes the life cycle of a driver object and `IOService`’s role in that life cycle. Then it summarizes each of the other major `IOService` APIs.

Driver Object Life Cycle

An I/O Kit driver can be loaded and unloaded, or activated and deactivated, at any time. Every driver's life cycle follows the same pattern, as laid out in a set of functions defined by the standard driver superclass, `IOService`. Some of these functions must be implemented by the driver; others are implemented by `IOService`, but can be overridden by the driver for additional flexibility.

Figure 5-2 Driver object life-cycle functions

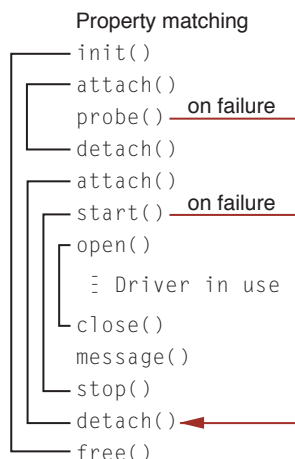


Figure 5-2 (page 64) shows the sequence of functions that gets invoked during the life of a driver object. The bracketing lines show how these functions are grouped into complementary pairs. A driver object class can override any of these functions, but must be sure to invoke the superclass's implementation of that same function at the appropriate point in its own implementation. For example, when you override the opening function of a complementary pair, such as `init` or `start`, your version must invoke the corresponding function of its superclass before doing its own initialization, as shown in Listing 5-1 (page 59). When you override a closing function, such as `free` or `stop`, you should perform your own cleanup before invoking the corresponding function in the superclass, as shown in Listing 5-3 (page 61).

Driver Matching and Loading

The first group of functions—`init`, `attach`, `probe`, and `detach`—is invoked during the process of driver matching and loading. This process occurs at boot time and at any time devices are added or removed. The following paragraphs summarize the matching process, paying special attention to the functions involved; see [Driver Matching and Loading](#) (page 48) for an extended discussion of the process.

The matching process is kicked off when a service provider detects a device. Usually this provider is the controller driver for a bus (such as a PCI bus), which detects the device by scanning its bus. The provider (usually through its family) then creates and registers any required nubs by calling the `IOService` function, `registerService`; this call, in turn, triggers the matching process.

As orchestrated by IOService, the I/O Kit finds and loads a driver for a nub in three distinct phases, using a subtractive process. In each phase, drivers that are *not* considered to be likely candidates for a match are subtracted from the total pool of possible candidates until a successful candidate is found. The phases are:

1. **Class matching**—the I/O Kit eliminates any drivers of the wrong provider (nub) class.
2. **Passive matching**—the I/O Kit examines the remaining drivers' personalities for family-specific properties.
3. **Active matching**—IOService calls each of the remaining drivers' probe functions with reference to the object the driver is being matched against. This function allows the driver to communicate with the device and verify that it can in fact drive that device. A probe score is returned that reflects how well suited the driver is to drive the device.

When a matching driver is found, its code is loaded and an instance of the principal class listed in the personality is created.

In the first two phases, none of the driver's life-cycle functions is called. It is only during the third stage of active matching, when the driver is asked to probe a device for suitability, that the first group of functions is invoked.

Whether a driver is loaded to drive a device or is merely asked to probe it, the first life-cycle function invoked is `init`, which is the libkern equivalent of the constructor function for the class. For I/O Kit drivers, this function takes as its sole argument an `OSDictionary` containing the matching properties from the personality in the XML file. The driver can use this to determine what specific personality it's been loaded for, determine the level of diagnostic output to produce, or otherwise establish basic operating parameters. However, I/O Kit drivers typically don't override the `init` function, performing their initialization in later stages, as described below. For more on `init`, and the related `free` function, see [Object Creation and Disposal \(OSObject\)](#) (page 54)

Before a driver object can either probe or start, it must be attached into the I/O Registry. To do this, the nub invokes the driver's `attach` function, which attaches the driver to the nub through the I/O Registry. The complementary function `detach` removes the driver from its nub. IOService gives both of these functions default implementations. A driver can override them, but rarely needs to do so.

If active matching is occurring, the nub next invokes the driver object's probe function. The probe function returns an IOService. This is usually the driver object itself, but the driver can return an instance of a different class, such as a specialized subclass included in the driver's kernel extension bundle. IOService's default implementation of probe simply returns the `this` pointer without altering the probe score. Overriding probe is optional; most drivers get enough information from property matching and don't need to override it. If you do override probe, however, you must make sure that the probe is not destructive, leaving the device in the state it found it. Hardware specifications generally define how to conduct non-destructive probes.

A driver's `start` function, just as with implementations of `probe`, should perform only the minimum necessary allocation of system resources to verify that it can operate the hardware. This conservative approach delays consumption of kernel resources until they're actually needed.

Each family, such as PCI, USB, or storage, defines a pair of activation and deactivation functions to indicate that the driver should prepare to service I/O requests and that the driver's services are no longer needed. These two functions are typically named `open` and `close`. Most drivers implement these functions to allocate and deallocate all of the necessary buffers and other structures in preparation for I/O processing.

Some families define additional levels of activation and deactivation. A networking driver, for example, does very little in `open` and `close`, instead performing setup and teardown in the `enable` and `disable` functions. Whatever the specific activation and deactivation functions, they can be invoked many times during a driver's life span; a driver should be able to function no matter how many times it gets activated or deactivated.

Driver Status Change

Another function that can be invoked many times during a driver's life span is the `message` function. This function informs the driver of important system status changes, such as when a disk is forcibly removed, when a power management change (sleep, wake-up) is occurring, or when the driver is being shut down. `IOService`'s implementation of this function does nothing and returns an "unsupported" result code. For more on the notification and messaging functionality provided by `IOService`, see [Notification and Messaging](#) (page 67)

Driver Shutdown

When a driver is going to be permanently shut down, its `message` function is invoked with a terminate message (`kIOMessageServiceIsTerminated`). If the driver accepts the termination, its `stop` function is then invoked. The driver should implement its `stop` function to close, release, or free any resources it opened or created in its `start` function, and to leave the hardware in the state the driver originally found it. Assuming the driver implements the activation and deactivation functions, there is usually little to do in the `stop` function. The final stage of driver shutdown is invocation of `free`, which occurs when the driver object's reference count reaches zero. In this function the driver can dispose of any resources it created in its `init` function.

Provider Matching

If you are implementing a provider driver object (that is, a subclass that is a member of an I/O Kit family) you may want to override `IOService`'s `matchPropertyTable` member function. When `IOService` performs matching for a driver object, it calls this method so the provider class can implement its own specific matching criteria in addition to the generic ones provided by `IOService`. The provider should examine the matching dictionary passed to see if it contains properties the family understands for matching and use them to match with the specific driver object if it does understand them.

Notification and Messaging

IOService provides two mechanisms for driver objects to communicate with each other and with the I/O Kit: notifications and messaging. Notifications are delivered to interested clients when a certain event occurs with an active service or driver that has properties matching a given dictionary. Messages are more targeted and flow in one direction, from provider to client. Any provider can send a message to any of its clients to notify it of some change in the runtime environment.

As discussed earlier in [Driver Object Life Cycle](#) (page 64) driver clients implement the `messagefunction` to receive and respond to messages from their providers. This function allows them to adapt to changes in the runtime environment. The messages can inform them of changes in system status, such as changes in power state, suspension of service, or impending service terminations. Providers implement the `messageClient`(or `messageClients`) functions to send messages by invoking their client's `message` methods. The I/O Kit defines some messages while others may be defined by families. See the header file `Kernel.framework/Headers/IOKit/IOMessage.h` for the generic messages that the `messageClient` and `messageClients` functions can deliver to a driver.

The broadcasting of notifications is a bit more complicated. Any driver object can install a notification handler through the `addNotification` or `installNotification` functions. The notification handler is set up to be invoked when a specific driver object (identified by a dictionary of matching properties) experiences a specific type of state change, such as when a driver is first published, matched at any time, or is terminated. Each notification handler is also given a priority number in case multiple notifications of the same type and for the same object are triggered at the same time.

The notification handler (of type `IOServiceNotificationHandler`) is invoked if any driver object whose personality matches the supplied matching dictionary changes to the specified state. For example, when a service provider calls `registerServices`, that not only starts the registration process but it also delivers notifications to all registered clients interested in the publication of the provider. The notification request is identified by an instance of an `IONotifier` object, through which the notification can be enabled, disabled, or removed.

Driver Accessors

IOService includes, as a convenience, a number of accessor member functions giving quick access to a driver object's state and the objects that are closely related to it. These functions return the following objects or values:

- The driver's state (`getState`), a bitfield indicating whether the driver is inactive, registered, matched, and so on
- The work loop being used by the driver (`getWorkLoop`) (see [Handling Events](#) (page 77)) for further information)

- The driver's primary provider (`getProvider`), as well as an `OSIterator` object for iterating over the driver's providers, if multiple (for example, a RAID device)
- The driver's primary client (`getClient`), as well as an `OSIterator` object for iterating over the driver's clients, if multiple

Other IOService Features

`IOService` incorporates functionality (other than that summarized above) that is useful for many categories of device driver. Most notably, this functionality includes the following features:

- **User client.** The `newUserClient` function creates an `IOUserClient`-based connection for communication with a non-kernel client; the client invokes this function by calling the `IOServiceOpen` function of the I/O Kit framework.
- **Device memory.** Several `IOService` member functions get, map, and set the physical memory ranges allocated to a memory-mapped device. These functions are intended for driver objects that are clients of PCI devices.
- **Interrupt handling.** `IOService` provides low-level functions for registering, unregistering, manipulating, and accessing interrupt handlers that are called at primary-interrupt time for a device's interrupt. The functions provide a mechanism for installing interrupt handlers that is not based on a work loop.

I/O Kit Families

In the I/O Kit, families are collections of classes that define and implement the abstractions common to all devices of a particular type. They provide the programmatic interfaces and generic support code for developing drivers that are members (providers) or clients of such families.

This chapter describes a number of concepts related to I/O Kit families:

- The relation of drivers to families
- Families as libraries, including the versioning and loading of libraries
- The programmatic structure of families and naming conventions

In addition, this chapter offers some tips for those who want to write their own I/O Kit families. For a reference to the current set of the I/O Kit families provided by Apple, see the appendix [I/O Kit Family Reference](#) (page 142)

Drivers and Families

An I/O Kit family is a library that implements some bus protocol (for example, PCI or USB) or some common set of services. But the support that a family provides is generic. A family does not include any of the details for getting at hardware because it cannot make assumptions about the specific hardware under the general layer it represents. It's the driver writer's responsibility to write code that bridges between the concrete and the abstract—that is, between the hardware and the abstraction defined by the family. A driver must extend a family to support specific hardware or to acquire specific features.

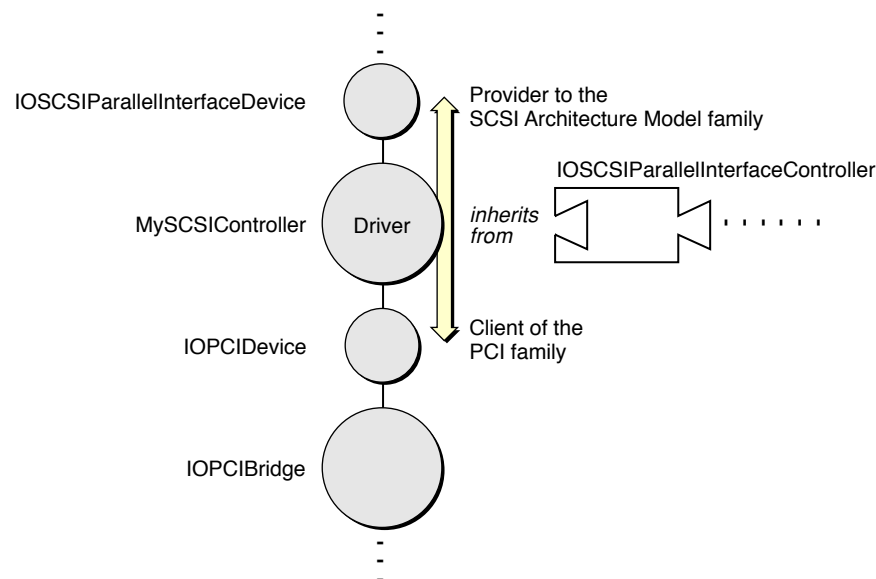
Take the SCSI Parallel family as an example. The SCSI Parallel family encapsulates the SCSI Parallel Interface-5 specification, which is well-defined. One of the things the specification describes is how to go about scanning the bus and detecting devices. Because this is an expensive operation, many SCSI Parallel controllers include firmware that can cache information about detected devices. To take advantage of this caching optimization, you could design your controller driver—member of the SCSI Parallel family—so that it overrides the scanning functionality to interact with the firmware.

Families commonly perform certain generic tasks, such as scanning buses, querying clients, queuing and validating commands, recovering from errors, and matching and loading drivers. Drivers do the tasks that impinge on hardware in some way. To continue with the example of the SCSI Parallel family, the primary job of the SCSI Parallel controller driver, as member of the SCSI Parallel family, is to receive SCSI commands from its family, execute each command on the hardware, and send a notification when the command completes.

Some I/O Kit families are clearly delimited by the specifications they encapsulate. Other families, such as the Audio family, are not as easily defined because there is no single specification prescribing what the family should include. In cases such as these, Apple carefully chose the set of abstractions to incorporate in the family to make it flexible and comprehensive enough. All families must advertise their capabilities and it is up to the higher levels of the driver stack to manage these capabilities.

A driver is both a provider and a client in its relationships to I/O Kit families. A driver that is a provider for a family (through its nub) is also a member of that family; it should inherit from a particular class in the family that describes the service it exports (however augmented). On the other hand, a driver is a client of the family whose service it imports (through a nub of the family). For example, a SCSI disk driver would inherit from the storage family rather than the SCSI Parallel family, to which it would be a client (see [Figure 6-1](#) (page 70)). A USB mouse driver would inherit from the HID (Human Interface Devices) family and would be a client of the USB family. A PCI Audio card driver would inherit from the Audio family and would be a client of the PCI family.

Figure 6-1 A driver's relationships with I/O Kit families



Families As Libraries

Families are implemented as libraries packaged as kernel extensions (KEXTs). They specify their defining attributes in an information property list and are installed in `/System/Library/Extensions`. Families are, mechanically, little different than ordinary drivers.

Two related characteristics distinguish a family from a driver. First, a driver expresses a dependency on a family using the `OSBundleLibraries` property; second, a family is loaded only as a byproduct of a driver listing it as a library. A driver specifies the libraries on which it depends as elements of the `OSBundleLibraries` dictionary. The I/O Kit guarantees that these libraries will be loaded into the kernel before it loads the driver and links it with its families. Note that libraries themselves declare the libraries (kernel extensions and the kernel itself) on which they depend using the `OSBundleLibraries` property.

You specify a library as a key-value pair in the `OSBundleLibraries` dictionary where the key is the bundle identifier (`CFBundleIdentifier`) of the library and the value is the earliest version of the library that the driver is compatible with. All versions are expressed in the 'vers' resource style. [Listing 6-1](#) (page 71) gives an example from the information property list of the `AppleUSBAudio` driver.

Listing 6-1 The `OSBundleLibraries` property

```
<key>OSBundleLibraries</key>
  <dict>
    <key>com.apple.iokit.IOAudioFamily</key>
    <string>1.0.0</string>
    <key>com.apple.iokit.IOUSBFamily</key>
    <string>1.8</string>
  </dict>
```

Although the I/O Kit loads the libraries before it loads the driver that specifies these dependencies, and loads libraries in proper dependency order, there is no guarantee about the order in which it loads libraries that have no interdependencies.

Generally, developers should declare dependencies for their device driver or any other kernel extension. (If the KEXT doesn't have an executable, dependency declaration is unnecessary.) What dependencies they need to declare depends on which symbols need to get resolved. If you include a header file of a family or other library, or if a header indirectly ends up including a library, you should declare that dependency. If you unsure whether a dependency exists, declare it anyway.

Library Versioning

To be available for loading and linking into the kernel, a family or other library has to declare its compatibility information using two properties: `CFBundleVersion` and `OSBundleCompatibleVersion`. The `CFBundleVersion` property defines the forward limit of compatibility—that is, the current version. The `OSBundleCompatibleVersion` property defines the backward limit of compatibility by identifying the last `CFBundleVersion`-defined version of the library that broke binary compatibility with prior versions.

Every time you revise a driver or a family, you should increment your `CFBundleVersion` value appropriately. You reset the `OSBundleCompatibleVersion` value (to the current `CFBundleVersion`) only when the revision makes the binary incompatible with prior versions, as when you remove a function or other symbol, or change a class such that the vtable layout changes. If you are writing an I/O Kit family, make sure that you specify an `OSBundleCompatibleVersion` property for your library; otherwise, drivers and other kernel extensions cannot declare a dependency on it and thus cannot link against it.

For both drivers and families (and, indeed, all kernel extensions), make sure that you also set the version in the kernel module and that this value is equivalent to the `CFBundleVersion` in the information property list. You set the version in the executable through the `MODULE_VERSION` setting in Xcode, in the target's Customized Settings list (you find this in the target's Build view).

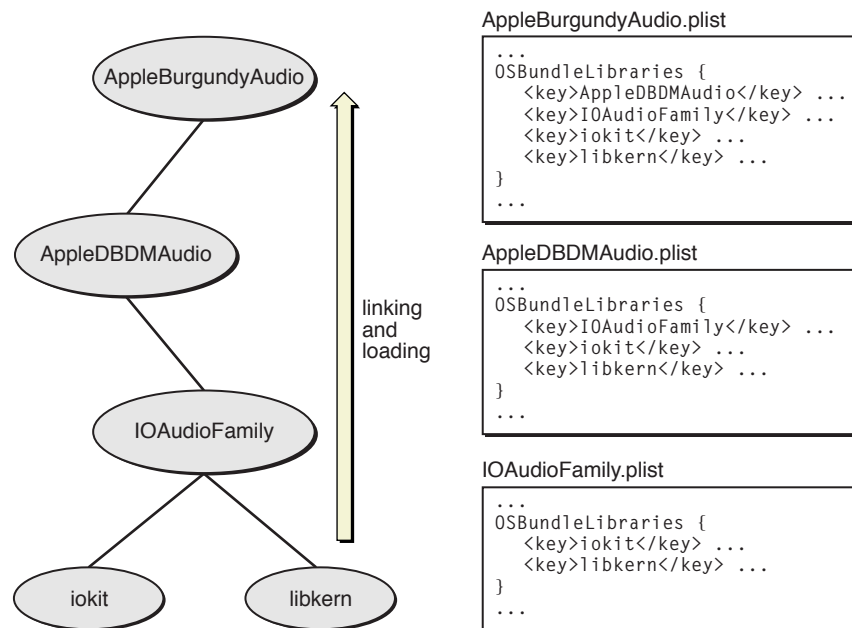
Library Loading

The KEXT manager functions as the kernel loader and linker. At boot time or whenever the system detects a newly attached device, the I/O Kit kicks off the matching process to find a suitable driver for a device. When such a driver is found, it is the KEXT manager's job to load the driver into the kernel and link it with the libraries on which it depends.

But before it can do this, the KEXT manager must ensure that those libraries, and all the other libraries on which those libraries depend, are loaded first. To do this, the manager builds a dependency tree of all libraries and other kernel modules required for the driver. It builds this tree using the contents of the `OSBundleLibraries` property, first of the driver and then of each required library.

After it builds the dependency tree, the KEXT manager checks if the libraries that are the most remote from the driver in the tree are already loaded. If any of these libraries is not loaded, the manager loads it and calls its start routine (the routine varies according to type of KEXT). It then proceeds up the dependency tree in similar fashion—linking, loading, and starting—until all required libraries have been linked and loaded. See [Figure 6-2](#) (page 73) for an illustration of this procedure.

Figure 6-2 OSBundleLibraries and the dependency tree



If the KEXT manager encounters a problem initializing a library, or it doesn't find a library with a compatible version (based on the value of `OSBundleCompatibleVersion`), it stops and (usually) returns a failure code. The modules already loaded stay loaded for awhile. Generally, unloading of modules does not happen immediately when they are not used. The I/O Kit includes a feature that tracks idle time and unloads modules after a certain period of idleness.

Important: The only way to load a kernel extension explicitly is to use the `kext load` command-line utility.

The Programmatic Structure of Families

Although I/O Kit families tend to be quite different from each other, they have some structural elements in common. First, `IOService` is the common superclass for all I/O Kit families; at least one important class in each family, and possibly more, inherits from `IOService` (see [The I/O Kit Base Classes](#) (page 61) for more information). And each family has one or more classes that present an interface to drivers.

Typical Classes

A family typically defines two classes for drivers:

- A class describing the nub interface for drivers that are clients of the family
- A superclass for drivers that are members of the family, and thus providers to its nubs

In other words, I/O Kit families usually define an upward interface and a downward interface. These interfaces are required for the layering of driver objects involved in an I/O connection. The upward interface—the nub interface—presents to the rest of the system the hardware abstractions and rule definitions encapsulated by the family. The downward interface provides the subclassing interface for member drivers. Together, the interfaces define the up calls into the family and the down calls that member drivers are expected to make.

In addition to these two classes, families typically define a number of utility classes and support classes. The appendix [I/O Kit Family Reference](#) (page 142) describes some of these classes.

Some families specify subclasses for particular varieties of client or member drivers. The Storage family, for example, defines a generic block storage class for nub objects (`IOBlockStorageDevice`) and then also provides specific subclasses for certain varieties: `IOCDBlockStorageDevice` and `IODVDBlockStorageDevice`. In addition, families can include classes for device interfaces (as subclasses of `IOUserClient`) as well as commands specific to the family (as subclasses of `IOCommand`). Families can also have various helper classes and header files for family-specific type definitions.

Some families do not include a public nub or provider class for drivers when there is little need for such drivers. And Apple has not provided families for all types of hardware. If you find that the I/O Kit does not have a family or interface for your needs, you can always create a driver that inherits directly from `IOService`. Such “family-less” drivers are sometimes necessary if the potential applications for the driver are few. They must incorporate the abstractions and range of functionality found in families as well as the hardware-specific code typical of drivers. Besides directly inheriting from `IOService`, family-less drivers frequently make use of the I/O Kit helper classes such as `IOWorkLoop`, the event-source classes, `IOMemoryCursor`, and `IOMemoryDescriptor`.

Naming and Coding Conventions

Generally, Apple’s position on class naming within families is that the name should indicate what the class represents. Often, this name is dictated by the specification for the hardware. For example, the PCI family defines the `IOPCIBridge` class for drivers that are providers for the family. The reason for this name is simple: the PCI bridge (as the specification makes clear) is what the PCI controller drivers control. When there is no clear naming precedent for a family’s classes, the I/O Kit follows a naming convention of `IOFamilyNameDevice` for nub (client) classes and `IOFamilyNameController` for provider classes.

Important: The general guideline of naming classes for what they represent applies equally to drivers. Drivers should be named for the device they control (but should not have the redundant suffix “Driver”).

If you are writing your own I/O Kit family, Apple recommends that you follow the same naming guidelines for your classes. And there are a few other general naming conventions to be aware of. Each class, function, type, and so on should have prefix that designates the vendor writing the software. Be sure not to use any of the prefixes that Apple reserves for itself (Table 6-1 (page 75)).

Table 6-1 API prefixes reserved by Apple

Prefix	Meaning
OS, os	libkern or other kernel service
IO, io	I/O Kit or I/O Kit family
MK, mk, mach_	Mach kernel
Apple, APPLE, apple, AAPL, aapl, com_apple_	Apple hardware support (for example, Apple-provided drivers)

In addition, private, internal symbols should have an underscore “_” prefix, following the convention used by Apple. Do not access these private APIs from a KEXT. As with drivers, use of reverse DNS notation (substituting underscores for periods) is highly recommended to avoid naming conflicts.

Creating An I/O Kit Family

There might be occasions when you deem it worthwhile to write your own I/O Kit family. Usually this happens when there is a standard or protocol for which no family exists, and you discern a need for interoperability among drivers for devices based on this protocol or standard. An example might be the IEEE488 standard for plotters and lab equipment.

If you decide to implement a family, here are a few guidelines to help you:

- At the beginning, write the family and driver code together; don’t worry yet about the division of functionality and interface between driver and family. Just concentrate on coming up with a good object-oriented design, determining what objects are necessary and what relationships they should have.
- After you have a working driver and have solved the stack for a particular device, separate the family code from the hardware-specific code. One approach that might be useful for locating family-generic code, especially for complex families, is to write two or more drivers for different hardware and then abstract away the common code.

- Define what the family's nub objects look like to drivers—that is, the APIs your clients will see. To do this, look at the specification and encapsulate the important features (it's not necessary to include rarely or never-used features). Keep in mind that the nubs of most families do very little. Most often they encapsulate addressing and arbitration details.
- Define the superclass for drivers that will be members of your family.
- Keep the layering separation of a family airtight. A family should not include headers from any other family or driver and should not define the superclass of clients.

There can also be situations that might call for the creation of a “superfamily”: a family that extends an existing family in a way similar to a subclass, but with a big difference; its aim is generality rather than specificity. Third-party vendors might want to have a superfamily to contain the code common to drivers based on different bus protocols. This would eliminate the need to load code that isn't needed. For example, a mouse vendor might have a driver capable of driving both USB and ADB mice. If a system requires a USB mouse, you don't want to have the ADB-specific code loaded as well. Thus the vendor might write a superfamily that acts as a service library; it would separate out the layers of code specific to a bus protocol into subfamilies and put the remaining code into the superfamily. Only the code specific to the currently used bus would be loaded.

Handling Events

A device driver works in perhaps the most chaotic of environments within an operating system. Many devices require a sequence of commands to perform a single I/O operation. However, multiple threads can enter the driver's code at any place and at any time, whether through client I/O requests, hardware interrupts, or timeout events; on a multiprocessor system, threads can even be running concurrently. The driver must be prepared to handle them all.

This flurry of incoming threads, each with its own event, poses some problems for a driver. The driver needs a way to protect its data structures from access by different threads because such simultaneous access can lead to data corruption, or worse. It needs to guarantee exclusive access to a thread for any single command or operation that must complete in order to preserve the integrity of the driver's data, or to prevent a deadlock or race condition. In the I/O Kit, this protection is provided by the `IOWorkLoop` class and its attendant event-source classes.

Work Loops

An `IOWorkLoop` object (or simply, a work loop) is primarily a gating mechanism that ensures single-threaded access to the data structures used by hardware. For some event contexts, a work loop is also a thread. In essence, a work loop is a mutually exclusive (mutex) lock associated with a thread. It does several things:

- Its gating mechanism synchronizes the actions among event sources.
- It provides a stackable environment for event handling.
- It spawns a dedicated thread for the completion of indirect interrupts delivered by the interrupt controller. This mechanism serializes interrupt handling for the work loop's driver, preventing simultaneous access to driver data by multiple interrupts.

To put the role of the work loop in perspective, it helps first to consider the event sources that it is designed for. In the I/O Kit there are five broad categories of asynchronous events:

- Interrupt events—indirect (secondary) interrupts originating from devices
- Timer events—events delivered periodically by timers, such as timeouts
- I/O commands—I/O requests issued by driver clients to their providers
- Power events—typically generated through calls down the driver stack

- Structural events—typically events involving the I/O Registry

The I/O Kit provides classes to handle these event sources: `IOInterruptEventSource`, `IOTimerEventSource`, and `IOCommandGate`. (You handle power and structural events using the mechanism provided by `IOCommandGate` objects.) Each of the event-source classes defines a mechanism specific to an event type for invoking a single function within the protected context of the work loop. If a thread carrying an event needs access to a driver's critical data, it must do so through an object of one of these classes.

Generally, client drivers set up their work loops, event sources, and event handlers in their `start` function. In order to avoid deadlocks and race conditions, all code that accesses the same data should share a single work loop, registering their event sources with it so that a single gating mechanism is used. Work loops can be safely shared among unrelated objects, of course, and often are shared by objects at different levels in a single driver stack. Work loops can also be dedicated for use by a particular driver and its clients. See [Shared and Dedicated Work Loops](#) (page 79) for more information.

The I/O Kit work-loop mechanism offers functionality roughly similar to that of the Vertical Retrace Manager, the Time Manager, and the Deferred Task Manager of Mac OS 9.

Work Loop Architecture

The I/O Kit's work-loop mechanism mitigates the performance penalty exacted by context switching, a by-product of the underlying event-handling model commonly used in some operating systems. To guarantee a single-threaded context for event handling, this model completes everything on one thread. Unfortunately, the transfer of work to the thread requires a switch in the context of the event-bearing thread. More precisely, when this thread goes from a running context to a non-running context, its current register state must be saved. When the secure thread completes its work, the state of the originating thread is restored and control branches back to the function originally referenced by the thread. This switching back and forth between thread contexts consumes cycles.

The work-loop model works quite differently for I/O commands and timer events. In these instances, the thread of the respective event source simply grabs the mutex lock held by the work loop. No other event from any source can be processed until the `Action` routine for the current event returns. Although the lock is mutually exclusive, it doesn't prevent reentrancy. Also, you can have multiple work loops in the same driver stack, and this increases the possibility of deadlock. However, work loops do avoid self-deadlocks because they are based on a recursive lock: They always check to see if they are the thread that currently own the lock.

The way the I/O Kit manages interrupt event sources does involve context switching. The completion routines for interrupts run on the work loop's thread and not on the thread delivering the interrupt. Context switching is required in this case because the interrupt controller must immediately dispatch direct (primary) interrupts to other threads to run the completion routines for those interrupts. See [Handling Interrupts](#) (page 83) for more information.

Two factors influence the order in which a work loop queries its event sources for work. The relative priority of threads is the main determinant, with timers having the highest priority. A client thread can modify its own priority and thereby expedite the delivery of I/O requests (it might not affect how soon they are processed, however, because I/O requests are usually queued in FIFO order). For interrupt event sources, which also have a relatively high priority, the order in which they are added to the work loop determines the order in which they are queried for work. See [Handling Interrupts](#) (page 83) for further details.

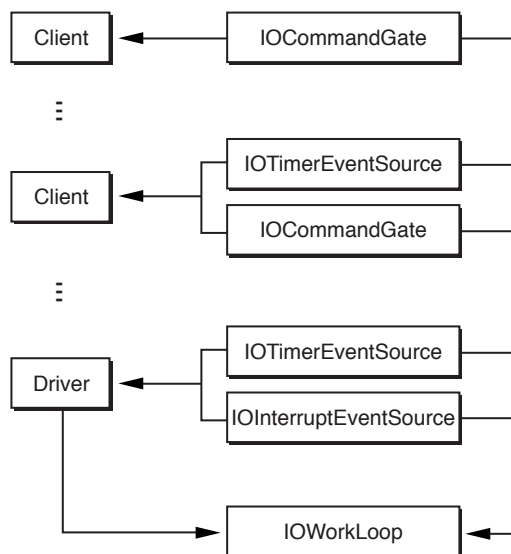
Regardless of event source and mechanism, a work loop is primarily used for one thing: to run the completion or `Action` routines specified by the event source. It guarantees that the routine handling an event is the only one running at any given time. This aspect of work loops raises a design point. When a thread is running code to handle an event, other events can be asynchronously delivered to their event sources, but they cannot be processed until the handler returns. Therefore event handlers should not attempt to complete large chunks of work or do anything that might *block* (that is, wait for some other process to complete), such as allocating memory or other resources. Instead they should, if possible, queue up the work or otherwise defer it for later processing.

Shared and Dedicated Work Loops

All I/O Kit services can easily share their provider's work loop. The base of the driver Registry, representing the logic board of the computer, always contains a work loop, so a driver is assured of having a work loop even if it doesn't create one itself. All a driver needs to do is call the IOService function `getWorkLoop` to access its provider's work loop.

In this way, an entire stack of driver objects, or a subset of such objects, can share one work loop. [Figure 7-1](#) (page 80) shows how a work loop shared by multiple driver objects uses event sources to manage access to its gating mechanism.

Figure 7-1 Driver objects sharing a work loop



Most drivers won't create their own work loop. If hardware doesn't directly raise interrupts in your driver, or if interrupts rarely occur in your driver, then you don't need your own work loop. However, a driver that takes direct interrupts—in other words, that interacts directly with the interrupt controller—should create its own dedicated work loop. Examples of such drivers are PCI controller drivers (or any similar driver with a provider class of `IOPCIDevice`) and RAID controller drivers. Even these work loops may be shared by the driver's clients, however, so it's important to realize that in either case, the driver must not assume that it has exclusive use of the work loop. This means that a driver should rarely enable or disable all events on its work loop, since doing so may affect other I/O Kit services using the work loop.

If a driver handles interrupts or for some other reason needs its own work loop, it should override the `IOService` function `getWorkLoop` to create a dedicated work loop, used by just the driver and its clients. If `getWorkLoop` isn't overridden, a driver object gets the next work loop down in its stack.

Examples of Obtaining Work Loops

To obtain a work loop for your client driver, you should usually use your provider's work loop or, if necessary, create your own. To obtain your provider's work loop, all you have to do is call the `IOService` function `getWorkLoop` and retain the returned object. Immediately after getting your work loop you should create your event sources and add them to the work loop (making sure they are enabled).

To create a dedicated work loop for your driver, override the `getWorkLoop` function. [Listing 7-1](#) (page 81) illustrates a thread-safe implementation of `getWorkLoop` that creates the work loop lazily and safely.

Listing 7-1 Creating a dedicated work loop

```
protected:
    IOWorkLoop *cntrlSync; /* Controllers Synchronizing context */
// ...
IOWorkLoop * AppleDeviceDriver::getWorkLoop()
{
    // Do we have a work loop already?, if so return it NOW.
    if ((vm_address_t) cntrlSync >> 1)
        return cntrlSync;

    if (OSCompareAndSwap(0, 1, (UInt32 *) &cntrlSync)) {
        // Construct the workloop and set the cntrlSync variable
        // to whatever the result is and return
        cntrlSync = IOWorkLoop::workLoop();
    }
    else while ((IOWorkLoop *) cntrlSync == (IOWorkLoop *) 1)
        // Spin around the cntrlSync variable until the
        // initialization finishes.
        thread_block(0);

    return cntrlSync;
}
```

This code first checks if `cntrlSync` is a valid memory address; if it is, a work loop already exists, so the code returns it. Then it tests to see if some other thread is trying to create a work loop by atomically trying to compare and swap the controller synchronizer variable from 0 to 1 (1 cannot be a valid address for a work loop). If no swap occurred, then some other thread is initializing the work loop and so the function waits for the `cntrlSync` variable to stop being 1. If the swap occurred then no work loop exists and no other thread is in the process of creating one. In this case, the function creates and returns the work loop, which unblocks any other threads that might be waiting.

As you would when getting a shared work loop, invoke `getWorkLoop` in `start` to get your work-loop object (and then retain it). After creating and initializing a work loop, you must create and add your event sources to it. See the following section for more on event sources in the I/O Kit.

Event Sources

A work loop can have any number of event sources added to it. An event source is an object that corresponds to a type of event that a device driver can be expected to handle; there are currently event sources for hardware interrupts, timer events, and I/O commands. The I/O Kit defines a class for each of these event types:

`IOInterruptEventSource`, `IOTimerEventSource`, and `IOCommandGate`, respectively. Each of these classes directly inherits from the abstract class `IOEventSource`.

An event-source object acts as a queue for events arriving from a particular event source and hands off those events to the work-loop context when it asks them for work. When you create an event-source object, you specify a callback function (also known as an “action” function) to be invoked to handle the event. Similar to the Cocoa environment’s target/action mechanism, the I/O Kit stores as instance variables in an event source the target of the event (the driver object, usually) and the action to perform. The handler’s signature must conform to an `Action` prototype declared in the header file of the event-source class. As required, the work loop asks each of its event sources in turn (by invoking their `checkForWork` function) for events to process. If an event source has a queued event, the work loop runs the handler code for that event in its own protected context. Note that when you register an event source with a work loop, the event source is provided with the work loop’s signaling semaphore, which it uses to wake the work loop. (For more information on how the work loop sleeps and wakes, see the `threadMain` function in `IOWorkLoop` documentation.)

A client driver, in its activation phase (usually the `start` function), creates the event sources it needs and adds them to its work loop. The driver must also implement an event handler for each event source, ensuring that the function’s signature conforms to the `Action` function prototype defined for the event-source class. For performance reasons, the event handler should avoid doing anything that might block (such as allocating memory) and defer processing of large amounts of data. See [Work Loops](#) (page 77) for further information on event priority and deferring work in event handlers.

The procedure for adding event sources to a work loop is similar for each type of event source. It involves four simple steps:

1. Obtain your work loop.
2. Create the event-source object.
3. Add the object to the work loop.
4. Enable the event source.

Disposing of an event source also has a common procedural pattern:

1. Disable the event source.
2. Remove it from the work loop.
3. Release the event source.

The following sections discuss the particulars of each event source and give examples specific to each kind.

Handling Interrupts

Interrupts are typically the most important type of event that drivers handle. They are the way that devices attached to a computer inform the operating system that an asynchronous action has occurred and that, consequently, they have some data. For example, when the user moves a mouse or plugs a Zip drive into a USB port, a hardware interrupt is generated and the affected driver is notified of this event. This section discusses interrupt handling in the I/O Kit, with particular attention to the role played by objects of `IOInterruptEventSource` and its subclasses.

Interrupt Handling in the I/O Kit

The I/O Kit's model for interrupt handling does not conform to the standard UNIX model. I/O Kit drivers nearly always work in the indirect-interrupt context instead of dealing with direct interrupts, as does the UNIX model. Indirect interrupts are less restrictive and permit the Mach scheduler to do its job. (Indirect interrupts are sometimes known as secondary interrupts and direct interrupts as primary interrupts.) The difference between the two types of interrupts has to do with the context in which the interrupt is dealt with.

Two types of events trigger an interrupt:

- Command-based events, such as incoming networking packets and reads of storage media
- Asynchronous events, such as keyboard presses

When an interrupt occurs, a specific interrupt line is set and, once the interrupted thread finishes the current instruction, control branches to the interrupt controller registered with the Platform Expert. When the interrupt controller receives the interrupt, its thread becomes that of the direct (primary) interrupt. There is typically only one direct interrupt in the system at any one time, and the direct-interrupt context has the highest priority in the system. The following list indicates the relative priorities of threads in the system:

1. Direct interrupt
2. Timers and page-out
3. Real time (multimedia)
4. Indirect interrupts (drivers)
5. Window Manager
6. User threads (including I/O requests)

Because of its extremely high priority, the direct-interrupt context has a design responsibility to hand off the interrupt to lower-priority threads as soon as possible. The interrupt controller must decode why the interrupt was taken, assign it to the appropriate driver object, and return.

In the direct-interrupt model, the target driver assumes the context carrying the direct interrupt. It must handle the interrupt in this highest-priority context. The problem with direct interrupts is that they can be neither lowered in priority nor preempted. All other interrupts are effectively disabled until the current interrupt is handled. Direct interrupts especially don't scale well in the OS X multiprocessing environment.

With indirect interrupts, the interrupt controller dispatches the interrupt it reads off the interrupt line to the appropriate interrupt event-source object of the target driver, effectively causing it to schedule on the driver's work-loop thread. The completion (or `Action`) routine defined by the event source is then run on the work-loop thread to handle the interrupt. The priority of the work-loop thread, although high compared to most client threads, is lower than the thread carrying the direct interrupt. Thus the completion routine running in the work-loop thread can be preempted by another direct interrupt.

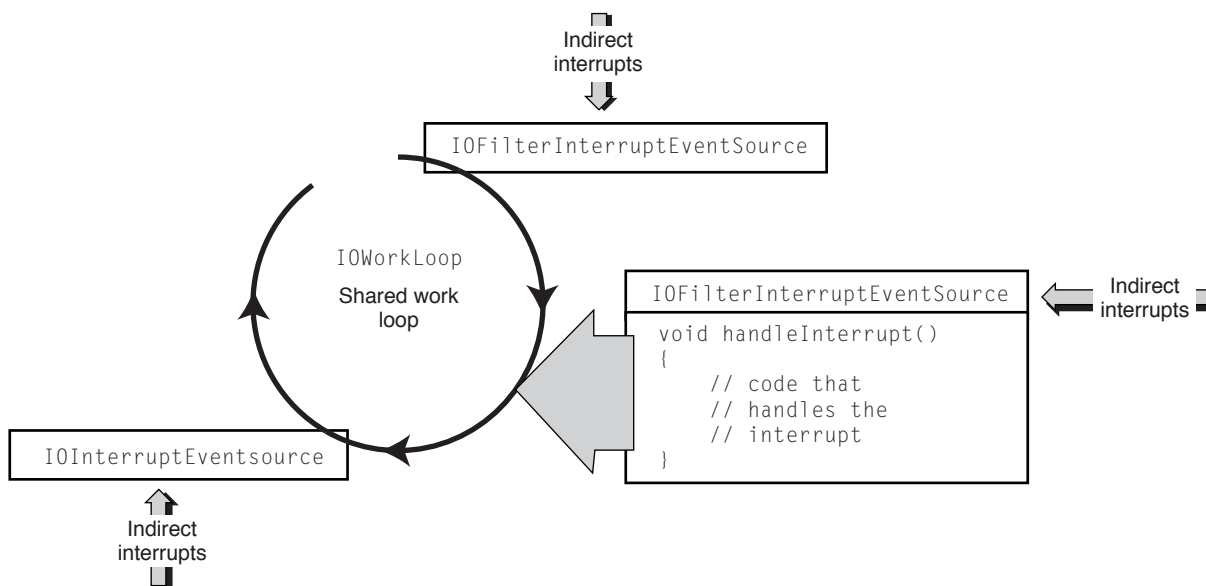
The I/O Kit does not prohibit access to the direct-interrupt context, and in fact provides a separate programming interface for this purpose (see [Using Interrupt Handlers With No Work Loops](#) (page 88)). However, use of direct interrupts is strongly discouraged.

A work loop can have several `IOInterruptEventSource` objects attached to it. The order in which these objects are added to the work loop (through `IOWorkLoop's addEventSource` function) determines the general order in which interrupts from different sources are handled.

[Figure 7-2](#) (page 85) illustrates some of these concepts. It shows events originating from different sources being delivered to the corresponding event-source objects "attached" to the work loop. As with any event-source object, each interrupt event source acts as a queue for events of that type; when there is an event in the queue, the object signals the work loop that it has work for it. The work loop (that is, the dedicated thread) awakes and queries each installed event source in turn. If an event source has work, the work loop runs the completion routine for the event (in this case, an interrupt) in its own protected thread. The previous thread—the client

thread running the event-source code—is blocked until the routine finishes processing the event. Then the work loop moves to the next interrupt event source and, if there is work, runs the completion routine for that interrupt in its protected context. When there is no more work to do, the work loop sleeps.

Figure 7-2 A work loop and its event sources



Remember that the order in which you add interrupt event sources to a work loop determines the order of handling for specific interrupt events.

Setting Up an Interrupt Handler Attached to a Work Loop

A driver typically creates an interrupt event-source object—generally of the `IOInterruptEventSource` or `IOFilterInterruptEventSource` class—in its start function by calling the factory creation method for the class (for example, `interruptEventSource`). This method specifies the driver itself as a target and identifies an action member function (conforming to the `Action` type defined for the event-source class) to be invoked as the completion routine for the event source. The factory method also associates the driver with a provider that deals with the hardware interrupt facility (usually a nub such as an `IOPCIDevice`). The driver then registers the event source with the work loop through `IOWorkLoop`'s `addEventSource` function.

[Listing 7-2](#) (page 85) provides an example for setting up an interrupt event source.

Listing 7-2 Adding an interrupt event source to a work loop

```
myWorkLoop = (IOWorkLoop *)getWorkLoop();

interruptSource = IOInterruptEventSource::interruptEventSource(this,
```

```
(IOInterruptEventAction)&MyDriver::interruptOccurred,  
provider);  
  
if (!interruptSource) {  
    IOLog("%s: Failed to create interrupt event source!\n", getName());  
    // Handle error (typically by returning a failure result).  
}  
  
if (myWorkLoop->addEventSource(interruptSource) != kIOReturnSuccess) {  
    IOLog("%s: Failed to add interrupt event source to work loop!\n",  
        getName());  
    // Handle error (typically by returning a failure result).  
}
```

In this example, if you do not specify a provider in the `interruptEventSource` call, the event source assumes that the client will call `IOInterruptEventSource`'s `interruptOccurred` method explicitly. Invocation of this function causes the safe delivery of asynchronous events to the driver's `IOInterruptEventSource`.

Events originating from direct interrupts are handled within the work loop's thread, which should never block indefinitely. This specifically means that the completion routines that handle interrupts, and any function they invoke, must not allocate memory or create objects, as allocation can block for unbounded periods of time.

You destroy an interrupt event source in a driver's deactivation function (usually `stop`). Before you release the `IOInterruptEventSource` object, you should disable it and then remove it from the work loop. [Listing 7-3](#) (page 86) gives an example of how to do this.

Listing 7-3 Disposing of an `IOInterruptEventSource` object

```
if (interruptSource) {  
    interruptSource->disable();  
    myWorkLoop->removeEventSource(interruptSource);  
    interruptSource->release();  
    interruptSource = 0;  
}
```

Filter Interrupt Event Sources

The I/O Kit supports shared interrupts, where drivers share a single interrupt line. For this purpose it defines the `IOFilterInterruptEventSource` class, a subclass of `IOInterruptEventSource`. Apple highly recommends that third-party device driver writers base their interrupt event sources on the `IOFilterInterruptEventSource` class instead of the `IOInterruptEventSource` class. The latter class does not ensure that the sharing of interrupt lines is safe.

The `IOFilterInterruptEventSource` class follows the same model as its superclass except that it defines, in addition to the `Action` completion routine, a special callback function. When an interrupt occurs the interrupt invokes this function for each driver sharing the interrupt line. In this function, the driver responds by indicating whether the interrupt is something that it should handle.

Important: The filter routine runs in the *primary (hardware) interrupt context*. You should do as little work as possible in this routine—just enough to determine whether your device produced the interrupt. For more information, see the documentation for `IOFilterInterruptEventSource`.

[Listing 7-4](#) (page 87) shows how to set up and use an `IOFilterInterruptEventSource`.

Listing 7-4 Setting up an `IOFilterInterruptEventSource`

```
bool myDriver::start(IOService * provider)
{
    // stuff happens here

    IOWorkLoop * myWorkLoop = (IOWorkLoop *) getWorkLoop();
    if (!myWorkLoop)
        return false;

    // Create and register an interrupt event source. The provider will
    // take care of the low-level interrupt registration stuff.
    //
    interruptSrc =
        IOFilterInterruptEventSource::filterInterruptEventSource(this,
            (IOInterruptEventAction) &myDriver::interruptOccurred,
            (IOFilterInterruptAction) &myDriver::checkForInterrupt,
            provider);
    if (myWorkLoop->addEventSource(interruptSrc) != kIOReturnSuccess) {
        IOLog("%s: Failed to add FIES to work loop.\n", getName());
    }
}
```

```
    }  
    // and more stuff here...  
}  
  
bool myDriver::checkForInterrupt(IOWFilterInterruptEventSource * src)  
{  
    // check if this interrupt belongs to me  
  
    return true; // go ahead and invoke completion routine  
}  
  
void myDriver::interruptOccurred(IOWInterruptEventSource * src, int cnt)  
{  
    // handle the interrupt  
}
```

If your filter routine (the `checkForInterrupt` routine in [Listing 7-4](#) (page 87)) returns `true`, the I/O Kit will automatically start your interrupt handler routine on your work loop. The interrupt will remain disabled in hardware until your interrupt service routine (`interruptOccurred` in [Listing 7-4](#) (page 87)) completes.

Note: In some cases, such as the implementation of pseudo-DMA, this behavior may not be desirable. In this case, you may choose to have your filter routine schedule the work on the work loop itself and then return `false`. If you do this, the interrupt will not be disabled in hardware and you could receive additional primary interrupts before your work loop–level service routine completes. Because this scheme has implications for synchronization between your filter routine and your interrupt service routine, you should avoid doing this unless your driver requires pseudo-DMA.

Using Interrupt Handlers With No Work Loops

The `IOService` class provides member functions for registering interrupt handlers that operate outside of the work-loop mechanism. These handlers can be invoked in a direct interrupt context and must call the interrupt management code of a provider such as an `IOPCIDevice` nub. Only one handler can be installed per interrupt source. It must be prepared to create and run its own threads and do its own locking.

Few drivers need to use interrupt handlers that are created and controlled in this way. One example where such an interrupt handler is justified is a multifunction card that needs to route direct interrupts to drivers. If you take this course, be careful. Very few system APIs are safe to call in the direct-interrupt context.

More information on this subject is forthcoming.

Handling Timer Events

Device drivers occasionally need to set timers, usually to implement a timeout so the driver can determine if an I/O request doesn't complete within a reasonable period. The `IOTimerEventSource` class is designed for that purpose.

Important: The absolute accuracy of timeouts in the I/O Kit cannot be guaranteed. The Mach scheduler can always run a higher priority thread, which might delay the execution of the timer `Action` routine.

A driver creates an `IOTimerEventSource` with a callback `Action` function and a time at which to invoke that function, and then registers it with the work loop to run on. When the timeout passes, the event source is scheduled with the work loop. When the work loop queries it for work, the event source closes the work-loop gate (by taking the work loop's lock), invokes the callback function, and then releases the work-loop lock to open the gate. [Listing 7-5](#) (page 89) show how to create and register a timer event source.

Listing 7-5 Creating and registering a timer event source

```
myWorkLoop = (IOWorkLoop *)getWorkLoop();

timerSource = IOTimerEventSource::timerEventSource(this,
    (IOTimerEventSource::Action)&MyDriver::timeoutOccurred);

if (!timerSource) {
    IOLog("%s: Failed to create timer event source!\n", getName());
    // Handle error (typically by returning a failure result).
}

if (myWorkLoop->addEventSource(timerSource) != kIOReturnSuccess) {
    IOLog("%s: Failed to add timer event source to work loop!\n", getName());
    // Handle error (typically by returning a failure result).
}

timerSource->setTimeoutMS(MYDRIVER_TIMEOUT);
```

Often a driver wants to set the timer and issue an I/O request at the same time. If the I/O request completes before the timer event is triggered, the driver should cancel the timer immediately. If a timer event is triggered first, the driver typically reissues the time-out I/O request (at which time it resets the timer).

If you want the timer event to be recurrent, you should reset the timer to the desired interval in the `Action` handler. The `IOTimerEventSource` class does not have a mechanism for setting periodic timers. The class does provide a few functions for setting relative and absolute timer intervals at various granularities (nanoseconds, microseconds, and so on). The code fragment in [Listing 7-5](#) (page 89) uses `setTimeoutMS` to set the timer with a specific time-out millisecond interval.

Events originating from timers are handled by the driver's `Action` routine. As with other event handlers, this routine should never block indefinitely. This specifically means that timer handlers, and any function they invoke, must not allocate memory or create objects, as allocation can block for unbounded periods of time.

To dispose of a timer event source, you should cancel the pending timer event before removing the event source from the work loop and releasing it. [Listing 7-6](#) (page 90) illustrates how you might do this in your driver's deactivation function.

Listing 7-6 Disposing of a timer event source

```
if (timerSource) {  
    timerSource->cancelTimeout();  
    myWorkLoop->removeEventSource(timerSource);  
    timerSource->release();  
    timerSource = 0;  
}
```

I/O Requests and Command Gates

Driver clients use `IOCommandGate` objects to issue I/O requests to a driver. A command gate controls access to the work-loop lock, and in this way it serializes access to the data involved in I/O requests. It does not require a thread context switch to ensure single-threaded access. An `IOCommandGate` event source simply takes the work-loop lock before it runs its `Action` routine; by doing so, it prevents other event sources on the same work loop from scheduling. This makes it an efficient mechanism for I/O transfers.

Note that nub classes usually define `Action` functions for their own clients to use, so that driver classes don't have to use command gates themselves.

Up Calls and Down Calls

Calls originated by clients through a command gate are known as down calls. These always originate in a thread other than the work loop's context, and so they may safely block without causing a deadlock (as long as they don't hold the work-loop gate). All allocation should occur on the down-call side of an I/O request before the command gate is closed.

Up calls, which are originated by an interrupt or timer event, occur within the work loop's context and should never block indefinitely. This specifically means that interrupt and timeout handlers, and any function they invoke, must not allocate memory or create objects, as allocation can block and, as a potential consequence, cause a paging deadlock.

It's possible for an up call to result in a client notification that immediately results in another I/O request through the command gate. A work loop can handle recursive closing of its gate by the same thread, so this situation never results in deadlock. However, because the new request is occurring on the context of an up call, that request cannot block; this concern belongs to the system client making the I/O request, though, so you need never worry about this as a driver developer.

Setting Up and Using Command Gates

Prior to closing a command gate, you should adequately prepare the I/O request. An I/O request involves three things: the command itself (which is family-specific), the memory involved in the transfer (defined as an `IOMemoryDescriptor` object), and the function to call to process the request within the context of the command gate. See [Managing Data](#) (page 95) for information on `IOMemoryDescriptors` and related objects.

Command gates should be closed for the briefest possible period, during which the least amount of work possible is performed. The longer a command gate holds the work-loop lock, the greater the likelihood of contention. As with all event sources, the command-gate function should not allocate memory or any other unbounded resource because of the danger of blocking. Instead, the client should preallocate the required resources before control is transferred to the work-loop context. For example, it could allocate a pool of resources in its `start` function.

You create an `IIOCommandGate` object by calling the `commandGate` factory method, specifying as parameters the object "owner" of the event source (usually `this`) and a pointer to a function conforming to the `Action` prototype. You then register the command gate with the client's work loop using `IOWorkLoop`'s `addEventSource` function. [Listing 7-7](#) (page 91) gives an example of this procedure.

Listing 7-7 Creating and registering a command gate

```
workLoop = (IOWorkLoop *)getWorkLoop();
commandGate = IIOCommandGate::commandGate(this,
                                           (IIOCommandGate::Action)receiveMsg);
```

```
if (!commandGate ||
    (workLoop->addEventSource(commandGate) != kIOReturnSuccess) ) {
    kprintf("can't create or add commandGate\n");
    return false;
}
```

The `IOCommandGate` class provides two alternatives for initiating the execution of an I/O request within the command gate. One is the `runCommand` member function and the other is the `runAction` member function. These functions work similarly. When the client wishes to invoke the `Action` function, rather than invoking it directly, it invokes the command gate's `runCommand` or `runAction` function, passing in all required arguments. The command gate then grabs the work-loop lock (that is, it closes the command gate), invokes the `Action` function, and then opens the gate.

Where the two functions differ is in their flexibility. The `runCommand` function makes use of the same target/action mechanism used by the other event-source classes. In this mechanism, the created `IOCommandGate` object encapsulates (a pointer to) an `Action` function as well as the target (or “owner”) object that implements this function. In this model, only one `Action` function can be invoked for an I/O request.

However, a driver often has to deal with multiple sources of I/O requests. If this is the case, you can use the `runAction` function to issue I/O requests in multiple command gates. This function lets you define the function to be called within the command-gate context; you must specify a pointer to this function as the first parameter.

Important: Do not call the `runAction` or `runCommand` function from interrupt context.

[Listing 7-8](#) (page 92) illustrates one the use of the `runCommand` function to issue an I/O request.

Listing 7-8 Issuing an I/O request through the command gate

```
void ApplePMU::enqueueCommand ( PMUrequest * request )
{
    commandGate->runCommand(request);
}

void receiveMsg ( OSObject * theDriver, void * newRequest, void *, void *, void *
)
{
    ApplePMU * PMUdriver = (ApplePMU *) theDriver;
    PMUrequest * theRequest = (PMUrequest*)newRequest;
```

```
// Inserts the request in the queue:
theRequest->prev = PMUdriver->queueTail;
theRequest->next = NULL;
if ( PMUdriver->queueTail != NULL ) {
    PMUdriver->queueTail->next = theRequest;
}
else {
    PMUdriver->queueHead = theRequest;
}
PMUdriver->queueTail = theRequest;

// If we can, we process the next request in the queue:
if ( (PMUdriver->PGE_ISR_state == kPMUidle) && !PMUdriver->adb_reading) {
    PMUdriver->CheckRequestQueue();
}
}
```

In this example, the `runCommand` function is used to indirectly invoke the command gate's `Action` function, `receiveMsg`. One important tactic that this example shows is how to defer processing I/O requests—when allocation of memory and other resources might be necessary—until no more events are queued at the command gate. The `receiveMsg` function queues up each incoming request and, if no more requests are pending, calls a `CheckRequestQueue` function to do the actual I/O work.

A typical procedure is to set a timeout (using an `IOTimerEventSource` object) at the same time you issue an I/O request. If the I/O request does not complete within a reasonable period, the timer is triggered, giving you the opportunity to correct any problem (if possible) and reissue the I/O request. If the I/O request is successful, remember to disable the timer. See [Handling Timer Events](#) (page 89) for details on using `IOTimerEventSource` objects.

You destroy an command-gate event source in a driver's deactivation function (usually `stop`). Before you release the `IOCommandGate` object, you should remove it from the work loop. [Listing 7-9](#) (page 93) gives an example of how to do this.

Listing 7-9 Disposing of an `IOCommandGate` object

```
if (commandGate) {
    myWorkLoop->removeEventSource(commandGate);
}
```

```
    commandGate->release();  
    commandGate = 0;  
}
```

Managing Data

A driver's essential work is to shuffle data in and out of the system in response to client requests as well as events such as hardware-generated interrupts. The I/O Kit defines standard mechanisms for drivers to do this using a handful of classes. Drivers use the `IOMemoryDescriptor` and `IOMemoryCursor` classes (and, in OS X v10.4.7 and later, the `IODMACommand` class) to perform pre-I/O and post-I/O processing on data buffers, such as translating them between client formats and hardware-specific formats. This chapter discusses these classes and various issues related to device drivers and data management.

Drivers handle client requests and other events using the `IOWorkLoop` and `IOEventSource` classes to serialize access and thereby protect their critical data. Because of these mechanisms, drivers rarely have to worry about such issues as protecting critical data or disabling and enabling interrupts during the normal course of handling a request. See the chapter [Handling Events](#) (page 77) for information.

Handling I/O Transfers

An I/O transfer is little more than a movement of data between one or more buffers in system memory and a device. “System memory” in this context, however, refers to actual physical memory, not the virtual memory address space used by both user and kernel tasks in OS X. Because I/O transfers at the level of device drivers are sensitive to the limitations of hardware, and hardware can “see” only physical memory, they require special treatment.

Input/output operations in OS X occur within the context of a virtual memory system with preemptive scheduling of tasks and threads. In this context, a data buffer with a stable virtual address can be located anywhere in physical memory, and that physical memory location can change as virtual memory is paged in and out. It's possible for that data buffer to not be in physical memory at all at any given time. Even kernel memory, which isn't subject to relocation, is accessed by the CPU in a virtual address space.

For a write operation, where data in the system is being sent out, a data buffer must be paged in if necessary from the virtual memory store. For a read operation, where a buffer will be filled by data brought into the system, the existing contents of the data buffer are irrelevant, so no page-in is necessary; a new page is simply allocated in physical memory, with the previous contents of that page being overwritten (after being paged out, if necessary).

Yet a requirement for I/O transfers is that the data in system memory not be relocated during the duration of the transfer. To guarantee that a device can access the data in buffers, the buffers must be resident in physical memory and must be wired down so that they don't get paged out or relocated. Then the physical addresses of the buffers must be made available to the device. After the device is finished with the buffers, they must be unwired so that they can once again be paged out by the virtual memory system.

To help you deal with these and other constraints of hardware, the I/O Kit puts several classes at your disposal.

Memory Descriptors and Memory Cursors

In a preemptive multitasking operating system with built-in virtual memory, I/O transfers require special preparation and post-completion processing:

- The space needed for an I/O transfer must reside in physical memory and must be wired down so it can't be paged out until the transfer completes.
- The virtual memory addresses used by software must be converted to physical addresses, and the buffer addresses and lengths must be collected into the scatter/gather lists that describe the data to be transferred.
- After the transfer completes, the memory must be unwired so it can be paged out.

In the I/O Kit, all of this work is performed by objects of the `IOMemoryDescriptor` and `IOMemoryCursor` classes (see [Supporting DMA on 64-Bit System Architectures](#) (page 101) for information on the `IODMACommand` class, which supersedes `IOMemoryCursor` in OS X v10.4.7 and later). An I/O request typically includes an `IOMemoryDescriptor` object that describes the areas of memory involved in the transfer. Initially, the description takes the form of an array of structures, each consisting of a client task identifier (`task_t`), an offset into the client's virtual address space, and a length in bytes. A driver uses the memory descriptor to prepare the memory pages—paging them into physical memory, if necessary, and wiring them down—by invoking the descriptor's `prepare` function.

When the memory is prepared, a driver at a lower level of the stack—typically a driver that controls a DMA (Direct Memory Access) engine—then uses a memory-cursor object to get the memory descriptor's buffer segments and with them generate a scatter/gather list suitable for use with the hardware. It does this by invoking the `getPhysicalSegments` function of the memory cursor and doing any necessary processing on the segments it receives. When the I/O transfer is complete, the driver originating the I/O request invokes the memory descriptor's `complete` function to unwire the memory and update the virtual-memory state. When all this is done, it informs the client of the completed request.

Beginning in OS X v10.2, `IOBufferMemoryDescriptor` (a subclass of `IOMemoryDescriptor`) allows a buffer to be allocated in any task for I/O or sharing through mapping. In previous versions of OS X, an `IOBufferMemoryDescriptor` object could only represent a buffer allocated in the kernel's address space. In OS X v10.2 and later, however, the changes to the `IOBufferMemoryDescriptor` API support a better way to handle

I/O generated at the behest of a nonkernel client. Apple recommends that such I/O be sent to buffers the kernel allocates in the client's address space using the `IOBufferMemoryDescriptor` API. This gives control of the allocation method to the kernel, ensuring that the buffer is allocated in accordance with internal guidelines dictated by the virtual-memory system. The user task can still specify the kind of buffer to be allocated, such as pageable or sharable. And, the user task can still access the buffer using the `vm_address_t` and `vm_size_t` variables it receives from the user client. For software running in OS X v10.2 and later, Apple recommends that:

- User tasks no longer use `malloc` or other user-level library functions to allocate I/O buffers in their own address space.
- User clients use `IOBufferMemoryDescriptor` objects to represent kernel-allocated buffers instead of the `IOMemoryDescriptor` objects that represented user task-allocated buffers.

Network drivers are the exception to the use of `IOMemoryDescriptor` objects. The Network family instead uses the `mbuf` structure defined by the BSD kernel networking stacks. BSD `mbuf` structures are already optimized for handling network packets, and translating between them and memory descriptors would merely introduce unnecessary overhead. The network family defines subclasses of `IOMemoryCursor` for extracting scatter/gather lists directly from `mbuf` structures, essentially making this aspect of I/O handling the same for network drivers as for any other kind of driver.

`IOMemoryDescriptor` and `IOMemoryCursor` are capable of handling the management and reuse of memory buffers for most drivers. Drivers with special requirements, such as that all buffer memory be contiguous or located within a particular area in physical memory, must perform the extra work necessary to meet these requirements, but still use memory descriptors and cursors for their interaction with the I/O Kit.

Memory in an I/O Request

Although the preceding section discusses how I/O Kit drivers transfer data between device and system memory by using objects of the `IOMemoryDescriptor` and `IOMemoryCursor` classes, it does so at a fairly general level. It's also instructive to consider an I/O request at a more detailed level: What happens from the moment a user process makes a call to write or read data to the instant the data is transferred to or from a device? What are the exact roles played by `IOMemoryDescriptors`, `IOMemoryCursors`, and other I/O Kit objects in this chain of events?

All read and write operations in user space—that is, made by applications or other non-kernel processes—are based ultimately on I/O vectors. An I/O vector is an array of structures, each of which gives the address and length of a contiguous chunk of memory of a particular process; this memory is expressed in the virtual address space of the process. An I/O vector is sometimes known as a scatter/gather list.

To start with a more specific example, a process makes a call to write some data to a device. It must pass in a minimum set of parameters: a handle to the target of the call (`io_service_t`), a command ("write"), the base address of the I/O vector array, and the number of array elements. These parameters get passed down to the

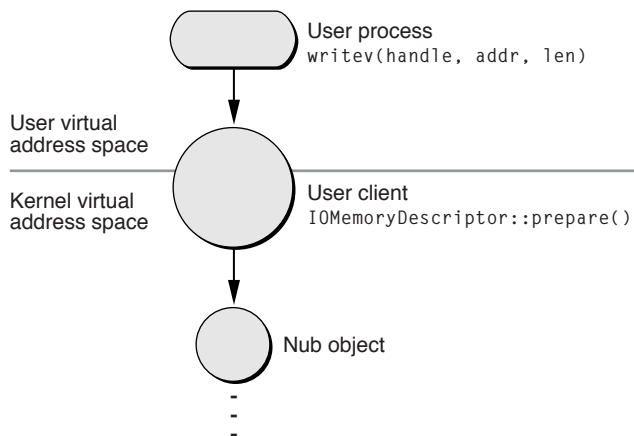
kernel. But how is the kernel to make sense of them, particularly the base address of the array, which is typed as `void *`. The kernel lives in its own virtual address space, and the virtual address space of a user process, untranslated, means nothing to it.

Before going further, let's review the memory maps maintained by the operating system. In OS X there are three different kinds of address space:

- The virtual address space of individual user processes (such as applications and daemons)
- The virtual address space of the kernel (which includes the I/O Kit)
- Physical address space (system memory)

When a user process issues an I/O call, the call (with its parameters) percolates down to the kernel. There the I/O Kit converts the handle parameter to an object derived from an appropriate `IOUserClient` subclass. The user-client object logically sits astride the boundary separating the kernel and user space. There is one user client per user process per device. From the perspective of the kernel, the user client is a client driver at the top of the stack that communicates with the nub object below it (see [Figure 8-1](#) (page 98)).

Figure 8-1 The role of the user client in an I/O command

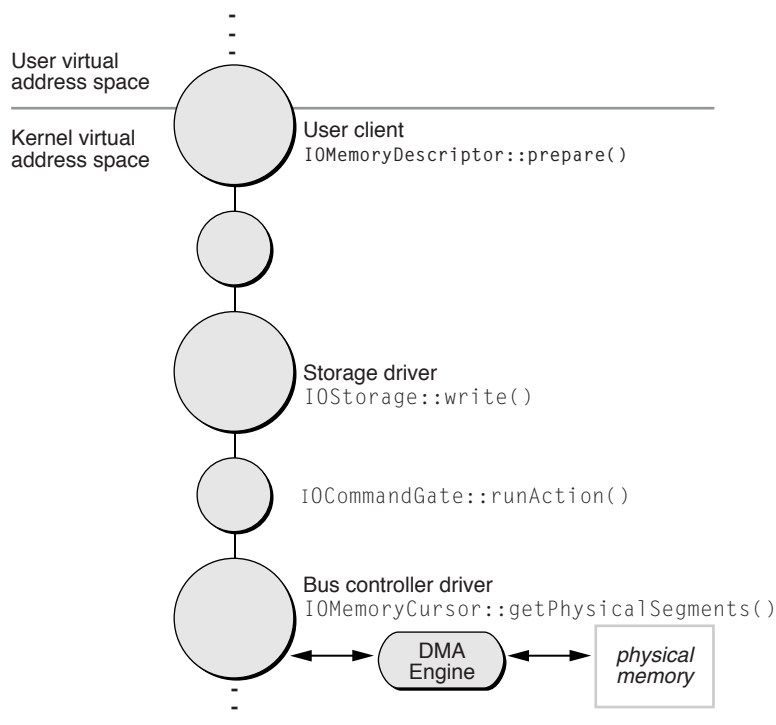


The user client also translates between the address spaces, mapping the user-process buffers into the kernel's virtual address space. To do this, it creates an `IOMemoryDescriptor` object from the other parameters specified in the original call. The particular power of an `IOMemoryDescriptor` object is that it can describe a piece of memory throughout its use in an I/O operation and in each of the system's address spaces. It can be passed between various drivers up and down the driver stack, giving each driver an opportunity to refine or manipulate the command. A driver can also reuse a memory descriptor; in some circumstances, it is better—if continuous allocation is going to involve a performance hit—to create a pool of `IOMemoryDescriptors` ahead of time and recycle them.

After the user client creates (or reuses) the `IOMemoryDescriptor` object, it immediately invokes the memory descriptor's `prepare` member function. The `prepare` call first makes sure physical memory is available for the I/O transfer. Because this is a write operation, the virtual memory system may have to page in the data from its store. If it were a read operation, the virtual-memory (VM) pager might have to swap out some other pages from physical memory to make room for the I/O transfer. In either case, when the VM pager is involved in preparing physical memory, there are implications for bus controller drivers that program DMA engines or otherwise must deal directly with the requirements of devices (see [DMA and System Memory](#) (page 105)). After sufficient physical memory is secured for the I/O transfer, the `prepare` function wires the memory down so it cannot be paged out.

The `prepare` call must be made before the I/O command crosses formally into the I/O Kit, and it should also be made on the requesting client's thread and before taking any locks. Never invoke `prepare` within the command-gate context because `prepare` is a synchronous call and can *block* (that is, wait for some other process to complete) indefinitely. When `prepare` returns, the user client (typically) initiates the call to schedule the I/O request, passing the memory descriptor to a driver down the stack (through established interfaces, as described in [Relaying I/O Requests](#) (page 103)). The memory descriptor might be passed along to other drivers, each of which might manipulate its contents, until it eventually reaches an object that is close to the hardware itself—typically a bus controller driver or that driver's client nub. This object schedules the I/O request and takes the command (see [I/O Requests and Command Gates](#) (page 90) for information on command gates). Within the command gate it usually queues up the request for processing as soon as the hardware is free.

Figure 8-2 The principal I/O Kit objects in an I/O transfer



Ultimately receiving the request at the lower levels of the driver stack is a bus controller driver (such as for ATA or SCSI). This lower-level driver object must program the controller's DMA engine or otherwise create a scatter/gather list to move data directly into and out of the device. Using an `IOMemoryCursor` object, this driver generates physical addresses and lengths from the memory descriptor's buffer segments and does any other work necessary to create a scatter/gather list with the alignment, endian format, and size factors required by the hardware. To generate the segments, it must call the `getPhysicalSegments` function of the memory cursor. This whole procedure is run in the work-loop context. At the completion of the transfer, a hardware interrupt is typically generated to initiate an I/O transfer in the other direction.

When the I/O transfer is complete, the object that called `prepare` invokes the memory descriptor's `complete` function to unwire the memory and update the virtual-memory state. It's important to balance each `prepare` with a corresponding `complete`. When all this is done, the user client informs the originating process of the completed request.

Issues With 64-Bit System Architectures

Beginning with OS X v10.3 and later, Apple introduced some changes to allow existing device drivers to work with the new 64-bit system architectures. The problem to be solved involved the communication between devices on the PCI bus, which can handle 32-bit addresses, and the 64-bit main memory.

Then, in OS X v10.4.7, Apple introduced the `IODMACommand` class to allow device drivers that perform DMA to address 64-bit main memory in Intel-based Macintosh computers. The following sections describe these changes.

Address Translation on 64-Bit System Architectures

Apple solved the problem with address translation which “maps” blocks of memory into the 32-bit address space of a PCI device. In this scheme, the PCI device still sees a 4-gigabyte space, but that space can be made up of noncontiguous blocks of memory. A part of the memory controller called the DART (device address resolution table) translates between the PCI address space and the much larger main memory address space. The DART handles this by keeping a table of translations to use when mapping between the physical addresses the processor sees and the addresses the PCI device sees (called I/O addresses).

The address-translation process is transparent if your driver adheres to documented, Apple-provided APIs. For example, when your driver calls the `IOMemoryDescriptor` method `prepare`, a mapping is automatically placed in the DART. Conversely, when your driver calls `IOMemoryDescriptor`'s `release` method, the mapping is removed. Although this has always been the recommended procedure, failure to do this in a driver running on OS X v10.3 or later may result in random data corruption or panics. Be aware that the `release` method does not take the place of the `complete` method. As always, every invocation of `prepare` should be balanced with an invocation of `complete`.

If your driver experiences difficulty on an OS X v10.3 system, you should first ensure that you are following these guidelines:

- Always call `IOMemoryDescriptor::prepare` to prepare the physical memory for the I/O transfer (this also places a mapping into the DART).
- Balance each `IOMemoryDescriptor::prepare` with an `IOMemoryDescriptor::complete` to unwind the memory.
- Always call `IOMemoryDescriptor::release` to remove the mapping from the DART.
- On hardware that includes a DART, pay attention to the DMA direction for reads and writes. On a 64-bit system, a driver that attempts to write to a memory region whose DMA direction is set up for reading will cause a kernel panic.

One side effect of these changes in the memory subsystem is that OS X is likely to return physically contiguous page ranges in memory regions. In earlier versions of OS X, the system returned multi-page memory regions in reverse order, beginning with the last page and moving towards the first page. Because of this, a multi-page memory region seldom contained a physically contiguous range of pages.

The greatly increased likelihood of seeing physically contiguous blocks of memory in memory regions might expose latent bugs in drivers that did not previously have to handle physically contiguous pages. Be sure to check for this possibility if your driver is behaving incorrectly or panicking.

Another result of the memory-subsystem changes concerns physical addresses a driver might obtain directly from the pmap layer. Because there is not a one-to-one correspondence between physical addresses and I/O addresses, physical addresses obtained from the pmap layer have no purpose outside the virtual memory system itself. Drivers that use pmap calls (such as `pmap_extract`) to get such addresses will fail to work on systems with a DART. To prevent the use of these calls, OS X v10.3 will refuse to load a kernel extension that uses them, even on systems without a DART.

Supporting DMA on 64-Bit System Architectures

As described in [Address Translation on 64-Bit System Architectures](#) drivers running in OS X v10.3 and later that use documented, Apple-provided APIs experience few (if any) problems when addressing physical memory on 64-bit systems, because the address translation performed by the DART is transparent to them. However, current Intel-based Macintosh computers do not include a DART and this has ramifications for device drivers that need to use physical addresses, such as those that perform DMA.

Because there is no hardware-supported address translation, or remapping, performed in current Intel-based Macintosh computers, device drivers that need to access physical memory must be able to address memory above 4 gigabytes. In OS X v10.4.7 and above, you can use the `IODMACommand` class to do this.

Important: A device driver running in OS X v10.4.7 and later that targets Intel-based Macintosh computers *must* be updated to use `IODMACommand` if it uses physical addresses. A device driver running in OS X v10.4.7 and later that targets PowerPC-based Macintosh computers is not required to do so, but it's recommended, especially if the driver uses physical addresses.

The `IODMACommand` class supersedes the `IOMemoryCursor` class: it provides all the functionality of `IOMemoryCursor` and adds a way for you to specify your hardware's addressing capability and functions to copy memory to a bounce buffer when necessary. When you instantiate an `IODMACommand` object, you can specify the following attributes:

- The number of address bits your hardware can support (for example, 32, 40, or 64)
- The maximum segment size
- Any alignment restrictions required by your hardware
- The maximum I/O transfer size
- The format of the physical address segments `IODMACommand` returns (for example, 32-bit or 64-bit and big-endian, little-endian, or host-endian)

In the typical case, you use an `IODMACommand` object in the following way:

1. Create an `IODMACommand` object per I/O transaction (you can create a pool of `IODMACommand` objects when your driver starts).
2. When an I/O request arrives, use `IODMACommand::setMemoryDescriptor` to target the `IOMemoryDescriptor` object representing the request.
3. Call `IODMACommand::prepare` (among other things, this function allocates the mapping resources that may be required for the transfer).
4. Use `IODMACommand` functions to generate the appropriate physical addresses and lengths (`IODMACommand::gen64IOVMSegments` returns 64-bit addresses and lengths and `IODMACommand::gen32IOVMSegments` returns 32-bit addresses and lengths).
5. Start the hardware I/O.
6. When the I/O is finished, call `IODMACommand::complete` (to complete the processing of DMA mappings), followed by `IODMACommand::clearMemoryDescriptor` (to copy data from the bounce buffer, if necessary, and release resources).

Note: The `IODMACommand` prepare and complete functions are distinct from the `IOMemoryDescriptor` prepare and complete functions. The `IODMACommand` prepare and complete functions bracket the start and end of the DMA transaction, whereas the `IOMemoryDescriptor` prepare and complete functions wire and unwire the memory, and must be called as usual.

If your DMA engine does complicated things, such as performing partial I/Os or synchronizing multiple accesses to a single `IOMemoryDescriptor`, you should write your driver assuming that the memory will be bounced. You don't need to add code that checks for bouncing, because `IODMACommand` functions, such as `synchronize`, are no-ops when they are unnecessary.

Relaying I/O Requests

Client requests are delivered to drivers through specific functions defined by the driver's family. A storage family driver, for example, handles read requests by implementing the function `read`. Similarly, a network family driver handles requests for transmitting network packets by implementing the function `outputPacket`.

An I/O request always includes a buffer containing data to write or providing space for data to be read. In the I/O Kit this buffer takes the form of an `IOMemoryDescriptor` object for all families except networking, which uses the `mbuf` structure defined by BSD for network packets. These two buffer mechanisms provide all drivers with optimized management of data buffers up and down the driver stacks, minimizing the copying of data and performing all the steps required to prepare and complete buffers for I/O operations.

An I/O Kit family defines I/O and other request interfaces. You typically don't have to worry about protection of resources in a reentrant context for your driver, unless it specifically forgoes the protection offered by the family or the I/O Kit generally.

More on Memory Descriptors

A memory descriptor is an object inheriting from the `IOMemoryDescriptor` class that describes how a stream of data, depending on direction, should either be laid into memory or extracted from memory. It represents a segment of memory holding the data involved in an I/O transfer and is specified as one or more physical or virtual address ranges (a range being a starting address and a length in bytes).

An `IOMemoryDescriptor` object permits objects at various levels of a driver stack to refer to the same piece of data as mapped into physical memory, the kernel's virtual address space, or the virtual address space of a user process. The memory descriptor provides functions that can translate between the various address spaces. In a sense, it encapsulates the various mappings throughout the life of some piece of data involved in an I/O transfer.

`IOMemoryDescriptor` is an abstract base class defining common methods for describing physical or virtual memory. Although it is an abstract class, the I/O Kit provides a concrete general-purpose implementation of `IOMemoryDescriptor` for objects that are directly instantiated from the class. The I/O Kit also provides two specialized public subclasses of `IOMemoryDescriptor`: `IOMultiMemoryDescriptor` and `IOSubMemoryDescriptor`. [Table 8-1](#) (page 104) describes these classes.

Table 8-1 Subclasses of `IOMemoryDescriptor`

Class	Description
<code>IOMultiMemoryDescriptor</code>	Wraps multiple general-purpose memory descriptors into a single memory descriptor. This is commonly done to conform to a bus protocol.
<code>IOSubMemoryDescriptor</code>	Represents a memory area which comes from a specific subrange of some other <code>IOMemoryDescriptor</code> .

The `IOMemoryDescriptor` base class itself defines several methods that can be usefully invoked from objects of all subclasses. Some methods return the descriptor's physically contiguous memory segments (for use with an `IOMemoryCursor` object) and other methods map the memory into any address space with caching and placed-mapping options.

A related and generally useful class is `IOMemoryMap`. When you invoke `IOMemoryDescriptor`'s `map` method to map a memory descriptor in a particular address space, an `IOMemoryMap` object is returned. Objects of the `IOMemoryMap` class represent a mapped range of memory as described by an `IOMemoryDescriptor`. The mapping may be in the kernel or a non-kernel task, or it may be in physical memory. The mapping can have various attributes, including processor cache mode.

More on Memory Cursors

An `IOMemoryCursor` lays out the buffer ranges in an `IOMemoryDescriptor` object in physical memory. By properly initializing a memory cursor and then invoking that object's `getPhysicalSegments` function on an `IOMemoryDescriptor`, a driver can build a scatter/gather list suitable for a particular device or DMA engine. The generation of the scatter/gather list can be made to satisfy the requirements of segment length, transfer length, endian format, and alignment imposed by the hardware.

A controller driver for a bus such as USB, ATA, or FireWire is typically the object that uses `IOMemoryCursors`. Such drivers should create an `IOMemoryCursor` and configure the memory cursor to the limitations of the driver's DMA hardware or (if PIO is being used instead) the limitations of the device itself. For instance, the memory cursor used for the FireWire SBP-2 protocol should be configured to a maximum physical segment size of 65535 and an unlimited transfer size.

You can configure an `IOMemoryCursor` in a variety of ways. The most obvious way is to supply the initialization parameters: maximum segment size, maximum transfer length, and a pointer to a segment function. This callback function, typed `SegmentFunction`, writes out a single physical segment to an element in a vector array defining the scatter/gather list. Your driver can also perform post-processing on the extracted segments, swapping bytes or otherwise manipulating the contents of the segments. Finally, you can create a subclass of `IOMemoryCursor` or use one of the subclasses provided by Apple. See [IOMemoryCursor Subclasses](#) (page 107) for more on this topic.

DMA and System Memory

Writers of bus controller drivers have two basic considerations when they effect I/O transfers. They are receiving data laid out in memory in a particular way and they must send that data to a destination that might expect a radically different memory layout. Depending on direction, the source and destination can be either a DMA engine (or a specific device) or the client memory represented by an `IOMemoryDescriptor`. Additionally, memory coming from the system might be conditioned by the Unified Buffer Cache (UBC) and this has implications for driver writers. This section discusses some of these factors.

Direct Memory Access (DMA) is a built-in capability of certain bus controllers for directly transferring data between a device attached to the bus and system memory—that is, the physical memory on the computer's motherboard. DMA enhances system performance by freeing the microprocessor from having to do the transfer of data itself. The microprocessor can work on other tasks while the DMA engine takes care of moving data in and out of the system.

Each bus on an OS X system has its own DMA engine and they all are different. The PCI bus controller on OS X uses bus master DMA, a type of DMA wherein the controller controls all I/O operations on behalf of the microprocessor. Other bus controllers (ATA, SCSI, or USB, for example) implement different DMA engines. Each engine can have its own alignment, endian format, and size restrictions.

An alternative to DMA is the Programmed Input/Output (PIO) interface, usually found on older or under-designed hardware. In PIO, all data transmitted between devices and system memory goes through the microprocessor. PIO is slower than DMA because it consumes more bus cycles to accomplish the same transfer of data.

OS X supports both bus master DMA and PIO for moving data in and out of the system. In fact, some drivers could conceivably make use of both models for the same I/O transfer—for example, processing most bytes using DMA and then processing the last few bytes using PIO, or using PIO to handle error conditions.

The Unified Buffer Cache (UBC) is a kernel optimization that combines the file-system cache and the virtual-memory (VM) cache. The UBC eliminates the situation where the same page is duplicated in both caches. Instead, there is just one image in memory and pointers to it from both the file system and the VM system. The underlying structure of the UBC is the Universal Page List (UPL). The VM pager deals with memory as defined by UPLs; when an I/O request is based on paged-in memory, it is called a conforming request. A non-conforming request is one that isn't UPL-defined. A UPL segment of memory has certain characteristics; it:

- Is page sized (4 kilobytes)
- Is page aligned
- Has a maximum segment size of 128 kilobytes
- Is already mapped into the kernel's virtual address space

The I/O Kit has adapted its APIs to the UPL model because it's more efficient for I/O transfers; in addition, it makes it easier to batch I/O requests. An `IOMemoryDescriptor` object might be backed—entirely or partially—by UPL-defined memory. If the object is backed by a UPL, then there cannot be more than a prearranged number of physical segments. The bus controller driver that extracts the segments (using an `IOMemoryCursor`) must allocate sufficient resources to issue the I/O request associated with the memory descriptor.

Dealing With Hardware Constraints

Apple's policy on how drivers should deal with hardware constraints is generous toward clients of DMA controller drivers and places certain expectations on the drivers themselves.

- Clients of DMA controller drivers have no alignment restrictions placed on them. Generally, UPL-defined data (page-aligned and page-sized) should be optimal, but it is not required.
- If the memory is UPL-defined, then in order to avoid deadlocking the driver should not allocate memory. Controller drivers must be able to process a UPL without allocating any buffers, or they must preallocate sufficient resources prior to any particular UPL-based I/O transfer. Operations that don't meet the pager constraints (that is, that aren't UPL-based) can allocate buffers.
- A controller driver must do its best to honor any request that it receives. If necessary, it should even create a buffer (or even use a static buffer), or preallocate resources and copy the data. (Remember that allocation during an I/O request can cause a deadlock.) If the driver cannot execute the I/O request, it should return an appropriate error result code.

In summary, driver writers must be prepared to handle a request for any alignment, size, or other restriction. They should first attempt to process the request as conforming to UPL specifications; if the assumption of UPL proves true, they should never allocate memory because doing so could lead to deadlock. If the request is non-conforming, the driver should (and can) do whatever it has to do to satisfy the request, including allocating resources.

IOMemoryCursor Subclasses

Apple provides several subclasses of `IOMemoryCursor` for different situations. If your DMA engine requires a certain endian data format for its physical segments, your driver can use the subclasses that deal with big-endian and little-endian data formats (and thus will not have to perform this translation when it builds the scatter/gather lists for the DMA engine). Another subclass enables your driver to lay out data in the byte orientation expected by the system's processor. [Table 8-2](#) (page 107) describes these subclasses.

Table 8-2 Apple-provided subclasses of `IOMemoryCursor`

Subclass	Description
<code>IONaturalMemoryCursor</code>	Extracts and lays out a scatter/gather list of physical segments in the natural byte order for the given CPU.
<code>IOBigMemoryCursor</code>	Extracts and lays out a scatter/gather list of physical segments encoded in big-endian byte format. Use this memory cursor when the DMA engine requires a big-endian address and length for each segment.
<code>IOLittleMemoryCursor</code>	Extracts and lays out a scatter/gather list of physical segments encoded in little-endian byte format. Use this memory cursor when the DMA engine requires a little-endian address and length for each segment.

Of course, you can create your own subclass of the virtual `IOMemoryCursor` or of one of its subclasses to have your memory cursor accomplish exactly what you need it to do. But in many cases, you may not have to create a subclass to get the behavior you're looking for. Using one of the provided memory-cursor classes, you can implement your own `outputSegment` callback function (which must conform to the `SegmentFunction` prototype). This function is called by the memory cursor to write out a physical segment in the scatter/gather list being prepared for the DMA engine. In your implementation of this function, you can satisfy any special layouts required by the hardware, such as alignment boundaries.

Managing Power

The power-management functionality of the I/O Kit aims to minimize the power consumed by a computer system, behavior that is especially important for portable computers where battery life is a crucial feature. Power management also imposes an orderly sequence of actions, such as saving and restoring state, when a system (or a part of it) sleeps or wakes.

This chapter focuses on power management for in-kernel drivers that manage hardware. Read this chapter to learn about power management in OS X and to find out what level of power-management support you need to provide and how to implement it. Although power management is a complex technology, the majority of in-kernel drivers need to implement only the most basic functionality to participate successfully in OS X power management.

Note: If you're developing an application that accesses hardware, such as a user-space driver for a digital camera, scanner, webcam, or tape drive, you probably do not need to perform any power-management tasks. For more information on developing applications that behave as user-space drivers, including information on how to set up your application to receive power-event notifications, see *Accessing Hardware From Applications*.

The precise set of power-management responsibilities your driver must fulfill depends on factors such as how much support your driver's superclass provides, whether your device receives power from a system bus (such as PCI), and to what power events your driver needs to respond.

If you're unfamiliar with power management in OS X, you should begin by reading the following three sections:

- [Power Events](#) (page 109) which explains what power events are and how they affect your device
- [The Power Plane: A Hierarchy of Power Dependencies](#) (page 109) which describes how OS X monitors the power relationships among devices, drivers, and other objects
- [Devices and Power States](#) (page 111) which defines devices and power states in power-management terms

Then, all driver developers should read [Deciding How to Implement Power Management in Your Driver](#) (page 112) to find out what to do next. After you decide what type of power management you need to implement, read [Implementing Basic Power Management](#) (page 113) and, if appropriate, [Implementing Advanced Power Management](#) (page 117)

Power Events

Before you consider how to implement power management in your driver, you need to understand what power events are and how they can affect your device. In OS X, power events are transitions to and from the following states:

- Sleep
- Wake
- Shutdown or restart

All drivers must respond to sleep events. OS X defines different types of sleep, which can occur for different reasons. For example, **system sleep** occurs when the user chooses Sleep from the Apple menu or closes the lid of a laptop; **idle sleep** occurs when there has been no device or system activity during the interval the user selects in the Energy Saver preferences. To your driver, however, all sleep events appear identical. The important thing to understand about a sleep event is that your device may be powered off when the system sleeps, so your driver must be prepared to initialize the device when it is awakened.

All drivers must respond to a system wake event by powering on. Wake can occur when the user hits a key on the keyboard, presses the power button, or when the computer receives a network administrator wake-up packet. On wake, drivers should perform the appropriate restoration of device state.

Device drivers do not have to respond to shutdown and restart events. A driver can choose to get notification of an impending shutdown or restart using the technique described in [Receiving Shutdown and Restart Notifications](#) (page 127) but it's important to understand that no driver can prevent a shutdown event.

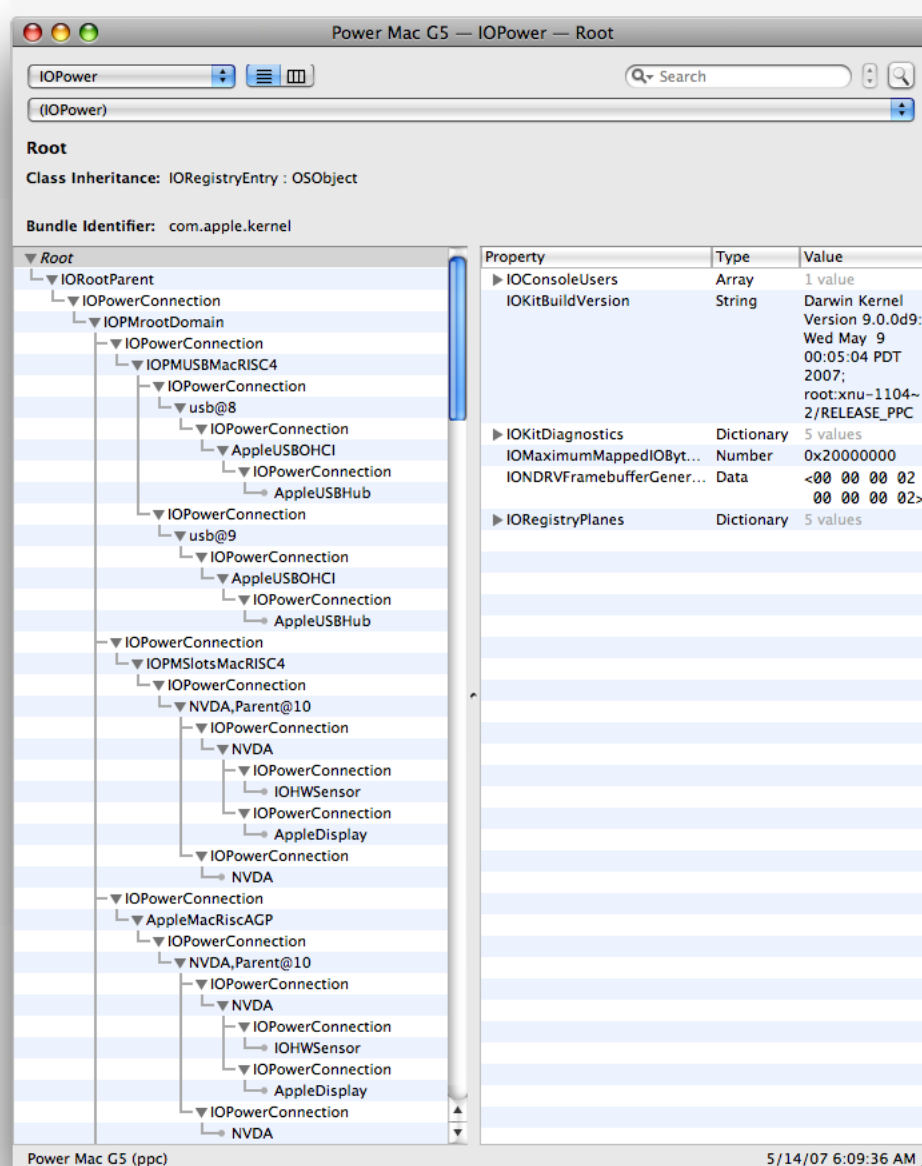
Another type of event is a device power-up request, which occurs when some object in the system requires an idle or powered-off device to be in a usable state. A device power-up request notification uses most of the same mechanisms as sleep and wake notifications. Although most drivers do not need to know about device power-up requests, some drivers might need to implement them and even make such requests themselves. For more information about this, see [Initiating a Power-State Change](#) (page 121)

The Power Plane: A Hierarchy of Power Dependencies

OS X tracks all power-managed devices in a tree-like structure, called the power plane, that captures the power dependencies among devices. A device, usually a leaf object in the power plane, generally receives power from its ancestors and may provide power to its children. For example, because a PCI card depends for power on the PCI bus to which it's attached, the PCI card is considered to be a **power child** of the PCI bus. Likewise, the PCI bus is considered to be the **power parent** of the devices attached to it.

The power plane is one of the planes of the I/O Registry. As described in [The I/O Registry](#) (page 39) the I/O Registry is a dynamic database of device and driver objects that expresses the various provider-client relationships among them. To view the power plane in a running system, open the I/O Registry Explorer application (located in `/Developer/Applications/Utilities`) and choose `IOPower` from the pop-up menu. You can also enter `ioreg -p IOPower` at the command line to see a representation of the current power plane. Figure 9-1 shows the power plane in a Power Mac G5 running OS X v10.5.

Figure 9-1 The power plane shown in I/O Registry Explorer



In Figure 9-1 you can see the root of the power plane, an object called `IOPMrootDomain`, and objects that represent devices and drivers. You can ignore the many `IOPowerConnection` objects, which represent power connections, because these objects are of interest only to internal power-management objects and processes.

Devices and Power States

The fundamental entity in power management is the device. From a power-management perspective, a device is a unit of hardware whose power consumption can be measured and controlled independently of system power. A device can also have some state that needs to be saved and restored across changes in power. In power-management terms, “device” is synonymous with the device-driver object that controls it.

A device must have at least two power states associated with it—off and on. A device may also have intermediate states that represent some level of power between full power and no power. These states are described in a power-state array you create in your driver. (You learn how to create this array and provide power-state information in step 3 in [Implementing Basic Power Management](#) (page 113)) The power-management functionality of the I/O Kit uses these states to ensure that all drivers in the power plane receive the power they require. Each power state is defined by the device’s capabilities when in that state:

- A device that is on uses maximum power and has complete functionality.
- A device that is off uses no power and has no functionality.
- A device can be in a reduced-power state in which it is still usable, but at a lower level of performance or functionality.
- A device can be in an intermediate state in which it is not usable, but retains some configuration or state.

The power-management functionality of the I/O Kit associates several attributes with each power state of a device. A device driver must set these attributes to ensure that accurate information about the device’s capabilities and requirements is available.

The power-state attributes provide the following information:

- The capability of the device while in a given state
- The device’s power requirements of its power parent
- The power characteristics the device can provide to its power children
- The version of the power-state structure the device uses to store its power-state information

Deciding How to Implement Power Management in Your Driver

To participate in OS X power management, most in-kernel drivers need only ensure that their devices respond appropriately to system sleep and wake events. Some in-kernel drivers might need to perform other tasks, such as implementing an idle state or taking action at system shutdown, but these drivers are not typical. Reflecting this distinction, OS X power management defines two types of drivers:

- A **passive** driver implements basic power management to respond to system power events; it does not initiate any power-related actions for its device.
- An **active** driver implements basic power management to respond to system power events, but it also implements advanced power management to perform tasks such as deciding when the device should become idle, changing the device's power state, or processing prior to system shutdown.

An example of a passive driver is the `AppleSmartBatteryManager` driver present in most Macintosh laptop computers. The `AppleSmartBatteryManager` driver provides battery-status information to the battery-status menu bar item; when the system is about to sleep, the driver simply stops polling the battery for status information. A good example of an active driver is the built-in audio chip driver, because it performs its own idleness determination to allow the audio hardware to power off when it is not in use. If there is no sound coming out of a laptop's or desktop's internal speakers, the audio hardware will drop into a low power mode until it is needed.

As you can imagine, a passive driver is much easier than an active driver to design and implement. Essentially, a passive driver implements one virtual method and makes between three and five calls to participate in power management. The responsibilities of an active driver, on the other hand, begin with those of a passive driver, but increase with each additional task the driver needs to perform.

Some I/O Kit families provide various levels of built-in power-management support to driver subclasses. For example, the Network family (`IONetworkingFamily`) performs some of the power-management initialization tasks for a subclass driver, leaving the driver to perform other device-specific power-management tasks.

Before you begin designing your driver's power-management implementation, you should look up your I/O Kit family in [I/O Kit Family Reference](#) (page 142) to find out if the family provides any power-management support or requires subclasses to perform different or additional tasks. Be aware, however, that any I/O Kit family that provides power-management functionality may still require you to implement some parts of it. The following I/O Kit families provide some type of power-management functionality:

- Audio family (described in [Audio](#) (page 145))
- FireWire family (described in [FireWire](#) (page 146))
- Network family (described in [Network](#) (page 152))
- PC card family, which includes Express Card devices (described in [PC Card](#) (page 156))

- PCI family (described in [PCI and AGP](#) (page 157))
- SCSI Architecture Model family (described in [SCSI Architecture Model](#) (page 161))
- USB family (described in [USB](#) (page 172))

Even if your driver is a subclass of an I/O Kit family that does not provide any power-management support, or if your driver is a direct subclass of `IOService`, it can still be a passive power-management participant as long as it only responds to system-initiated power events. If, on the other hand, your driver needs to determine when your device is idle or perform pre-shutdown tasks, you must implement advanced power management.

If you decide to develop a passive driver, you should read [Implementing Basic Power Management](#) (page 113) to learn how to participate in power management and respond to sleep and wake events. You do not need to read any other sections in this chapter.

If your driver needs to be an active power manager, you should also read [Implementing Basic Power Management](#) (page 113) Then you should read [Implementing Advanced Power Management](#) (page 117) for guidance on implementing specific tasks.

Implementing Basic Power Management

As defined in [Deciding How to Implement Power Management in Your Driver](#) (page 112) a passive driver only responds to sleep and wake events; it does not initiate any power state-changing activity. Your passive driver must do the following things to handle sleep and wake:

- Get attached into the power plane so you receive power-change notifications and to ensure that your device's power dependencies are considered when it is told to sleep and wake.

Power dependencies affect the ordering of sleep and wake notifications. Specifically, your driver is told to sleep before its power parent is told to sleep, and your driver is told to wake after its power parent is told to wake.

Note: It's possible for a device to have more than one power parent, but it's important to understand that, in this case, a specific ordering of power changes is not guaranteed. Specifically, your device is awakened after the first power parent wakes up, not after all power parents wake up.

- Save hardware state to memory before system sleep and restore state during wake.
You are responsible for writing code to do this.
- Prevent all hardware accesses while your device is preparing for sleep.

You can return an error to any I/O request you receive while your device is going to sleep or you can block all incoming threads using a gating mechanism, such as `IOWorkGate`, on your work loop (see [Work Loops](#) (page 77) to learn more about work loops).

To participate in power management so that you receive notifications of power events, ensure your driver is correctly attached into the power plane, and handle power-state changes, you make a few calls and implement one virtual method. The `IOService` class provides all the methods described in this section. Follow the steps listed below to implement basic power management in your driver.

1. Initialize power management using `PMInit`. The `()` method allocates internal power-management data structures that allow internal processes to track your driver.

In your driver's start routine, after the call to your superclass's start method, make the following call:

```
PMInit();
```

2. Get attached into the power plane using `joinPMtree`. The `(IOService*)` method attaches the passed-in driver object into the power plane as a child of its provider.

In your driver's start routine, after the call to `PMInit` and before the call to `registerPowerDriver` (shown in step 3), call `joinPMtree` as shown below:

```
provider->joinPMtree(this);
```

3. Provide information about your device's power states and register your driver with power management.
 - a. First, declare an array of two structures to contain information about your device's off and on states. The first element in the array must contain the structure that describes the off state and the second element of the array must contain the structure that describes the on state. Typically, a driver switches its device to the off state in response to a sleep event and to the on state in response to a wake event, as described in [Power Events](#) (page 109)

In your driver's start routine, after the call to `joinPMtree`, fill in two `IOPMPowerState` structures, as shown below:

```
// Declare an array of two IOPMPowerState structures (kMyNumberOfStates
= 2).
static IOPMPowerState myPowerStates[kMyNumberOfStates];
// Zero-fill the structures.
bzero (myPowerStates, sizeof(myPowerStates));
// Fill in the information about your device's off state:
```

```
myPowerStates[0].version = 1;
myPowerStates[0].capabilityFlags = kIOPMPowerOff;
myPowerStates[0].outputPowerCharacter = kIOPMPowerOff;
myPowerStates[0].inputPowerRequirement = kIOPMPowerOff;
// Fill in the information about your device's on state:
myPowerStates[1].version = 1;
myPowerStates[1].capabilityFlags = kIOPMPowerOn;
myPowerStates[1].outputPowerCharacter = kIOPMPowerOn;
myPowerStates[1].inputPowerRequirement = kIOPMPowerOn;
```

In some drivers, you might see this step implemented in code similar to the following:

```
static IOPMPowerState myPowerStates[kMyNumberOfStates] = {
    {1, kIOPMPowerOff, kIOPMPowerOff, kIOPMPowerOff, 0, 0, 0, 0, 0, 0,
    0, 0},
    {1, kIOPMPowerOn, kIOPMPowerOn, kIOPMPowerOn, 0, 0, 0, 0, 0, 0,
    0},
};
```

- b. Then, still in your driver's start routine, register your driver with power management using (`IOService*`, `IOPMPowerState*`, `unsigned long`). The `registerPowerDriver` method tells power management that the passed-in driver object can transition the device between the power states described in the passed-in array. After you fill in the `IOPMPowerState` structures, call `registerPowerDriver` with your power-state array as shown below:

```
registerPowerDriver (this, myPowerStates, kMyNumberOfStates);
```

4. Handle power-state changes using `setPowerState`. While your driver is running, you perform tasks that handle sleep and wake event notifications in your implementation of the virtual `IOService` method `setPowerState`. An example of how to do this is shown below:

```
IOReturn MyIOServiceDriver::setPowerState ( unsigned long whichState,
IOService * whatDevice )
// Note that it is safe to ignore the whatDevice parameter.
{
    if ( 0 == whichState ) {
        // Going to sleep. Perform state-saving tasks here.
```

```
    } else {  
        // Waking up. Perform device initialization here.  
    }  
    if ( done )  
        return kIOPMAckImplied;  
    else  
        return (/* a number of microseconds that represents the maximum time  
        required to prepare for the state change */);  
}
```

If you return `kIOPMAckImplied`, you signal that you've completed the transition to the new power state. If you do not return `kIOPMAckImplied` and instead return the maximum amount of time it takes to prepare your device for the power-state change, you must be sure to call `acknowledgeSetPowerState` when you have finished the power-state transition. If you do not call `acknowledgeSetPowerState` before the length of time you specify has elapsed, the system continues with its power-state change as if you had returned `kIOPMAckImplied` in the first place.

Note: If you are developing a driver for OS X v10.5 or later, you may perform all necessary processing to prepare for the state change in the `setPowerState` method before you return `kIOPMAckImplied`. In other words, you do not have to return an estimate of how long the processing will take, perform the processing in another method, and call `acknowledgeSetPowerState` when the processing is finished.

5. Unregister from power management when your driver unloads using `PMstop`. The `PMstop` method handles all the necessary cleanup, including the removal of your driver from the power plane. Because `PMstop` may put your hardware into its off state, be sure to complete all hardware accesses before you call it.

Important: This step is crucial. If you neglect to call `PMstop`, you will probably cause a leak and you might cause a system panic the next time the computer wakes up.

In your driver's stop routine, after you finish all calls that might access your hardware, call `PMstop` as shown below:

```
PMstop();
```

Implementing Advanced Power Management

This section delves deeper into the power-management functionality of the I/O Kit. The vast majority of driver developers do not need to understand the information in this section because basic power management (as described in [Deciding How to Implement Power Management in Your Driver](#) (page 112)) is sufficient for most devices. If your device can be passively power managed, read [Implementing Basic Power Management](#) (page 113) instead.

You should read this section if your driver needs to perform advanced power-management tasks, such as determining device idleness, taking action when the system is about to shutdown, or deciding to change the device's power state. Of course, active drivers share some tasks with passive drivers, namely the initialization and tear-down of power management. Before you read about the tasks in this section, therefore, you should glance at the steps in [Implementing Basic Power Management](#) (page 113) to learn how to initialize and terminate power management in your driver. Even if your driver must perform advanced power-management tasks, it still needs to call `PMinit`, `joinPMtree`, `registerPowerDriver`, and `PMstop` and implement `setPowerState`, as shown in [Implementing Basic Power Management](#)

This section covers several tasks an active driver might need to perform. Although few active drivers will perform all the tasks, most will perform at least one. Each task is accompanied by a code snippet to help you implement it in your driver.

Defining and Using Multiple Power States

As described in [Devices and Power States](#) (page 111) information about a device's power states and capabilities must be available to I/O Kit power management. Although most devices have only the two required power states, off and on, some devices have additional states. As shown in step 3 of [Implementing Basic Power Management](#) (page 113) you construct an array of `IOPMPowerState` structures, each of which contains information about the device's capabilities in each power state. Table 9-1 describes the fields in the `IOPMPowerState` structure, which is defined in the `IOPM.h` header file.

Table 9-1 Fields and appropriate values in the `IOPMPowerState` structure

Field	Description	Value
version	Version number of this structure.	1
capabilityFlags	The capability of the device in this state.	An <code>IOPMPowerFlags</code> flag.
outputPowerCharacter	The power supplied in this state.	An <code>IOPMPowerFlags</code> flag.

Field	Description	Value
inputPowerRequirement	The input power required in this state.	An IOPMPowerFlags flag.
staticPower	Average power consumption (in milliwatts) of a device in this state.	0
unbudgetedPower	Additional power consumption (in milliwatts) from a separate power supply, such as a battery.	0
powerToAttain	The power consumed by a device (in milliwatts) in entering this state from the next lowest state.	0
timeToAttain	The time (in microseconds) required for a device to enter this state from the next lower state; in other words, the time required to program the hardware.	0
settleUpTime	The time (in microseconds) required to allow power to settle after entering this state from the next lower state.	0
timeToLower	The time (in microseconds) required for a device to enter the next lower state from this state; in other words, the time required to program the hardware.	0
settleDownTime	The time (in microseconds) required to allow power to settle after entering the next lower state from this state.	0
powerDomainBudget	The power (in milliwatts) that a power parent in this state is electronically able to deliver to its children.	0

As shown in Table 9-1 the values of some fields may be provided by an IOPMPowerFlags flag. Table 9-2 shows the IOPMPowerFlags flags you are likely to use.

Table 9-2 Power flags that describe device capabilities

Flag	Description
kIOPMPowerOn	The device is in the full-power state.

Flag	Description
kIOPMDeviceUsable	The clients of the device can use it in this state.
kIOPMMaxPerformance	The device is capable of its highest performance in this state.
kIOPMAuxPowerOn	The PCI auxiliary power supply is on (used only by devices in the PCI family).

Power management has the following requirements for the array of `IOPMPowerState` structures you construct in your driver's `start` method:

- The `IOPMPowerState` structure describing your device's off state must be the first element in the array.
- The `IOPMPowerState` structure describing your device's on (that is, full power) state must be the last element in the array.
- You can define any number of intermediate power states, but the `IOPMPowerState` structures describing them must not be the first or last elements of the array.

After you construct the power-state array to these specifications, call `registerPowerDriver`, passing in a pointer to the array and the number of power states. Listing 9-1 shows one way to do this. It also shows the driver creating a work loop and setting up a command gate to synchronize the power state-change code, which is described in [Changing the Power State of a Device](#) (page 120)

Listing 9-1 Building the power-state array and registering the driver

```
enum {
    kMyOffPowerState = 0,
    kMyIdlePowerState = 1,
    kMyOnPowerState = 2
};

static IOPMPowerState myPowerStates[3] = {
    {1, kMyOffPowerState, kMyOffPowerState, kMyOffPowerState, 0, 0, 0, 0, 0, 0, 0, 0},
    {1, kMyIdlePowerState, kMyIdlePowerState, kMyIdlePowerState, 0, 0, 0, 0, 0, 0, 0, 0},
    {1, kMyOnPowerState, kMyOnPowerState, kMyOnPowerState, 0, 0, 0, 0, 0, 0, 0, 0}
};

bool PMExampleDriver::start(IOService * provider)
{
```

```
/*
 * Create a work loop and set up synchronization
 * using a command gate.
 */
fWorkloop = IOWorkLoop::workLoop();
fGate = IOCommandGate::commandGate(this);

if (fGate && fWorkloop) {
    fWorkloop->addEventSource(fGate);
}

/*
 * Initialize power management, join the power plane,
 * and register with power management.
 */
PMininit();
provider->joinPMtree(this);
registerPowerDriver(this, myPowerStates, 3);
}
```

Changing the Power State of a Device

A driver is responsible for changing the power state of its device. Most power-state change requests come from power management when the system is about to sleep or wake. It's also possible for an active driver to become aware of the need to change its device's power state and initiate the request. The following sections describe both tasks.

Responding to a Power State–Change Request

As with a passive driver, an active driver must override the `setPowerState` method and change the power state of its device when it is instructed to do so. The ordinal value passed in to `setPowerState` is an index to the power-state array for the device.

If you're developing a driver to run in versions of OS X prior to v10.5, you must perform only the minimum processing required to change the power state of your device in your `setPowerState` method. Any additional processing must be performed outside of the `setPowerState` method and followed by a call to `acknowledgeSetPowerState` when it is finished. This is described in step 4 of [Implementing Basic Power Management](#) (page 113)

If, on the other hand, your driver will run in OS X v10.5 and later, you can perform all necessary processing in your `setPowerState` method before you return `kIOPMAckImplied`. It's important to understand, however, that power management calls the `setPowerState` method from a thread-call context. In other words, power management does not perform any automatic synchronization using your driver's work loop. Therefore, it's essential that you continue to use a command gate or other locking primitive to ensure that access to your device's state is serialized.

As soon as your driver returns `kIOPMAckImplied` or calls `acknowledgeSetPowerState` after additional processing, power management marks the power change as completed. Thus it's important for all drivers, regardless of the version of OS X they target, to avoid reporting a power change as complete until the power-state of the device has actually changed. It's possible that other power changes depend on your hardware having completed its power change before you call `acknowledgeSetPowerState`.

Initiating a Power-State Change

An active driver might become aware of the need to change its device's power state, either through mechanisms of its own or through some other object. The `IOService` class provides three methods that assist in this task:

- `makeUsable`
- `changePowerStateTo`
- `changePowerStateToPriv`

Any object in the driver stack, including a user client (described in [The Device-Interface Mechanism](#) (page 36)), can request that a dormant device be made active by calling the `makeUsable` method on the device's driver. The `makeUsable` method is interpreted as a request to put the device in its highest power state.

An active driver typically calls the `changePowerStateTo` method once in its `start` method, to set an initial power state. Later, when it wants to change its device's power state, an active driver calls the `changePowerStateToPriv` method, passing in the desired power state. An active driver might do this to shut down parts of the hardware that are not currently being used.

Important: Even though the `makeUsable`, `changePowerStateTo`, and `changePowerStateToPriv` methods are asynchronous and return immediately, you must block all hardware access until you receive the call to your `setPowerState` method. Before your `setPowerState` method is called, you cannot be certain that your device is in a usable state. Because of this, these functions should be called in a work-loop context.

Power management uses the states passed in to `changePowerStateTo` and `changePowerStateToPriv` to determine the device's new power state. Specifically, power management selects as the new power state the highest value of the following three values:

- The power state set by `changePowerStateToPriv`
- The power state set by `changePowerStateTo`
- The highest of all power states required by the driver's power children

The following code snippet shows how a driver can get the device's current power state (using the `getPowerState` method introduced in OS X v10.5) and then request a power-state change with `changePowerStateToPriv`.

```
enum {
    kMyOffState = 0,
    kMyOnState = 1
};
/*
 * Make sure the hardware is in the ON state
 * before accessing it. If it's powered off, call changePowerStateToPriv
 * to put the device in the ON state.
 */
if (getPowerState() == kMyOnState)
{
    /* Device is ON. OK to access hardware. */
} else {
    changePowerStateToPriv( kMyOnState );
}
/*
 * Note: If your device has been powered off for a system sleep, you cannot
 * try to adjust your power state upwards. You are locked in your OFF or
 * low-power state until system power is restored on wake.
```

```
    */

    /*
    * Although changePowerStateToPriv returns immediately,
    * it is _NOT_ safe to touch the hardware yet. You must wait until you
    * receive your setPowerState() call before you can safely modify
    * the hardware.
    */
}
```

Implementing Idleness Determination and Idle Power Saving

When a device is idle, it can be powered down to conserve system power, which is especially important for laptop computers running on battery power. You should implement idle power saving in your device if:

- Access to your device is intermittent, and the device is often left unused for minutes, hours, or days at a time.
- Your device consumes a significant amount of power, and putting it in a low power state when possible results in substantial power savings.

To implement idle power saving, you must determine when your device is idle and specify how long the period of idleness should last before your device powers off. You determine idleness by supplying device-access information to the `IOService` superclass, which uses this information, in conjunction with the idleness period you specify, to tell your device to power off at the appropriate time. The `IOService` class provides two methods an active driver uses to do this:

- `activityTickle`. On your device's access path, you call `activityTickle` every time your driver or any other client (including an application) triggers a hardware access. This allows power management to confirm that your device is in a usable state and to track the most recent access times for your device.
- `setIdleTimerPeriod`. You call `setIdleTimerPeriod` to specify the duration of a watchdog timer that tracks how long your device can be idle at full power before it should be powered down. By setting the duration of the idle period, you effectively start a countdown that begins after each device access.

When the idle period expires without any device activity, power management calls your implementation of the `setPowerState` method to lower your device's power state. See [Changing the Power State of a Device](#) (page 120) for more information on how to implement this method.

Of course, you must respond to any device-access request you receive while your device is powered off by first setting your device to its full-power state. Because you call `activityTickle` on your device's access path, power management is immediately alerted to the fact that some entity is requesting access to a device that is currently powered off. When this happens, the `IOService` superclass automatically calls `makeUsable` on your device, which ultimately results in a call to your implementation of the `setPowerState` method.

Important: As described in [Initiating a Power-State Change](#) (page 121) you must block all hardware accesses until your `setPowerState` implementation is called, because you cannot be certain that your device is in a usable state until that time. The best way to do this is to use a work loop to serialize hardware access.

The following steps outline the process of idleness determination:

1. Specify how long your device should remain in a high-power state while idle. Typically, one minute is an appropriate interval.

Call `setIdleTimerPeriod`, passing in the idle interval in seconds, as shown below:

```
setIdleTimerPeriod ( 60 );
```

2. Inform power management every time an entity (including your driver) initiates a device access.

In your driver's device-access path, call `activityTickle`, as shown below:

```
activityTickle ( kIOPMSuperclassPolicy1, myDevicePowerOn );
```

As shown above, the first parameter to `activityTickle` is `kIOPMSuperclassPolicy1`, which indicates that the `IOService` superclass will track device activity and take action when the idle period expires. The second parameter specifies the power state required for this activity, typically the on state.

3. When the idle timer expires, the `IOService` superclass checks whether there has been any device activity since the last idle timer expiration. The superclass determines this by checking when `activityTickle (kIOPMSuperclassPolicy1)` was last called.
4. If there has been device activity since the last timer expiration, the `IOService` superclass restarts the timer. If no device activity has occurred, the `IOService` superclass calls `setPowerState` on your driver to power down the device to the next lowest state.

Optional: A driver may implement variable idle timeout behaviors by overriding the `IOService` method `nextIdleTimeout`. To do this, your implementation of `nextIdleTimeout` should return how many "seconds from now" the device should move into its next lowest power state.

For example, the Graphics family uses `nextIdleTimeout` to dynamically adjust the display's idle-sleep timeouts. If the user moves the mouse very soon after the display dims, the display driver remembers this and increases the timeout period, effectively waiting a longer time before it initiates the next several display-dim events.

After your device has been powered down to a lower state through this process, a new `activityTickle` invocation causes power management to raise the device's power to the level required for the activity. If the device is already in the correct state, the superclass simply returns `true` from the call to `activityTickle (kIOPMSuperclassPolicy1)`; otherwise, the superclass returns `false` and proceeds to make the device usable.

Although the return value of `activityTickle` indicates whether the device is in a usable power state, it's better to keep track of your device's current power state in your driver than to rely on the `activityTickle` return value for this information. This is because `activityTickle` is not called on the power management work loop and a device's power state might change before `activityTickle` returns.

Receiving Notification of Power-State Changes in Other Devices

In some cases, your driver might need to be notified when another driver changes its device's power state. I/O Kit power management brackets each change to the power state of a device with a pair of notifications. These notifications are delivered through invocations of the `IOService` virtual methods `powerStateWillChangeTo` and `powerStateDidChangeTo`. You can implement these methods to receive the notifications and prepare for the changes.

Your driver can register its interest in another driver, as long as the following is true:

- The driver in which your driver is interested must be attached into the power plane.
- Your driver must be a C++ subclass of `IOService`, but it does not have to be attached into the power plane itself.

To find out when another driver changes its device's power state, follow these steps in your driver:

1. Call the `IOService` method `registerInterestedDriver`. This ensures that power management will notify your driver when it sends out power-change notifications.
2. Implement the virtual `IOService` method `powerStateWillChangeTo`. This method is called by the device's driver when it is about to change the device's power state.

If your driver is prepared for the change, it should return `kIOPMAckImplied`; if it needs more time to prepare, it should return an upper limit on the time required (in microseconds).

If your driver returns a number representing the maximum preparation time needed, it should call the `acknowledgePowerChange` method when it is prepared. If it does not do this, and the time requested for preparation elapses, the other driver carries on as if your driver had acknowledged the change. This behavior prevents power-state changes from stalling because of failing drivers.

Important: The `powerStateWillChangeTo` method is not the place to perform any tasks related to the actual changing of the power state. Such tasks should be performed in the `setPowerState` method, which is described in [Responding to a Power State–Change Request](#) (page 120)

3. Implement the virtual `IOService` method `powerStateDidChangeTo`. This method is called by the device's driver after the power-state change is complete.

After the change to power occurs and the power has settled to its new level, power management broadcasts this fact to all interested objects via the `powerStateDidChangeTo` method. If a device is going to a reduced power state, interested drivers generally don't need to do much with this notification. However, if the device is going to a higher power state, interested drivers would use this notification to prepare for the change by, for example, restoring state or programming a device.

In your implementation of `powerStateDidChangeTo`, your driver can examine the `IOPMPowerFlags` bitfield (described in [Table 9-2](#) (page 118)) passed in to make its determination; this bitfield is derived from the `capabilityFlags` field of the power-state array, which is described in [Table 9-1](#) (page 117). As with `powerStateWillChangeTo`, your driver should return `kIOPMAckImplied` if it has prepared for the change. If it needs time to prepare, it should return the maximum time required (in microseconds); when your driver is finally ready for the change, it should call the `acknowledgePowerChange` method.

Important: The `powerStateDidChangeTo` method is not the place to perform any tasks related to the actual changing of the power state. Such tasks should be performed in the `setPowerState` method, which is described in [Responding to a Power State–Change Request](#) (page 120)

When your driver is no longer interested in the power changes of other drivers, it should deregister itself to stop receiving notifications. To do this, call the `IOService` method `deRegisterInterestedDriver`, usually in your driver's `stop` method.

Receiving Shutdown and Restart Notifications

In a driver targeting OS X v10.5 and later, you can implement the `systemWillShutdown` method to receive notification of an impending shutdown or restart. It's important to understand, however, that there is nothing your driver can do to prevent a shutdown, regardless of the notification it receives. Your driver is capable of delaying shutdown, but that is strongly discouraged because it can severely degrade the user's experience.

Do not assume that your driver should implement `systemWillShutdown` so that it can respond to shutdown and restart notifications by shutting down your hardware. At shutdown time, power is about to be removed from your device regardless of its current state. Similarly, if the system is restarting, your device will be reinitialized shortly and, again, its current state is not important. Most built-in device drivers in OS X do not shut down their devices when the system is about to shutdown or restart and most third-party device drivers should do the same.

Although the majority of device drivers do not need to handle shutdown or restart in any way at all, there are two valid reasons for a driver to run at shutdown or restart time:

- The architecture requires the driver to execute code at shutdown time. For example, all drivers that perform DMA in an Intel-based Macintosh *must* stop active DMA before shutdown can complete.
- The driver must run at shutdown or restart time to avoid negative user experience. For example, an audio driver might need to turn off its device's amplifiers to avoid an audible "pop" when power is removed.

Note: There are other places that you may run code on the shutdown path. For example, if your software has a user-space daemon that runs until shutdown, that daemon can catch `SIGTERM`, which the kernel sends to all processes at shutdown. In general, try to run your shutdown code as early as possible in the shutdown path.

The `systemWillShutdown` method is called on all members of the power plane, in leaf-to-root order. A driver's `systemWillShutdown` method is invoked only after all its power children have completed their shutdown tasks. This ensures that a child object can handle its shutdown or restart tasks before its parent powers off. Note that it is not necessary to call your driver's `free` method when the system is about to restart or shutdown, because all drivers are unloaded and destroyed at this time.

When a driver receives the `systemWillShutdown` call, it performs the necessary tasks to prepare for the shutdown or restart and then invokes its superclass's implementation of the method. This is essential, because system shutdown will stall until all drivers have finished handling their `systemWillShutdown` notifications. Other than postponing the call to `super::systemWillShutdown` until an in-flight I/O request completes, you should do everything possible to avoid delaying shutdown. Listing 9-2 shows how to override `systemWillShutdown` and receive notification of shutdown or restart.

Listing 9-2 Getting notification of system shutdown or restart

```
void MyExampleDriver::systemWillShutdown( IOOptionBits specifier )
{
    if ( kIOMessageSystemWillPowerOff == specifier ) {
        // System is shutting down; perform appropriate processing.
    } else if ( kIOMessageSystemWillRestart == specifier ) {
        // System is restarting; perform appropriate processing.
    }
    /*
     * You must call your superclass's implementation of systemWillShutdown as
     * soon as you're finished processing your shutdown or restart
     * because the shutdown will not proceed until you do.
     */
    super::systemWillShutdown( specifier );
}
```

Keeping Power On for Future Device Attachment

To conserve power, a device whose children have all disappeared is usually considered idle and is told to power off. However, your device might need to stay powered on to allow new children to attach at any time. For example, a bus might need to remain powered on even when there are no devices attached to it, because a new device trying to attach can cause a crash by attempting to access hardware that's turned off.

The `IOService` class provides a method that allows you to keep your device's power on, even if all its power children have disappeared. The `clampPowerOn` method allows you to specify a length of time to keep the device in its highest power state. If you need to do this in your driver, call the `clampPowerOn` method before the last power child disappears, as shown below:

```
// timeToStayOn is a length of time in milliseconds.
clampPowerOn ( timeToStayOn );
```


Managing Device Removal

OS X is an operating system that includes hot-swapping as a feature. Users can plug in and remove external devices (for example, mass-storage drives, CD-RW drives, modems, and scanners) and the system immediately does what is necessary to make the device usable or, in the case of removal, to register the absence of the device. No system restart or shutdown is necessary.

This chapter describes how your driver should respond to the removal of its device.

The Phases of Device Removal

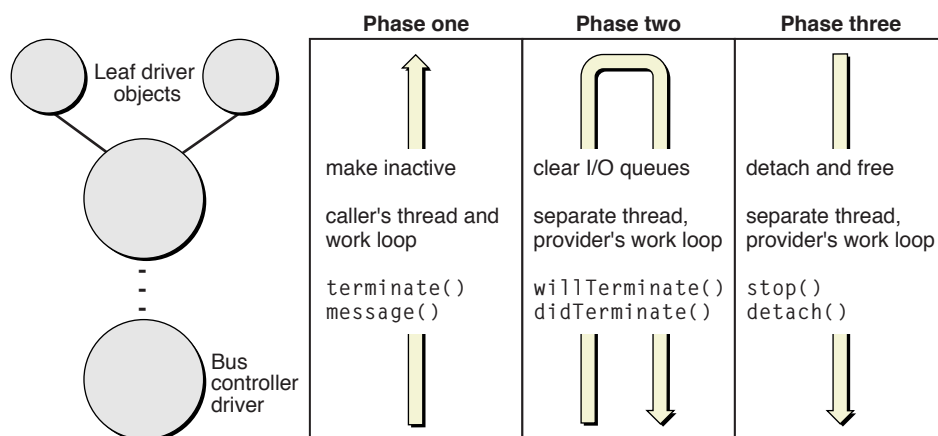
When a user plugs a device into the system, the I/O Kit responds to the event using the normal process for discovering and loading drivers. A low-level driver scans its bus, notices a new device, and kicks off the matching process to find a suitable driver. The I/O Kit then loads the driver into the kernel and the device is usable.

When a user removes a device, however, the situation is different. A driver stack must be torn down rather than built up. Before the drivers in the stack can be released, they must, in a coordinated manner, stop accepting new requests and clear out all queued and in-progress work; this requires a special programming interface and procedure.

The I/O Kit performs an orderly tear-down of a driver stack upon device removal in three phases. The first phase makes the driver objects in the stack inactive so they receive no new I/O requests. The second phase clears out pending and in-progress I/O requests from driver queues. Finally, in the third phase, the I/O Kit invokes

the appropriate driver life-cycle methods on drivers to clean up allocated resources and detach the objects from the I/O Registry before freeing them. [Figure 10-1](#) (page 130) summarizes what happens during the three phases, including the calling direction within the driver stack.

Figure 10-1 Phases of device removal



Making Drivers Inactive

Just as a bus controller driver scans its bus to detect a newly inserted device, it also detects devices that have just been removed. When this happens, it calls `terminate` on its client nub; the `terminate` method has the default behavior of making the called object inactive immediately. The `terminate` method is also recursively invoked on clients; it is called on each object in the stack above the bus controller until all objects in the stack are made inactive.

As a consequence of being made inactive, each object also sends its clients (or, in rare cases, providers) a `kIOWServicesIsTerminated` message via the `message` method. When the `terminate` call returns to the original caller (the bus controller driver), all objects in the stack are inactive, but the stack is still attached to the I/O Registry.

The I/O Kit assumes that objects that have multiple providers (drivers of RAID devices, for instance) do not want to be torn down and thus does not call `terminate` on them. If these objects do want to receive the `terminate` message, they should implement the `requestTerminate` method to return `true`.

The `terminate` call is asynchronous to avoid deadlocks and, in this first phase, takes place in the thread and work-loop context of the caller, the bus controller driver.

Clearing I/O Queues

The I/O Kit itself coordinates the second phase of the device-removal procedure. It starts with the newly inactive client of the bus controller driver and, as in the first phase, moves up the driver stack until it reaches the leaf objects. It calls the `willTerminate` method on each object it encounters. Drivers should implement the `willTerminate` method to clear out any queues of I/O requests they have. To do this, they should return an appropriate error to the requester for each request in a queue.

After `willTerminate` has been called on each object, I/O Kit then reverses direction, going from the leaf object (or objects) to the root object of the driver stack, calling `didTerminate` on each object. Certain objects at the top of the stack—particularly user clients—may decide to keep a count of I/O requests they have issued and haven't received a response for ("in-flight" I/O requests). (To ensure the validity of this count, the object should increment and decrement the count in a work-loop context.) By this tracking mechanism, they can determine if any I/O request hasn't been completed. If this is the case, they can implement `didTerminate` to return a deferral response of `true`, thereby deferring the termination until the final I/O request completes. At this point, they can signal that termination should proceed by invoking `didTerminate` on themselves and returning a deferral response of `false`.

If a driver attaches to a client nub and has it open, the I/O Kit assumes a deferred response (`true`) to `didTerminate`. The termination continues to be deferred until the client driver closes its provider.

At the end of this second phase, there shouldn't be any I/O requests queued or "in flight." The I/O Kit completes this phase of the device-removal procedure on its own separate thread and makes all calls to clients on the work-loop context of the provider.

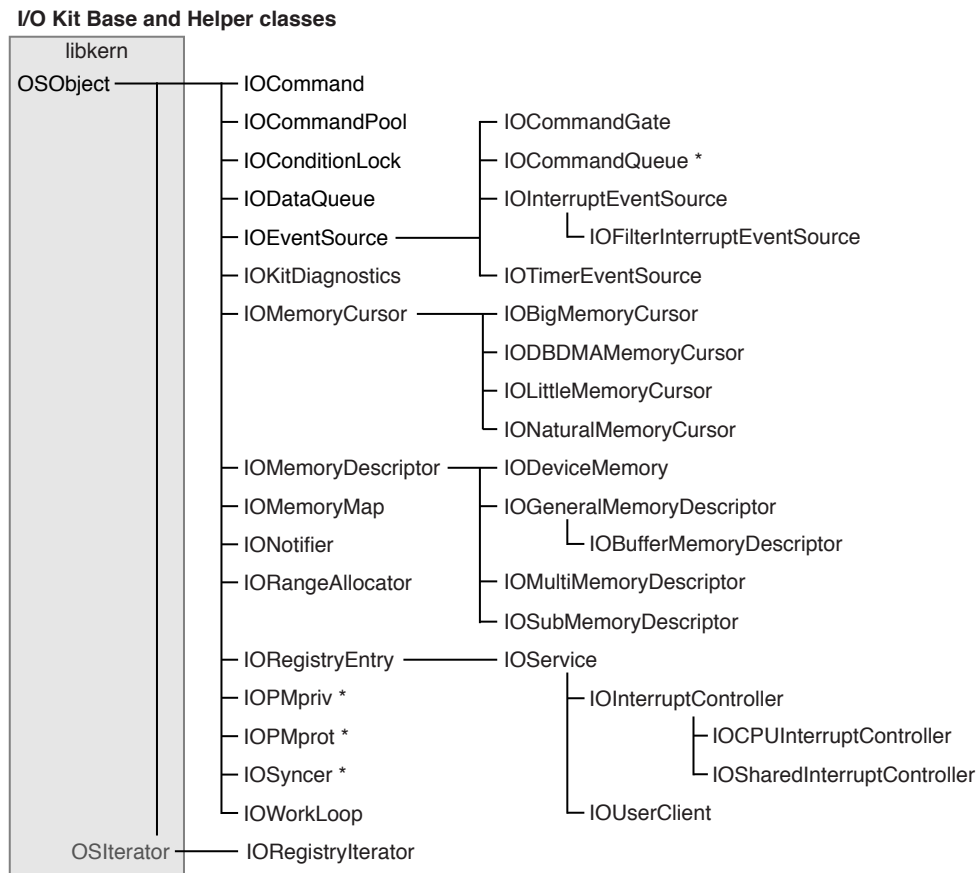
Detaching and Releasing Objects

In the third phase of the procedure for device removal, the I/O Kit calls the driver life-cycle methods `stop` and `detach` (in that order) in each driver object in the stack to be torn down, starting from the leaf object (or objects). The driver should implement its `stop` function to close, release, or free any resources it opened or created in its `start` function, and to leave the hardware in the state the driver originally found it. The driver can implement `detach` to remove itself from its nub through its entry in the I/O Registry; however, this behavior is the default, so a driver usually does not need to override `detach`. The `detach` method usually leads immediately to the freeing of the driver object because the I/O Registry typically has the final `retain` on the object.

The I/O Kit completes this phase of the device-removal procedure on its own separate thread and makes all calls to clients on the work-loop context of the provider.

Base and Helper Class Hierarchy

The following chart presents the class hierarchy of all I/O Kit classes that are not members of a specific family. See the preceding appendix, [I/O Kit Family Reference](#) (page 142) for the class hierarchy charts of most families.



* Classes that are special-purpose or to be obsoleted; do not subclass.

Bibliography

System Internals

Advanced Topics in UNIX: Processes, Files, & Systems, Ronald J. Leach, Wiley, 1996, ISBN 1-57176-159-4

The Design and Implementation of the 4.4BSD UNIX Operating System, Marshall Kirk McKusick, et al, Addison-Wesley, 1996, ISBN 0-201-54979-4

Panic! UNIX System Crash Dump Analysis Chris Drake, Kimberly Brown Prentice Hall, 1995, ISBN 0-13-149386-8

UNIX Internals: The New Frontiers, Uresh Vahalia, Prentice-Hall, 1995, ISBN 0-13-101908-2

Websites - Online Resources

Apple Computer's developer website (<http://developer.apple.com/index.html>) is a general repository for developer documentation. Additionally, the following sites provide more domain-specific information.

Apple's Public Source projects and Darwin

<http://developer.apple.com/darwin/projects/darwin/>

CVS (Concurrent Versions System)

<http://developer.apple.com/opensource/cvs.html>

GDB, GNUPro Toolkit 99r1 Documentation

<http://www.redhat.com/docs/manuals/gnupro/>

The Internet Engineering Task Force (IETF)

<http://www.ietf.org/>

jam

<http://www.perforce.com/jam/jam.html>

The USENIX Association

<http://www.usenix.org/>

USENIX Proceedings

<http://www.usenix.org/publications/library/>

Glossary

active driver A device driver that implements advanced power management tasks, such as determining device idleness and performing pre-shutdown tasks. See also [passive driver](#)

base class In C++, the class from which another class (a subclass) inherits. It can also be used to specify a class from which all classes in a hierarchy ultimately derive (also known as a root class).

BSD Berkeley Software Distribution. Formerly known as the Berkeley version of UNIX, BSD is now simply called the BSD operating system. The BSD portion of Darwin is based on 4.4BSD Lite 2 and FreeBSD, a flavor of 4.4BSD.

bundle A directory in the file system that typically stores executable code and the software resources related to that code. (A bundle can store only resources.) Applications, plug-ins, frameworks, and kernel extensions are types of bundles. Except for frameworks, bundles are file packages, presented by the Finder as a single file instead of a folder. See also [kernel extension](#)

bus A transmission path on which signals can be dropped off or picked up by devices attached to it. Only devices addressed by the signals pay attention to them; the others discard the signals. Buses both exist within the CPU and connect it to physical memory and peripheral devices. Examples of I/O buses on Darwin are PCI, SCSI, USB, and FireWire.

bus master A program, usually in a separate I/O controller, that directs traffic on the computer bus or input/output paths. The bus master actually controls the bus paths on which the address and

control signals flow. DMA is a simple form of bus mastering where the bus master controls I/O transfers between a device and system memory and then signals to the CPU when it has done so. See also [DMA](#)

client A driver object that consumes services of some kind supplied by its provider. In a driver stack, the client in a provider/client relationship is farther away from the Platform Expert. See also [provider](#)

command gate A mechanism that controls access to the lock of a work loop, thereby serializing access to the data involved in I/O requests. A command gate does not require a thread context switch to ensure single-threaded access. IOCommandGate event-source objects represent command gates in the I/O Kit.

Darwin Another name for the OS X core operating system, or kernel environment. The Darwin kernel environment is equivalent to the OS X kernel plus the BSD libraries and commands essential to the BSD Commands environment. Darwin is Open Source technology.

DMA (Direct Memory Access) A capability of some bus architectures that enables a bus controller to transfer data directly between a device (such as a disk drive) and a device with physically addressable memory, such as that on a computer's motherboard. The microprocessor is freed from involvement with the data transfer, thus speeding up overall computer operation. See also [bus master](#)

device Computer hardware, typically excluding the CPU and system memory, which can be controlled and can send and receive data. Examples of devices include monitors, disk drives, buses, and keyboards.

device driver A component of an operating system that deals with getting data to and from a device, as well as the control of that device. A driver written with the I/O Kit is an object that implements the appropriate I/O Kit abstractions for controlling hardware.

device file In BSD, a device file is a special file located in `/dev` that represents a block or character device such as a terminal, disk drive, or printer. If a program knows the name of a device file, it can use POSIX functions to access and control the associated device. The program can obtain the device name (which is not persistent across reboots or device removal) from the I/O Kit.

device interface In the I/O Kit, a mechanism that uses a plug-in architecture to allow a program in user space to communicate with a nub in the kernel that is appropriate to the type of device the program wishes to control. Through the nub the program gains access to I/O Kit services and to the device itself. From the perspective of the kernel, the device interface appears as a driver object called a user client.

device matching In the I/O Kit, a process by which an application finds an appropriate device interface to load. The application calls a special I/O Kit function that uses a “matching dictionary” to search the I/O Registry. The function returns one or more matching driver objects that the application can then use to load an appropriate device interface. Also referred to as device discovery.

driver See [device driver](#)

driver matching In the I/O Kit, a process in which a nub, after discovering a specific hardware device, searches for the driver or drivers most suited to drive that device. Matching requires that a driver have one or more personalities that specify whether it is a candidate for a particular device. Driver matching is a subtractive process involving three phases: class matching, passive matching, and active matching. See also [personality](#)

driver stack In an I/O connection, the series of driver objects (drivers and nubs) in client/provider relationships with each other. A driver stack often refers to the entire collection of software between a device and its client application (or applications).

event source An I/O object that corresponds to a type of event that a device driver can be expected to handle; there are currently event sources for hardware interrupts, timer events, and I/O commands. The I/O Kit defines a class for each of these event types, respectively `IOInterruptEventSource`, `IOTimerEventSource`, and `IOCommandGate`.

family A collection of software abstractions that are common to all devices of a particular category. Families provide functionality and services to drivers. Examples of families include protocol families (such as SCSI, USB, and Firewire), storage families (disk drives), network families, and families that describe human interface devices (mouse and keyboard).

fault In the virtual-memory system, faults are the mechanism for initiating page-in activity. They are interrupts that occur when code tries to access data at a virtual address that is not mapped to physical memory. See also [page](#); [virtual memory](#)

framework A type of bundle that packages a dynamic shared library with the resources that the library requires, including header files and reference

documentation. Note that the Kernel framework (which contains the I/O Kit headers) contains no dynamic shared library. All library-type linking for the Kernel framework is done using the `mach_kernel` file itself and kernel extensions. This linking is actually static (with vtable patch-ups) in implementation.

idle sleep A sleep state that occurs when there has been no device or system activity for the period of time the user specifies in the Energy Saver pane of System Preferences. See also [system sleep](#)

information property list A property list that contains essential configuration information for bundles such as kernel extensions. A file named `Info.plist` (or a platform-specific variant of that filename) contains the information property list and is packaged inside the bundle.

interrupt An asynchronous event that suspends the currently scheduled process and temporarily diverts the flow of control through an interrupt handler routine. Interrupts can be caused by both hardware (I/O, timer, machine check) and software (supervisor, system call, or trap instruction).

interrupt handler A routine executed when an interrupt occurs. Interrupt handlers typically deal with low-level events in the hardware of a computer system, such as a character arriving at a serial port or a tick of a real-time clock.

I/O Catalog A dynamic database that maintains entries for all available drivers on a Darwin system. Driver matching searches the I/O Catalog to produce an initial list of candidate drivers.

I/O Kit A kernel-resident, object-oriented environment in Darwin that provides a model of system hardware. Each type of service or device is

represented by one or more C++ classes in a family; each available service or device is represented by an instance (object) of that class.

I/O Kit framework The framework that includes IOKitLib and makes the I/O Registry, user client plug-ins, and other I/O Kit services available from user space. It lets applications and other user processes access common I/O Kit object types and services. See also [framework](#)

I/O Registry A dynamic database that describes a collection of driver objects, each of which represents an I/O Kit entity. As hardware is added to or removed from the system, the I/O Registry changes to accommodate the addition or removal.

kernel The complete OS X core operating-system environment, which includes Mach, BSD, the I/O Kit, drivers, file systems, and networking components. The kernel resides in its own protected memory partition. The kernel includes all code executed in the kernel task, which consists of the file `mach_kernel` (at file-system root) and all loaded kernel extensions. Also called the kernel environment.

kernel extension (KEXT) A dynamically loaded bundle that extends the functionality of the kernel. A KEXT can contain zero or one kernel modules as well as other (sub) KEXTs, each of which can contain zero or one kernel modules. The I/O Kit, file system, and networking components of Darwin can be extended by KEXTs. See also [kernel module](#)

kernel module (KMOD) A binary in Mach-O format that is packaged in a kernel extension. A KMOD is the minimum unit of code that can be loaded into the kernel. See also [kernel extension](#)

lock A data structure used to synchronize access to a shared resource. The most common use for a lock is in multithreaded programs where multiple threads

need access to global data. Only one thread can hold the lock at a time; by convention, this thread is the only one that can modify the data during this period. See also [mutex](#)

Mach A central component of the kernel that provides such basic services and abstractions as threads, tasks, ports, interprocess communication (IPC), scheduling, physical and virtual address space management, virtual memory, and timers.

map To translate a range of memory in one address space (physical or virtual) to a range in another address space. The virtual-memory manager accomplishes this by adjusting its VM tables for the kernel and user processes.

matching See [device matching](#); [driver matching](#)

memory cursor An object that lays out the buffer ranges in a memory descriptor in physical memory, generating a scatter/gather list suitable for a particular device or DMA engine. The object is derived from the `IOMemoryCursor` class. See also [DMA](#); [memory descriptor](#)

memory descriptor An object that describes how a stream of data, depending on direction, should either be laid into memory or extracted from memory. It represents a segment of memory holding the data involved in an I/O transfer and is specified as one or more physical or virtual address ranges. The object is derived from the `IOMemoryDescriptor` class. See also [DMA](#); [memory cursor](#)

memory protection A system of memory management in which programs are prevented from being able to modify or corrupt the memory partition of another program. Although OS X has memory protection, Mac OS 8 and 9 do not.

mutex A mutual-exclusion locking object that allows multiple threads to synchronize access to shared resources. A mutex has two states: locked and unlocked. Once a mutex has been locked by a thread, other threads attempting to lock it will block. When the locking thread unlocks (releases) the mutex, one of the blocked threads (if any) acquires (locks) it and uses the resource. The thread that locks the mutex must be the one that unlocks it. The work-loop lock (which is used by a command gate) is based on a mutex. See also [lock](#); [work loop](#)

notification A programmatic mechanism for alerting interested recipients (sometimes called observers) that an event has occurred.

nub An I/O Kit object that represents a detected, controllable entity such as a device or logical service. A nub may represent a bus, disk, graphics adaptor, or any number of similar entities. A nub supports dynamic configuration by providing a bridge between two drivers (and, by extension, between two families). See also [device](#); [driver](#)

page (1) The smallest unit (in bytes) of information that the virtual memory system can transfer between physical memory and backing store. In Darwin, a page is currently 4 kilobytes. (2) As a verb, page refers to the transfer of pages between physical memory and backing store. Refer to `Kernel.framework/Headers/mach/machine/vm_params.h` for specifics. See also [fault](#); [virtual memory](#)

passive driver A device driver that performs only basic power-management tasks, such as joining the power plane and changing the device's power state. See also [active driver](#)

personality A set of properties specifying the kinds of devices a driver can support. This information is stored in an XML matching dictionary defined in the information property list (`Info.plist`) file in the

driver's KEXT bundle. A single driver may present one or more personalities for matching; each personality specifies a class to instantiate. Such instances are passed a reference to the personality dictionary at initialization.

physical memory Electronic circuitry contained in random-access memory (RAM) chips, used to temporarily hold information at execution time. Addresses in a process's virtual memory are mapped to addresses in physical memory. See also [virtual memory](#)

PIO (Programmed Input/Output) A way to move data between a device and system memory in which each byte is transferred under control of the host processor. See also [DMA](#)

plane A subset of driver (or service) objects in the I/O Registry that have a certain type of provider/client relationship connecting them. The most general plane is the Service plane, which displays the objects in the same hierarchy in which they are attached during Registry construction. There are also the Audio, Power, Device Tree, FireWire, and USB planes.

Platform Expert A driver object for a particular motherboard that knows the type of platform the system is running on. The Platform Expert serves as the root of the I/O Registry tree.

plug-in A module that can be dynamically added to a running system or application. Core Foundation Plug-in Services uses the basic code-loading facility of Core Foundation Bundle Services to provide a standard plug-in architecture, known as the CFPlugIn architecture, for Mac apps. A kernel extension is a type of kernel plug-in.

port A heavily overloaded term which in Darwin has two particular meanings: (1) In Mach, a secure unidirectional channel for communication between

tasks running on a single system; (2) In IP transport protocols, an integer identifier used to select a receiver for an incoming packet or to specify the sender of an outgoing packet.

POSIX The Portable Operating System Interface. An operating-system interface standardization effort supported by ISO/IEC, IEEE, and The Open Group.

power child In the power plane, a driver for a device that relies on another object for its power. See also [power parent](#); [plane](#)

power parent In the power plane, an object that provides power for a device. See also [power child](#); [plane](#)

preemptive multitasking A type of multitasking in which the operating system can interrupt a currently running program in order to run another program, as needed.

probe A phase of active matching in which a candidate driver communicates with a device and verifies whether it can drive it. The driver's `probe` member function is invoked to kick off this phase. The driver returns a probe score that reflects its ability to drive the device. See also [driver matching](#)

process A BSD abstraction for a running program. A process' resources include a virtual address space, threads, and file descriptors. In OS X, a process is based on one Mach task and one or more Mach threads.

provider A driver object that provides services of some kind to its client. In a driver stack, the provider in a provider/client relationship is closer to the Platform Expert. See also [client](#)

release Decrementing the reference count of an object. When an object's reference count reaches zero, it is freed. When your code no longer needs to reference a retained object, it should release it. Some

APIs automatically execute a release on the caller's behalf, particularly in cases where the object in question is being "handed off." Retains and releases must be carefully balanced; too many releases can cause panics and other unexpected failures due to accesses of freed memory. See also [retain](#)

retain Incrementing the reference count of an object. An object with a positive reference count is not freed. (A newly created object has a reference count of one.) Drivers can ensure the persistence of an object beyond the present scope by retaining it. Many APIs automatically execute a retain on the caller's behalf, particularly APIs used to create or gain access to objects. Retains and releases must be carefully balanced; too many retains will result in wired memory leak. See also [release](#)

scheduler That part of Mach that determines when each program (or program thread) runs, including assignment of start times. The priority of a program's thread can affect its scheduling. See also [task](#); [thread](#)

service A service is an I/O Kit entity, based on a subclass of `IOService`, that has been published with the `registerService` method and provides certain capabilities to other I/O Kit objects. In the I/O Kit's layered architecture, each layer is a client of the layer below it and a provider of services to the layer above it. A service type is identified by a matching dictionary that describes properties of the service. A nub or driver can provide services to other I/O Kit objects.

socket In BSD-derived systems such as Darwin, a socket refers to different entities in user and kernel operations. For a user process, a socket is a file descriptor that has been allocated using `socket(2)`. For the kernel, a socket is the data structure that is allocated when the kernel's implementation of the `socket(2)` call is made.

system sleep A sleep state that occurs when the user chooses Sleep from the Apple menu or closes the lid of a laptop computer. See also [idle sleep](#)

task A Mach abstraction consisting of a virtual address space and a port name space. A task itself performs no computation; rather, it is the context in which threads run. See also [process](#); [thread](#)

thread In Mach, the unit of CPU utilization. A thread consists of a program counter, a set of registers, and a stack pointer. See also [task](#)

timer A kernel resource that triggers an event at a specified interval. The event can occur only once or can be recurring. Timers are one of the event sources for work loops.

user client An interface provided by an I/O Kit family, that enables a user process (which can't call a kernel-resident driver or other service directly) to access hardware. In the kernel, this interface appears as a driver object called a user client; in user space, it is called a device interface and is implemented as a Core Foundation Plug-in Services (CFPlugin) object. See also [device interface](#)

user space Virtual memory outside the protected partition in which the kernel resides. Applications, plug-ins, and other types of modules typically run in user space.

virtual address A memory address that is usable by software. Each task has its own range of virtual addresses, which begins at address zero. The Mach operating system makes the CPU hardware map these addresses onto physical memory only when necessary, using disk memory at other times.

virtual memory The use of a disk partition or a file on disk to provide the same facilities usually provided by RAM. The virtual-memory manager in

OS X provides 32-bit (minimum) protected address space for each task and facilitates efficient sharing of that address space.

wired memory A range of memory that the virtual-memory system will not page out or move. The memory involved in an I/O transfer must be wired down to prevent the physical relocation of data being accessed by hardware. In the I/O Kit memory is wired when the memory descriptor describing the memory prepares the memory for I/O (which happens when its `prepare` method is invoked).

work loop A gating mechanism that ensures single-threaded access to the data structures and hardware registers used by a driver. Specifically, it is a mutex lock associated with a thread. A work loop typically has several event sources attached to it; they use the work loop to ensure a protected, gated context for processing events. See also [event source](#)

I/O Kit Family Reference

This appendix describes each of the I/O Kit families in detail, paying particular attention to client/provider relationships. For most families, it provides a class hierarchy chart. It also tells you if a family exports a device interface, thereby allowing applications to access devices represented by the family. You should seriously consider taking the device-interface approach before attempting to write a kernel-resident driver. For information on using device interfaces, see the document *Accessing Hardware From Applications*.

Some categories of devices are not currently supported by an I/O Kit family. If your device falls into an unsupported category, you might be able to write a “family-less” driver, use an SDK other than the I/O Kit, or create a new family. See [Devices Without I/O Kit Families](#) (page 176) for details.

You may find it helpful to examine the source code for an I/O Kit family or a specific device driver. To do this, visit [Darwin Releases](#), select the appropriate version of OS X, and click Source to view the available source projects.

ADB

The ADB family provides support for, and access to, devices attached to the Apple Desktop Bus (ADB). It provides an abstraction for ADB bus controller drivers (IOADBController) and another for drivers of ADB devices (IOADBDevice).

Bundle identifier:

- `com.apple.iokit.IOADBFamily`

Headers at:


- Kernel resident: `Kernel.framework/Headers/IOKit/adb/`
- Device interface: `IOKit.framework/Headers/adb`

References and specifications:

- Chapter 5 — “ADB Manager” of *Inside Macintosh: Devices* (<http://developer.apple.com/documentation/Hardware/DeviceManagers/adb/adb.html>)
- Technical note HW01 (http://developer.apple.com/technotes/hw/hw_01.html)

- *Guide to the Macintosh Family Hardware*, second edition, Apple Computer.

Class hierarchy:



Device Interface:

- Exports an interface for reading and writing registers on ADB devices. The interface is defined in `IOADBLib.h`. Only polled mode operations are supported through this library. Interrupt operations are only supported for kernel-resident clients.

Table A-1 Clients and providers of the ADB family

	Client of the nub	Provider for the nub
Action	Drives devices that plug into an ADB port.	Drives an ADB bus controller
Example	A driver for an ADB mouse is a <i>client</i> of the ADB family but is a <i>member</i> of the HID family (in the <code>IOHIDPointing</code> class).	
Classes	An instance of <code>IOADBDevice</code> matches your driver and loads it into the kernel. Your driver communicates with its family through an instance of <code>IOADBDevice</code> .	ADB bus controller drivers should inherit from the <code>IOADBController</code> class as defined in the header file <code>IOADBController.h</code>
Notes	Common client families include the HID family (<code>IOHIDPointing</code> and <code>IOHIDKeyboard</code> classes) and the Graphics family (<code>IODisplay</code> class).	All current ADB bus hardware produced by Apple is well supported by drivers that are included with OS X. Third-party developers should not need to write drivers for the ADB family, except for ADB-USB adaptors.

ATA and ATAPI

The ATA and ATAPI family provides support for ATA controllers, and access to ATA and ATAPI devices on the ATA bus.

Important: The ATA and ATAPI family is still under development. The information in this section is subject to change.

Bundle identifier:

- `com.apple.iokit.IOATAFamily`

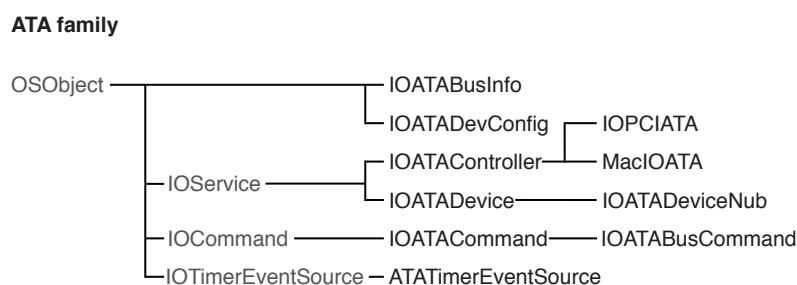
Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/ata`

References and specifications:

- American National Standards Institute (ANSI)—<http://www.ansi.org>
- National Committee for Information Technology Standards (NCITS)—<http://www.ncits.org>
- Technical Committee 13 (T13)—<http://www.t13.org>

Class hierarchy:



Device interface:

- Although this family does not itself export device interfaces, the [SCSI Architecture Model](#) (page 161) family does provide device-interface support for ATAPI devices.

Table A-2 Clients and providers of the ATA and ATAPI family

	Client of the nub	Provider for the nub
Action	Drives a device that is connected to an ATA bus.	Drives an ATA bus controller.
Example	An ATA hard drive or an ATAPI DVD-ROM drive. The Storage family is the most common family for clients.	

	Client of the nub	Provider for the nub
Classes	Clients of this family match against an instance of <code>IOATADevice</code> , one of which is created and published by the controller driver for every ATA or ATAPI device that is detected on the bus. They use the services provided by that object to communicate with the physical device on the bus.	
Notes	Clients of the ATA family must issue ATA/ATAPI commands encapsulated by an <code>IOATACCommand</code> object. This command object encapsulates all the information necessary to encode a single ATA/ATAPI command as well as the results of the command's execution.	Third-party developers should never need to create a member of the ATA/ATAPI family.

Audio

The Audio family provides support to enable access to devices that record or play back audio signals. It provides a flexible abstraction for audio devices that permits an unlimited number of channels as well as arbitrary sample rates, bit depths, and sample formats. The Audio family utilizes a high-resolution time base that is used as the basis for timing information for the entire audio and MIDI system in OS X. (The Audio family itself does not provide any MIDI services; these are provided by the Core MIDI framework.)

The Audio Hardware Abstraction Layer (Audio HAL) provides all audio services to applications in OS X. The Audio HAL is accessed through the Core Audio framework and has its programmatic interface defined in `AudioHardware.h` in that framework. The Audio family provides the link between an audio driver and the Audio HAL. Because the Audio HAL is a client of the Audio family, all audio device functionality is available to clients of the Audio HAL.

Bundle identifier:

- `com.apple.iokit.IOAudioFamily`

Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/audio/`

Device interface:

- Although this family does not directly export device interfaces, the Audio family does provide a device interface that is used by the Audio Hardware Abstraction Layer (Audio HAL) to access all of the abstractions provided by the Audio family (see description above).

Power management:

- The Audio family performs most power management tasks for subclassed device drivers. An audio driver does not have to call `PMinit`, `joinPMtree`, `registerPowerDriver`, or `PMstop`, because the Audio family takes care of initializing power management, attaching the driver into the power plane and registering it with power management, and terminating power management.

Although an audio driver does not have to implement the `IOService` method `setPowerState`, it does need to implement the `IOAudio` method `performPowerStateChange` to do the work of changing its device's power state.

- The Audio family implements idleness determination by keeping track of active audio engines, so a custom audio driver never needs to call `activityTickle` or determine idleness on its own.

Table A-3 Clients and providers of the Audio family

	Client of the nub	Provider for the nub
Action	A kernel-resident client is not necessary. Use the Core Audio framework to access Audio HAL.	Either records or plays back audio signals.
Example		PCI audio cards, external USB or FireWire audio devices and any other device that produces or consumes audio.
Classes		An audio driver must contain subclasses of both <code>IOAudioDevice</code> and <code>IOAudioEngine</code> .
Notes		An audio driver is a client of other families that provide access to the hardware that the driver supports. For example, a driver for a PCI audio card will be a client of the PCI family.

FireWire

The FireWire family provides support for, and access to, devices attached to the FireWire bus (FireWire is an Apple trademark applied to the IEEE 1394 standard, also sometimes known as i.LINK™).

The FireWire family has strong affinities with the SBP2 family. A driver that uses the SBP-2 transport protocol is the most common client of the FireWire family.

Bundle identifier:

- `com.apple.iokit.IOFireWireFamily`

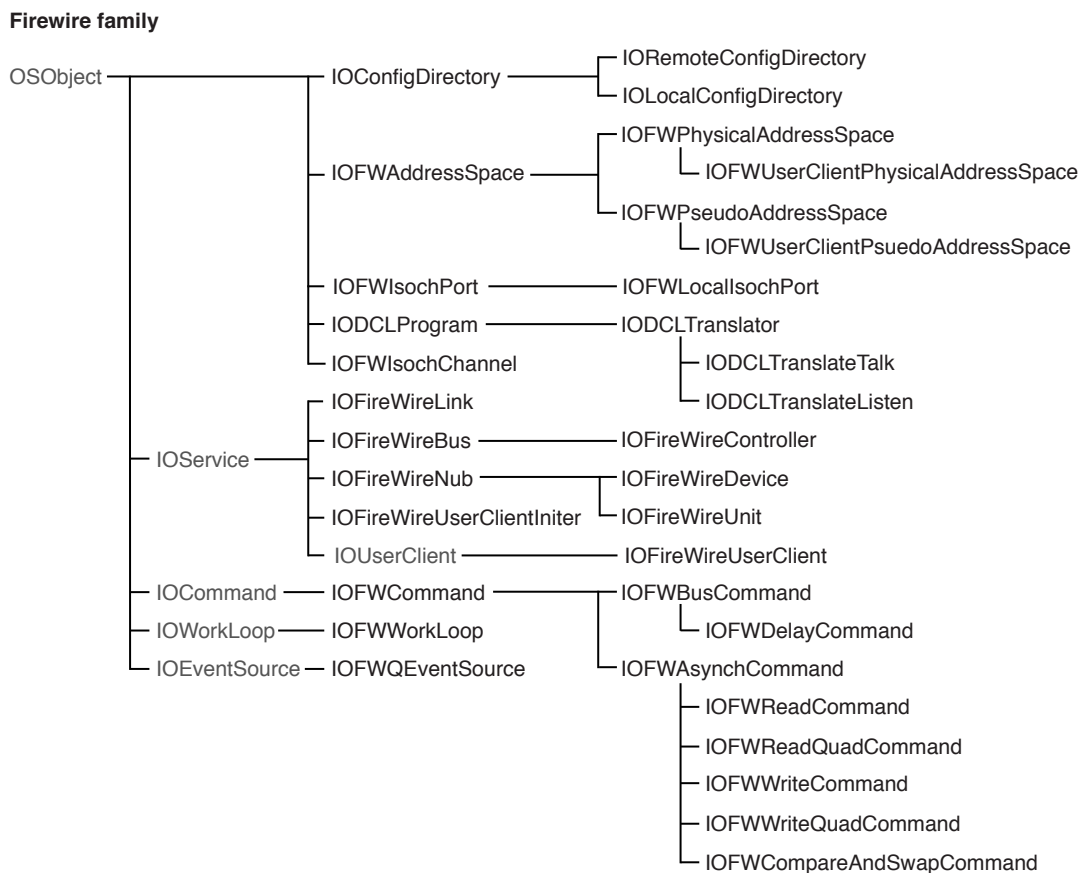
Headers in:

- Kernel-resident: `Kernel.framework/Headers/IOKit/firewire/`
- Device interface: `IOKit.framework/Headers/firewire`

References and specifications:

- 1394 Trade Association—<http://www.1394ta.org>
- IEEE—<http://www.ieee.org/portal/site>. Copies of the IEEE 1394 specification can be purchased here.
- Apple Developer Connection—<http://developer.apple.com/hardwaredrivers/firewire/index.html>

Class hierarchy



Device interface:

- Provides a device interface exporting an interface for sending and receiving packets on the FireWire bus and for adding entries into the computer's own FireWire config ROM.

Table A-4 Clients and providers of the FireWire family

	Client of the nub	Provider for the nub
Action	<i>Unit drivers:</i> Drives communication with a unit of a device that plugs into the FireWire bus. <i>Protocol drivers:</i> Adds support for a protocol enabling peer-to-peer communication or emulation over FireWire.	Drives a FireWire bus controller.
Example	<i>Unit drivers:</i> A driver for a FireWire speakers. <i>Protocol drivers:</i> A driver that provides TCP/IP over FireWire.	
Classes	<i>Unit drivers:</i> An instance of IOFireWireUnit, representing a unit found in a device's configuration ROM, is matched against the driver and it is loaded into the kernel. The driver communicates with the FireWire family through this instance. <i>Protocol drivers:</i> An instance of IOFireWireController matches a protocol driver and loads it into the kernel. The driver communicates with the FireWire family through this instance.	Driver classes should inherit from the IOFireWireController class as defined in the header file <code>IOFireWireController.h</code> .
Notes	The most common client family is the SBP2 family.	OS X ships drivers for all Open Host Controller Interface (OHCI) bus controllers. Third-party developers generally should not need to write bus-controller drivers for the FireWire family.

In some cases, you can write a driver for a FireWire device instead of for a unit. An example might be a driver for a device with a minimal config ROM (that is, with just a vendor ID). However, use of the minimal config ROM is strongly discouraged by Apple. Also, if your driver matches against a FireWire unit, it is often possible to do some things with the device.

Graphics

The Graphics family provides support for frame buffers and display devices (monitors).

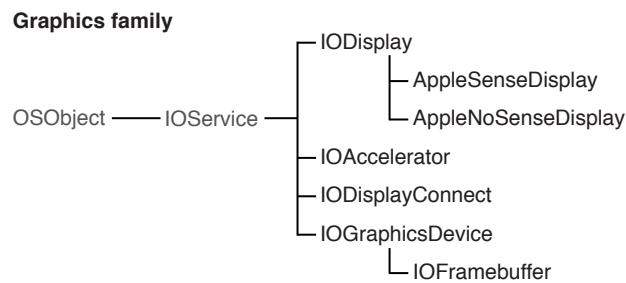
Bundle identifier:

- `com.apple.iokit.IOGraphicsFamily`

Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/graphics/`
- Device interface: `IOKit.framework/Headers/graphics`

Class hierarchy:



Device interface:

- The Graphics family exports several device interfaces since most access of graphics devices is from user space. However, the Quartz layer arbitrates access to frame buffers from user space through the windowing system or the CGDirectDisplay API. Other layers, such as Carbon Draw Sprockets, provide application access to graphics.
- Graphics acceleration is supplied by modules loaded into user address space. A CFPlugIn interface, defined in `IOGraphicsInterface.h`, implements two-dimensional acceleration. Similarly, OpenGL defines a loadable-bundle interface for three-dimensional rendering. Because there is no standard way to implement this functionality, hardware-specific code can exist in both user-space code and in a kernel-loaded driver.

Table A-5 Clients and providers of the Graphics family

	Client of the nub	Provider for the nub
Action		Implements support for a frame buffer.
Example		
Classes		Frame-buffer drivers must be subclasses of the <code>IOFramebuffer</code> class. The <code>IONDRVFramebuffer</code> class supports native Power PC Mac OS graphics drivers (known as “ndrv”s); this support is automatic, provided the drivers are written correctly to the specification.

	Client of the nub	Provider for the nub
Notes	Support for kernel-resident clients is limited. Because the Quartz layer owns the display, the kernel generally does not render graphics directly.	Apple provides generic support for displays and so displays should not generally require third-party drivers.

A Note on NDRV Compatibility

NDRV graphics drivers should function in OS X if they are correctly written. If they are not correctly written, the many differences in OS X's runtime environment could cause them to fail, be ignored, or even cause a crash. If you are writing an NDRV driver, follow these rules:

- Access the card hardware using virtual addressing. Do not assume the card is mapped into its physically assigned address. In Mac OS 9, NDRV cards are mapped one-to-one, but in OS X, this is not guaranteed. Obtain the virtual addresses for your card's hardware via the `AAPL, address` property as documented in "Designed Cards and Drivers for PCI Power Macintosh."
- Link only on the native driver libraries, which are `NameRegistryLib`, `DriverServicesLib`, and `VideoServicesLib`. If your card links on `InterfaceLib` or any other application-level library, it probably won't work on OS X.
- Do not access low memory; doing so causes a crash (kernel panic) in OS X.
- Name registry calls are not supported from interrupt level in Mac OS 9 or OS X. They return errors in OS X.
- Secondary interrupts are not supported in OS X. There is no need to fake vertical blank interrupts if your card does not support them—simply do not create a VBL service. Mac OS 9 continues to require a VBL service to be installed to move the cursor on your device.
- Stack size is limited to 16K on OS X; any NDRV invocation should consume no more than 4K of stack.

If you want to make runtime conditional changes to your NDRV code, the property `AAPL, iokit-ndrv` is set in the PCI device properties before OS X uses your driver.

OS X supports 32 bits-per-pixel, alpha-blended cursors in hardware. If your device supports an alpha-blended direct color cursor, it should call `VSLPrepareCursorForHardwareCursor` with these fields set in the `HardwareCursorDescriptor` record:

```
bitDepth = 32 maskBitDepth = 0 numColors = 0 colorEncodings = NULL
```

The `hardwareCursorData` buffer in the `HardwareCursorInfo` should point to a buffer of 32 bits per pixel, ARGB data. The data is not premultiplied by the alpha channel.

HID

The Human Interface Device (HID) class is one of several device classes described by the USB (Universal Serial Bus) architecture. The HID class consists primarily of devices humans use to control a computer system's operations. Examples of such HID class devices include:

- Keyboards and pointing devices such as mice, trackballs, and joysticks
- Front-panel controls such as knobs, switches, sliders, and buttons
- Controls that might be found on games or simulation devices such as data gloves, throttles, and steering wheels

Currently, OS X provides the HID Manager to allow applications to access joysticks, audio devices, and non-Apple displays. You can also use the HID Manager to get information from another type of HID class device, a UPS (uninterruptible power supply) device. UPS devices share the same report descriptor structure as other HID class devices and provide information such as voltage, current, and frequency. The OS X HID Manager consists of three layers:

- The HID Manager client API that provides definitions and functions your application can use to work with HID class devices
- the HID family that provides the in-kernel HID infrastructure such as the base classes, the kernel-user space memory mapping and queueing code, and the HID parser
- the HID drivers provided by Apple

Bundle identifier:

- `com.apple.iokit.IOHIDFamily`


Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/hid/` and `Kernel.framework/Headers/IOKit/hidsystem/`
- Device interface: `IOKit.framework/Headers/hid/` and `IOKit.framework/Headers/hidsystem/`

References and specifications:

- HID Information section of USB.org—Developers website (<http://www.usb.org/developers/hidpage>)
- *HID Class Device Interface Guide*

Class hierarchy:



Device interface:

- The HID family exports a device interface through the HID Manager client API. The HID Manager includes `IOHIDLib.h` and `IOHIDKeys.h` (located in `/System/Library/Frameworks/IOKit.framework/Headers/hid`) which define the property keys that describe a device, the element keys that describe a device's elements, and the device interface functions and data structures you use to communicate with a device. After you've created a device interface for a selected HID class device, you can use the device interface functions to open and close the device, get the most recent value of an element, or set an element value.

Table A-6 Clients and providers of the HID family

	Client of the nub	Provider for the nub
Action		Drives an input device such as a multi-button mouse, trackball, or joystick.
Example		
Classes		
Notes		Support for most simple input devices is provided by the generic driver.

Network

The Network family provides support for network controllers. The Network family consists of two logical layers:

- Controller layer—this layer represents the network controller.
- Interface layer—this layer represents the network interface published by the network controller.

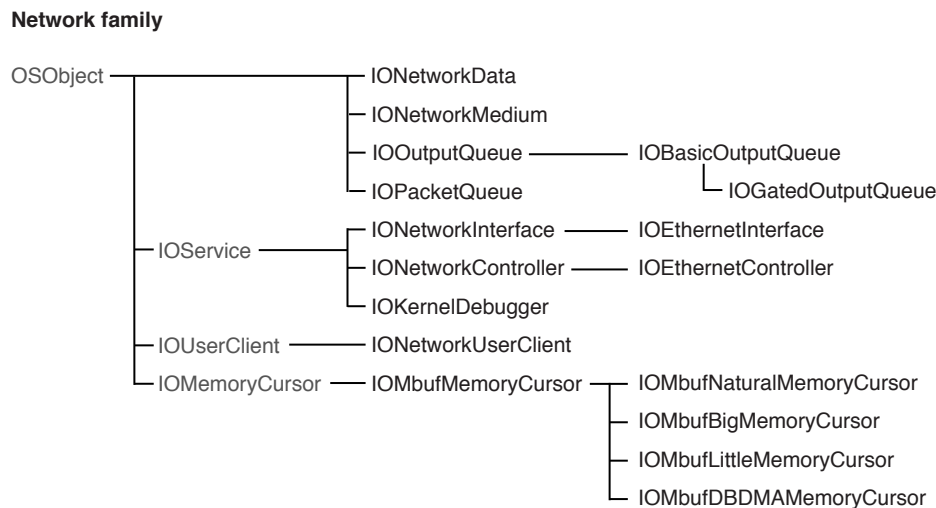
Bundle identifier:

- `com.apple.iokit.IONetworkingFamily`

Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/network/`
- Device interface: `IOKit.framework/Headers/network/`

Class hierarchy:



Device interface:

- The device interface for this family is usually the BSD network stack. Applications use the socket interface provided by the network stack to access indirectly the services provided by the Network family.

Power management:

The Network family performs most of the power-management set-up and tear-down tasks for subclassed device drivers. If you're developing a driver for a network device that can be passively power managed (which describes most network devices), you can meet most of your basic power-management needs by overriding the `IONetworkController` method `registerWithPolicyMaker` and calling the `IOService` method `registerPowerDriver`.

In your implementation of `registerWithPolicyMaker`, create an array of `IOPMPowerState` structures to define your device's power states and pass them in to `registerPowerDriver`. Then, return `kIOReturnSuccess` from `registerWithPolicyMaker` to tell the Network family that your driver can respond to power-management calls. (The default implementation of `registerPowerDriver` returns `kIOReturnUnsupported`.) The following code snippet shows one way to do this:

```
IOReturn MyEthernetDriver::registerWithPolicyMaker( IOService * policyMaker )
{
    IOReturn ioreturn;
    static IOPMPowerState powerStateArray[ kPowerStateCount ] = {
        { 1,0,0,0,0,0,0,0,0,0,0,0 },
        { 1,kIOPMDeviceUsable,kIOPMPowerOn,kIOPMPowerOn,0,0,0,0,0,0,0,0 }
    };
    fCurrentPowerState = kPowerStateOn;
    ioreturn = policyMaker->registerPowerDriver( this, powerStateArray, kPowerStateCount );
    return ioreturn;
}
```

Most network device drivers handle power changes related to sleep and wake in their implementations of the `IONetworkController` methods `enable` and `disable`. Note that the Network family enables a device when it transitions to a power state for which the `kIOPMDeviceUsable` flag is set. When a currently enabled device moves to a power state for which the `kIOPMDeviceUsable` flag is not set, the Network family disables it.

If you need to perform additional tasks to handle sleep and wake, you can override the `IOService` method `setPowerState`. Be aware, however, that the Network family will call `disable` *before* you receive a call to your `setPowerState` implementation if the new power state puts the device into an unusable state. Conversely, the Network family calls `enable` *after* you receive a `setPowerState` call to move the device to a usable state.

If your network device driver performs DMA, you should override the `IOService` method `systemWillShutdown`, which was introduced in OS X v10.5. This is especially important for drivers that run in Intel-based Macintosh computers. In your implementation of `systemWillShutdown`, you should make sure that the DMA engine is shut off, which results in the necessary disabling of the port.

Important: As described in [Receiving Shutdown and Restart Notifications](#) (page 127) the `systemWillShutdown` call is made to drivers in the power plane, in leaf-to-root order. If your driver returns `kIOReturnUnsupported` from `registerWithPolicyMaker`, it will not be attached to the power plane and will not receive a `systemWillShutdown` call.

Table A-7 Clients and providers of the Network family

	Client of the nub	Provider for the nub
Action		Drives a network controller.

	Client of the nub	Provider for the nub
Example		Controllers on Ethernet, Token Ring, and FDDI adapters.
Classes		Driver must be an instance of a subclass of a controller class that implements generic network controller functionality, such as <code>IONetworkController</code> or of a controller class that builds upon <code>IONetworkController</code> to specialize for Ethernet controller support (<code>IOEthernetController</code>). See discussion on Network family classes below for more information.
Notes	Drivers are typically not clients of the Network family. The primary system client of this family is the DLIL (Data Link Interface Layer) module in the BSD network stack.	

Member drivers must also create `IONetworkInterface` objects that are registered with the DLIL; such registration associates the driver with a network interface (for example, `en0`) in the system. You can create a Network Kernel Extension (NKE) and insert it at various locations above the IOKit/DLIL boundary to intercept the packets, commands, or other event traffic between an `IONetworkInterface` object and the upper layers.

Another client of this family is the KDP (Kernel Debugger Protocol) module in the kernel. A driver can create an `IOKernelDebugger` object that vends debugging services and allows kernel debugging through the network hardware managed by the driver. Only drivers that drive a built-in network controller are required to provide this support.

Other classes of the Network family include:

`IOOutputQueue`—assists in queuing outbound packets.

`IOPacketQueue`—represents a generic `mbuf` packet FIFO queue.

`IOMbufMemoryCursor`—translates `mbuf` packets into a scatter-gather list of physical addresses and length pairs.

`IONetworkData`—represents a container for a single group of statistics counters.

`IONetworkMedium`—represents a single network medium supported by the network device.

PC Card

The PC Card family supports 32-bit PC cards (CardBus), 16-bit PC cards (I/O and memory), and Zoom Video cards. This support encompasses controllers that are compatible with ExCA (Intel 82365) and Yenta register sets. Apple's PC Card family's Card Services are, for the most part, compliant with the 1997 PC Card™ standard.

CardBus cards are essentially PCI devices in a different form factor. If you are writing a driver for a CardBus card, you can choose to subclass from either `IOPCIDevice` or `IOCardBusDevice`. See the reference section “[PCI and AGP](#)” (page 157) for more information about the PCI family.

Other classes provided by the family include:

`IOPCCard16Device`—represents a 16-bit PC card device.

`IOPCCard16Enabler`—allows you to override the card configuration process for 16-bit cards. It is used mainly for cards with broken or missing CIS entries.

`IOZoomVideoDevice`—represents a Zoom Video device.

`IOPCCardBridge`—represents a PC Card bridge; this class is a subclass of `IOPCI2PCIBridge` which is a subclass of `IOPCIBridge`.

Bundle identifier:

- `com.apple.iokit.IOPCCardFamily`

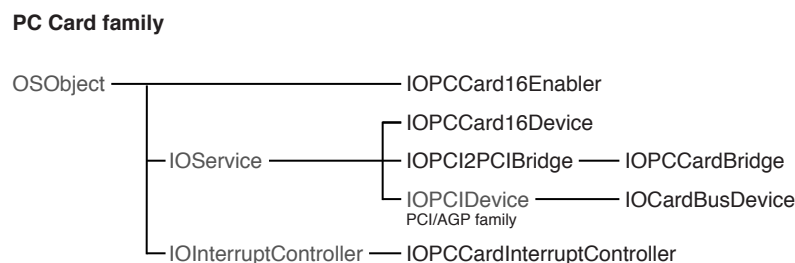
Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/pccard/`

References and specifications:

- Documentation for Card Services can be found in the `doc` directory of the PC Card family source (available via Darwin CVS). You can also find in the same location a sample 16-bit driver.
- Apple Developer Connection—<http://developer.apple.com/hardwaredrivers/pci/>

Class hierarchy:



Device interface:

- None.

Power management:

- Classes `IOPCCardBridge`, `IOPCCard16Device`, `IOPCCard16Enabler`, and `IOPCardBusDevice` provide some power-management support.

PCI and AGP

The PCI and AGP family provides support for, and access to, devices attached to PCI and AGP buses and PCI bridges.

Bundle identifier:

- `com.apple.iokit.IOPCIFamily`


Headers in:

- `Kernel.framework/Headers/IOKit/pci/`

References and specifications:

- This family supports all major features of the PCI Localbus 2.1 specification.
- PCI Special Interest Group—<http://www.pcisig.com>
- Apple Developer Connection—<http://developer.apple.com/hardwaredrivers/pci/>

Class hierarchy:



Device interface:

- None. Applications are not permitted direct access to PCI bus hardware for security reasons. Instead, applications should interact with higher-level services, such as those provided by device interfaces of the USB or other families.

Matching properties:

- The PCI/AGP family permits matching based on Open Firmware or on PCI registers. See the description of the `IOPCIDevice` class in the reference documentation for details.

Check the Darwin Open Source project for example PCI drivers.

Table A-8 Clients and providers of the PCI and AGP family

	Client of the nub	Provider for the nub
Action	Drives a device that plugs into a PCI bus.	Drives a PCI bus controller or a PCI bridge.
Example	A driver for a PCI SCSI controller card is a <i>client</i> of the PCI family but is a <i>member</i> of the SCSI Parallel family.	
Classes	The driver communicates with its family via an instance of <code>IOPCIDevice</code> or <code>IOAGPDevice</code> . An instance of one of these classes matches your driver and loads it into the kernel.	PCI family member drivers should inherit from the <code>IOPCIBridge</code> class.
Notes	The most common client families are the USB, Network, SCSI, Graphics, and Audio families.	OS X supports most PCI bus hardware with a set of generic drivers. In general, third-party developers do not need to write drivers for the PCI and AGP family unless they are building a PCI expansion chassis or developing drivers for a PCI bridge with special characteristics not addressed by the generic drivers.

SBP-2

The SBP2 family (Serial Bus Protocol 2) provides support for, and access to, devices attached to the FireWire bus that use the SBP-2 transport protocol. The SBP2 family is a client of the FireWire family. SBP-2 devices require FireWire to run.

Bundle identifier:

- `com.apple.iokit.IOFireWireSBP2`

Headers in:

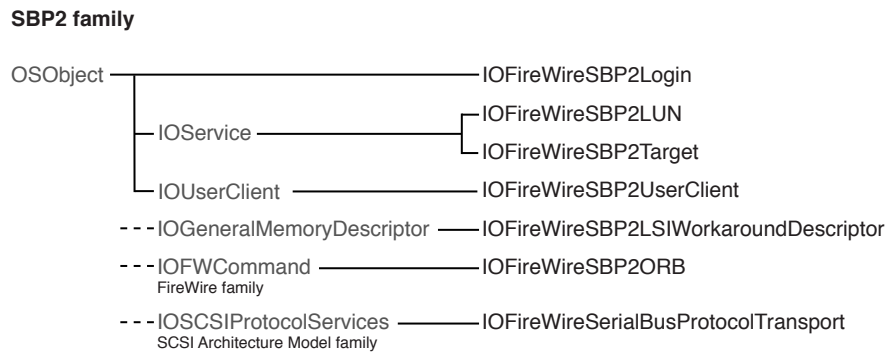
- Kernel resident: `Kernel.framework/Headers/IOKit/sbp2/`

- Device interface: `IOKit.framework/Headers/sbp2`

References and specifications:

- T10 Technical Committee — <http://www.t10.org>
- SBP-2 standard: <ftp://ftp.t10.org/t10/drafts/sbp2/>

Class hierarchy:



Device interface:

- The SBP2 family provides a device interface that exports an interface for sending SBP-2 ORBs to a device.

Table A-9 Clients and providers of the SBP2 family

	Client of the nub	Provider for the nub
Action	Drives a device that uses the SBP-2 transport protocol.	Provides SBP-2 transport services
Example	A driver for a typical FireWire hard disk drive is a <i>client</i> of the SBP2 family but is a <i>member</i> of the mass storage family.	
Classes	A client driver communicates with the SBP2 family through an instance of <code>IOFireWireSBP2LUN</code> . An instance of this class is created for each LUN (Logical Unit Number) found in a configuration ROM unit directory; this instance matches your driver and loads it into the kernel.	SBP2 family-member drivers should inherit from the <code>IOFireWireSBP2Target</code> class.
Notes	The most common client family is the Storage family.	

SCSI Parallel

The SCSI Parallel family provides support for, and access to, devices attached to a parallel SCSI bus. This family supports all major features of the SCSI Parallel Interface-5 (SPI-5) specification for parallel buses.

Bundle identifier:

- `com.apple.iokit.IOSCSIParallelFamily`

Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/scsi/spi`. The header file `IOSCSIParallelInterfaceController.h` supports the development of a driver for a SCSI card.

References and specifications:

- T10 Technical Committee — <http://www.t10.org>
- SCSI Technical Library of Information — <http://www.scsilibrary.com/>

Class hierarchy:



Device interface:

- Device interface support for parallel SCSI devices is provided by the SCSI Architecture Model family. See [SCSI Architecture Model](#) (page 161)

Table A-10 Clients and providers of the SCSI Parallel family

	Client of the nub	Provider for the nub
Action	Drives a device that plugs into a SCSI bus.	Drives a SCSI host adapter or controller chip.
Example	A driver for a SCSI disk drive is a <i>client</i> of the SCSI Parallel family but a <i>member</i> of the Storage family.	

	Client of the nub	Provider for the nub
Classes	A client driver communicates with the SCSI Parallel family through an instance of <code>IOSCSIDevice</code> . Your driver would match on this instance and be loaded into the kernel. Another class of interest is <code>IOSCSICommand</code> .	Because the SCSI Parallel family presently supports parallel buses, member drivers use the <code>IOSCSIParallelInterfaceController</code> class.
Notes	Common client families include the Transport Drivers for storage devices.	OS X includes generic drivers for most built-in SCSI hardware. In general, third-party developers do not need to write drivers that are members of the SCSI Parallel family unless they are developing drivers for an expansion card.

SCSI Architecture Model

The SCSI Architecture Model family provides common client support for SCSI, USB (Storage), FireWire SBP-2 and ATAPI devices. Many of the classes of this family belong to the Transport Driver layer, which is summarized in [Table A-11](#) (page 161)

Table A-11 SCSI Architecture Model family—Transport Driver layer

Type	Classes	Comments
Device Services Linkage	<code>IOBlockStorageServices</code> <code>IOReducedBlockServices</code> <code>IOCompactDiscServices</code> <code>IODVDServices</code>	Linkage objects that understand the APIs from the Device Services layer (Storage family) and export APIs that the Logical Unit drivers (in the Transport Driver layer) understand.
Logical Unit drivers	<code>IOSCSIPeripheralDeviceType00</code> <code>IOSCSIPeripheralDeviceType05</code> <code>IOSCSIPeripheralDeviceType07</code> <code>IOSCSIPeripheralDeviceType0E</code>	Drive mass storage devices that have file systems or are bootable.
Peripheral Device Type Nub	<code>IOSCSIPeripheralDeviceNub</code>	Not really an I/O Kit nub, but an object that queries the device and determines which Logical Unit driver is needed.

Type	Classes	Comments
SCSI Protocol drivers	IOUSBMassStorageClass IOATAPIProtocolTransport IOFireWireSerialBusProtocol- Transport IOCSIParallelInterface- ProtocolTransport	Classes for bus-specific drivers. Although these classes belong to other families, they are part of the SCSI Architecture Model layering.

The Logical Unit drivers and the Peripheral Device Type Nub are in the SCSI Application layer, which inherits (ultimately) from IOCSIPrimaryCommandsDevice. The SCSI Protocol drivers are in the SCSI Protocol layer, which inherits from IOCSIProtocolServices.

The SCSI Architecture Model family supports multimedia command set devices, such as CD-RW drives.

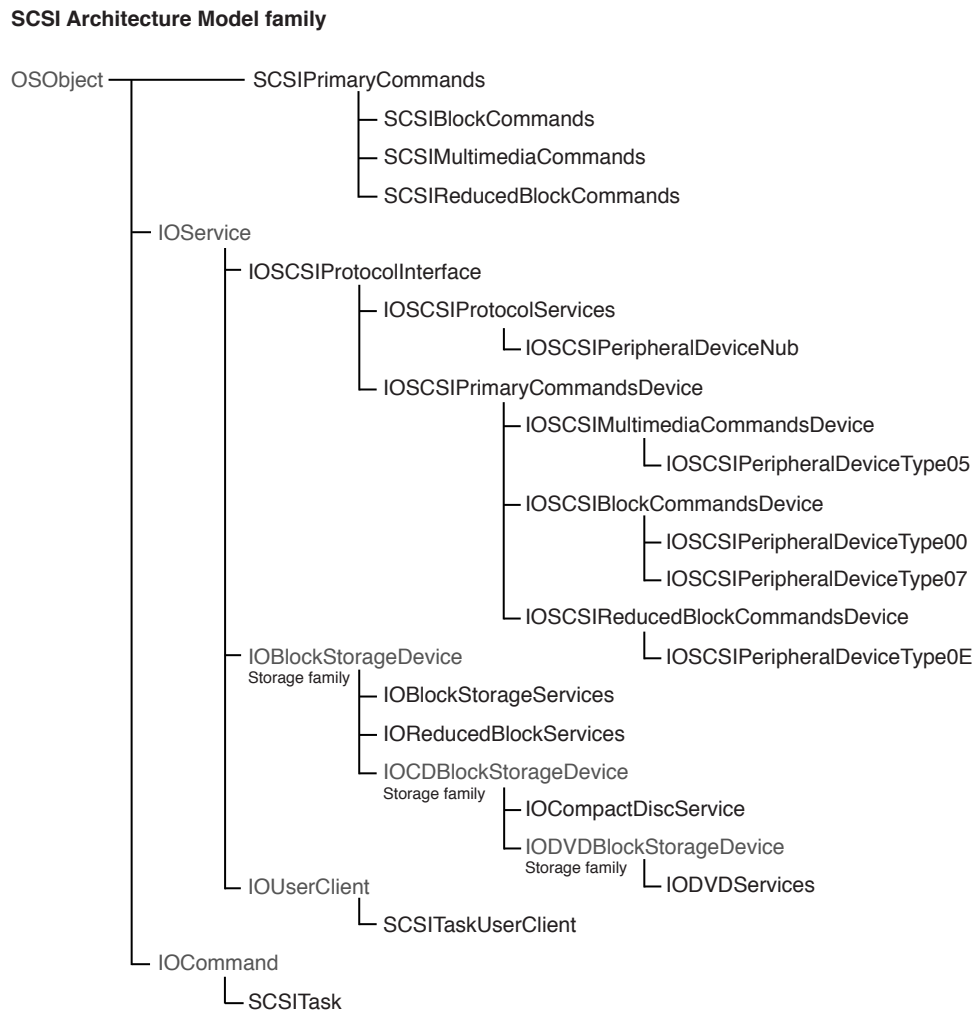
Bundle identifier:

- `com.apple.iokit.IOCSISArchitectureModelFamily`

Headers in:

- Kernel-resident: `Kernel.framework/Headers/IOKit/scsi-commands/`
- Device interface: `IOKit.framework/Headers/scsi-commands`

Class hierarchy:



Device interface:

- The library for the SCSI Architecture Model family is called `SCSITaskLib`. It includes three interfaces: `MMCDeviceInterface`, `SCSITaskDeviceInterface`, and `SCSITaskInterface`. The user-client class is `SCSITaskUserClient`. If you need to access a SCSI Parallel device *and* your application must run in versions of OS X prior to v10.2, see [Accessing SCSI Parallel Devices](#) for information on how to do this. Otherwise, you should use the device interfaces in the `SCSITaskLib` to access your device (see [Accessing SCSI Architecture Model Devices](#) for information on how to do this).

Power management:

The SCSI Architecture Model family performs most of the power-management set-up and tear-down tasks for both protocol services drivers and logical unit drivers. In general, a protocol services driver, such as `IOATAPIProtocolTransport`, must be able to transition the physical interconnect device between the off and

on states. On the other hand a logical unit driver, such as `IOCSIPeripheralDeviceType05`, must be able to manage a multimedia device that supports all the power states defined by the SCSI multimedia commands specification. In addition, a logical unit driver might need to determine if a power-state change is needed, block incoming I/O when the device is not in an appropriate power state, and specify the power state a device should enter at boot time.

Note: The SCSI Architecture Model family defines the system sleep power state, which is in addition to the power states defined by the command-set specifications. System sleep corresponds to the sleep that occurs when the user chooses Sleep in the Apple menu or closes the lid of a laptop. The sleep state defined in the command-set specifications corresponds to the sleep that occurs when the device is idle. Because power can be removed from devices in system sleep, the SCSI Architecture Model family handles it differently than sleep.

As shown in the class hierarchy diagram above, the SCSI Architecture Model family defines a common superclass for both types of drivers: `IOCSIProtocolInterface`. The `IOCSIProtocolInterface` class defines a number of power-management methods that subclass drivers can call or, less frequently, override.

- If you're developing a custom protocol services driver, you do not need to perform the steps outlined in [Implementing Basic Power Management](#) (page 113). Instead, you must call the `IOCSIProtocolInterface` method `InitializePowerManagement` in your driver's start routine. Then, if there is device-specific work to do to handle power-state changes, you can implement the methods `HandlePowerOn` and `HandlePowerOff`.
- Similarly, if you're developing a custom logical unit driver, you do not need to perform the steps outlined in [Implementing Basic Power Management](#) (page 113). If your device complies with the appropriate command-set specification, you do not need to override any methods in your driver unless you want to implement custom power-management functionality. For example, you might want your device to transition from the active state directly to the sleep state, instead of through the intermediate states between active and sleep.

In the `IOCSIProtocolInterface` class, the SCSI Architecture Model family provides the following methods a logical unit driver subclass can implement:

- `InitializePowerManagement`. The superclass implementation of this method performs power-management setup tasks. A subclass driver can override this method to provide information about the power states the device supports.
- `TicklePowerManager`. The superclass implementation of `TicklePowerManager` calls the `activityTickle` method, which results in a request to change your device to its active power state. A subclass driver can override this method to specify a state different from the active one.

- `GetInitialPowerState`. This method is used to specify the default state (usually active) a device should be in when the system boots. If a device should enter a different state when the system boots, the subclass driver can override this method and specify that state.
- `GetNumberOfPowerStateTransitions`. A subclass driver can override this method to report the number of transitions between the lowest and highest specification-defined power states the device supports. Note that system sleep is not counted in the transitions because it is not a power state the device enters voluntarily.
- `HandlePowerChange`. A subclass driver overrides this method to perform the work of changing the device's power state.

Serial

The Serial family provides support for serial byte character streams.

Bundle identifier:

- `com.apple.iokit.IOSerialFamily`


Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/serial/`
- Device interface: `IOKit.framework/Headers/serial/`

References and specifications:

- See `termios`. Also see the related header file `Kernel.framework/Headers/sys/termios.h`

Class hierarchy:



Device interface:

- Applications can access the Serial family through the BSD device nodes, the most common client of this family. An application can read and write data using the BSD device nodes in `/dev`. Data is also routed through to PPP via these device nodes. You can find keys and other properties for use in device access in `IOKit.framework/Headers/serial/IOSerialKeys.h`.

Table A-12 Clients and providers of the Serial family

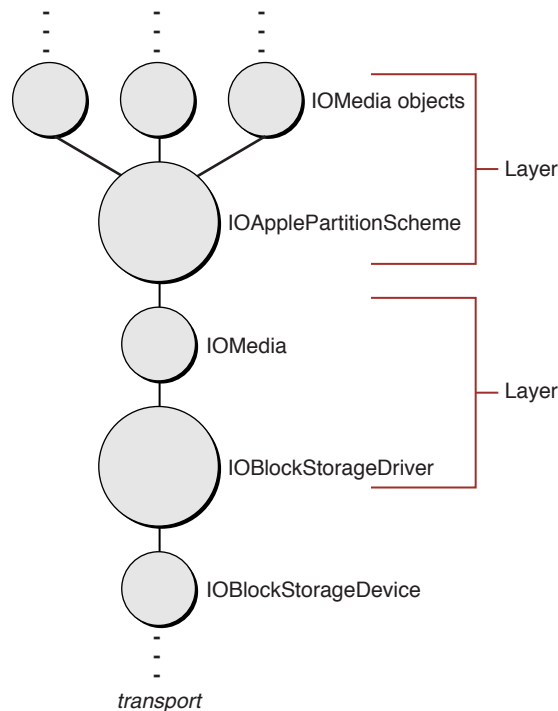
	Client of the nub	Provider for the nub
Action	<i>Requires</i> a single-banded data streaming service with elementary flow control.	<i>Provides</i> a single-banded streaming service; in other words, it cannot be packet-based. although it may be bi-directional. The driver may also implement flow control. The driver must describe the services it is capable of providing.
Classes		Serial port writers should subclass <code>IOSerialDriverSync</code> for their drivers, then publish as nubs objects of the <code>IOModemSerialStreamSync</code> or <code>IORS232SerialStreamSync</code> classes, or (if neither of these suffices) a concrete subclass of <code>IOSerialStreamSync</code> . The I/O Kit uses these objects to create the appropriate user-client interface for user-space access via BSD.
Notes	Developers should use the BSD device file mechanism documented in <i>Accessing Hardware From Applications</i> .	

Storage

The Storage family provides high-level support for random-access mass storage devices. It is separate from the underlying technology that transports the data to and from the represented storage space. The interface to the underlying transport technology is declared in the abstract class `IOBlockStorageDevice`. Storage driver objects communicate all mass-storage requests across this interface, without having to have knowledge of, or involvement with, the commands and mechanisms used to communicate with the device or bus.

The scope of the Storage Family encompasses the IOBlockStorageDevice interface, at one end (provider direction), and the BSD interface at the other end (client direction), with various driver and media layers in the between. [Figure A-1](#) (page 167) illustrates this stack.

Figure A-1 Storage family driver stack



Each layer consists of a set of two objects—a driver object and the child media object (or objects) it publishes. The IOStorage class is the common base class for both driver and media objects. It is an abstract class that declares the basic open, close, read, and write interfaces that subclasses are to implement. It establishes the protocol with which media objects can talk to driver objects without needing to be subclassed for each driver. The read and write interfaces provide byte-level access to the storage space.

The IOBlockStorageDriver class is the common base class for generic block storage drivers. It matches and communicates via an IOBlockStorageDevice interface, and connects to the remainder of the storage framework via the IOStorage protocol. It extends the IOStorage protocol by implementing the appropriate open and close semantics, deblocking for unaligned transfers, polling for ejectable media, implementing locking and ejection policies, creating and tearing down media object, and gathering and reporting statistics. The Storage family supports other basic types of drives, such as CD drives and DVD drives, through subclasses of the IOBlockStorageDriver. You rarely, if ever, need to subclass the generic block storage driver to handle device idiosyncrasies; rather, you should change the underlying transport drivers to correct any non-conforming behavior.

The `IOMedia` class is an abstraction of a random-access disk device. It is equivalent to the “device object” or “interface object” in other I/O Kit families.

- It represents the presence of a device, or a piece thereof.
- It presents an programmatic interface to that device.
- It is backed by a separate driver that implements the required functionality.

`IOMedia` provides a consistent interface for both real and virtual disk devices, for subdivisions of disks such as partitions, for supersets of disks such as RAID volumes, and so on. It extends the `IOStorage` class by implementing the appropriate open, close, read, write, and matching semantics for media objects. Its properties reflect the properties of real disk devices, including natural block size and writability.

The other driver and media layers in the Storage Family are known as filter schemes. These optional layers separate media objects, providing some kind of data manipulation or offset manipulation between media objects. These layers may stack arbitrarily on top one another, as they are both a client and a provider of media objects. Most third-parties will develop drivers for the filter scheme layer, if not for the underlying transport technology. Refer to [IOMedia Filter Schemes](#) (page 169) for more information on developing filter scheme drivers.

For more information on developing drivers for the underlying transport technology, refer to SCSI Architecture Model Family documentation. The SCSI Architecture Model Family is the common underlying transport technology for ATAPI, FireWire, SCSI, and USB. It provides a consistent CDB-based access model for application writers and driver writers, and a simple infrastructure for correcting device idiosyncrasies.

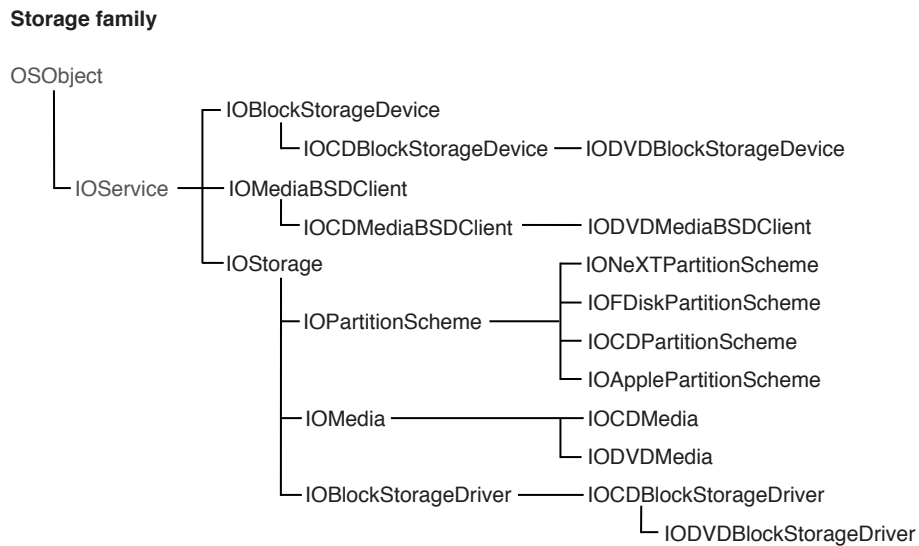
Bundle identifier:

- `com.apple.iokit.IOStorageFamily`

Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/storage/`
- Device interface: `IOKit.framework/Headers/storage`

Class hierarchy:



Device interface:

- See [Accessing IOMedia From Applications](#) (page 171)

IOMedia Filter Schemes

A filter scheme is a driver for IOMedia objects. It acts as both as a client and a provider of media objects. The filter-scheme driver receives mass storage requests through its abstract `read` and `write` member-function interfaces, in which it can perform the data manipulation or offset manipulation before passing the request on to its provider IOMedia object (or objects). There are several different kinds of media filter schemes:

- One-to-one—A block-level compression or encryption scheme, for example, would match against one IOMedia object and produce one child IOMedia object representing the uncompressed or unencrypted content.
- One-to-many—A partition scheme, for example, would match against one IOMedia object and produce multiple child IOMedia objects representing the content of each distinct partition (see [Partition Schemes](#) (page 170) for more information).
- Many-to-one—A RAID scheme, for example, would match against multiple IOMedia objects and produce one child IOMedia object representing the RAIDed content.
- Many-to-many—A many-to-many scheme would match against multiple IOMedia objects and produce multiple child IOMedia objects.

A filter-scheme driver inherits from the IOStorage class and necessarily participates in the IOStorage match category (IOMatchCategory's value in personality).

Partition Schemes

OS X includes two standard partition schemes:

- IOApplePartitionScheme, the standard Apple partition scheme driver
- IOFDiskPartitionScheme, the standard PC partition scheme driver

A partition-scheme driver inherits from the IOPartitionScheme class (which inherits from IOStorage), matches against a single IOMedia parent, and produces one or more IOMedia child objects for each partition it must represent. It participates in the IOStorage match category, as with all filter scheme drivers.

A partition scheme need not be subclassed in order to make use of developer-defined content within a partition. A partition's contents is represented by a distinct IOMedia object, published as a child of the partition-scheme driver. Each child media object has properties that further identify information known about the partition, such the content hint, size, and natural block size. The content hint field is a string formed similarly to the well-known "Apple_Driver" and "Apple_HFS" strings, or by definition, in the form "MyCompany_MyContent". It permits partitions with developer-defined contents to be identified uniquely (when the partition is created), and permits filter-scheme drivers to match against such content automatically without ever probing a disk.

IOMedia Properties

[Table A-13](#) (page 170) lists the standard set of properties for all IOMedia objects. These properties can be used as matching properties in I/O Kit's search and notification APIs, as well as for filter scheme driver matching purposes.

Table A-13 Storage family (IOMedia) properties

Key	Type	Description
kIOMediaEjectableKey	Boolean	Is the media ejectable?
kIOMediaLeafKey	Boolean	Is the media a leaf in the media tree? This is false whenever a client filter-scheme driver has matched against the media object.
kIOMediaPreferred-BlockSizeKey	Number	The media's natural block size in bytes.
kIOMediaSizeKey	Number	The media's entire size in bytes.
kIOMediaWholeKey	Boolean	Is the media at the root of the media tree? This is true for the physical media representation, a RAID media representation, and similar representations.

Key	Type	Description
kIOMediaWritableKey	Boolean	Is the media writable?
kIOMediaContentKey	String	The media's content description, as forced upon by the client filter-scheme driver. (This content description is copied automatically from the client filter-scheme driver's <code>kIOMediaContentMaskKey</code> property.) The string's format follows the "MyCompany_MyContent" convention, and defaults to the Content Hint string should no client filter scheme have matched against the object. Used for informational purposes in user disk utilities.
kIOMediaContent-HintKey	String	The media's content description, as hinted at the time of the object's creation. The string's format follows the "MyCompany_MyContent" convention. Used for matching purposes in filter scheme drivers.
kIOBSDNameKey	String	The media's BSD device node name. The name is dynamically assigned at the time of the object's creation. Used for read and write access to the media's contents (see Accessing IOMedia From Applications (page 171)).

A media object also has a unique I/O Kit path, which can be obtained via standard I/O Kit APIs.

Accessing IOMedia From Applications

The standard user-space mechanism for accessing data on a piece of media is the BSD device interface. The BSD device interface is abstracted at the file-system layer through distinct read and write APIs, while general user application access is provided via the `read` and `write` system calls (see the corresponding man pages for more documentation). The properties and structure of a physical disk are represented in the I/O Kit Registry object hierarchy and in each media object's properties, including the BSD device interface name assigned to the media object.

A specific media object can be found via the standard I/O Kit search and notification APIs. The dictionary used to describe the IOMedia might refer to specific property values, to a specific service path or device tree path, or to a specific subclass of media. A CD, for example, appears as an `IOCDMedia` subclass, with properties appropriate to a CD, such as `kIOCDMediaTOKKey`. Such properties can be combined to describe a media object uniquely in the system, or generalized to identify a certain kind of media with multiple possible matches (returned in an iterator).

The Carbon APIs and Cocoa APIs provide mechanisms for sending notifications of file system mount and unmount events, as well as for file access within the file system. An application can obtain the associated BSD device interface name for a given file system through `GetVolParms` (`VMDeviceID`) in Carbon and through `getmntinfo` in BSD. The associated `IOMedia` object can be obtained for a given BSD device interface name through I/O Kit APIs using the `kIOBSDNameKey` property.

USB

The USB family provides support for, and access to, devices attached to a Universal Serial Bus (USB).

Two basic types of drivers are clients of this family: kernel-mode drivers and user-mode drivers. Kernel-mode drivers are required when the clients of the driver also reside in the kernel (such as HID devices, mass storage devices, or networking devices). User-mode drivers are preferred when only one process has access to the device (for example printers and scanners).

Bundle identifier:

- `com.apple.iokit.IOUSBFamily`

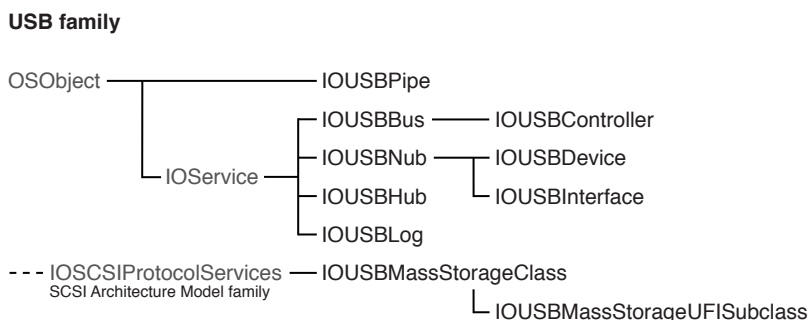
Headers in:

- Kernel resident: `Kernel.framework/Headers/IOKit/usb/`
- Device interface: `IOKit.framework/Headers/usb`

References and specifications:

- USB.org—<http://www.usb.org>; see especially USB Common Class Specification, revision 1.0, available for download at http://www.usb.org/developers/devclass_docs/usbccs10.pdf
- Apple Developer Connection—<http://developer.apple.com/hardwaredrivers/usb/index.html>

Class hierarchy:



Device interface:

- User mode clients use an API which is part of the I/O Kit framework; this API is defined in `IOKit/usb/IOUSBLib.h`. Clients use user mode abstractions of the `IOUSBDevice` and `IOUSBInterface` classes found in the kernel in order to communicate with the USB device or USB interface.

Kernel-resident drivers:

- Kernel drivers for physical USB devices can be written for either USB devices or for USB interfaces. A physical USB device consists of a device descriptor that can describe any number of interfaces. When writing a kernel-resident driver, you need to decide if the driver is to control the whole USB device or if it is to control only an interface of a USB device.

The USB family for kernel-resident drivers consists of three main classes:

- **IOUSBDevice:** The `IOUSBDevice` class is an abstraction of a physical USB device. There is one `IOUSBDevice` class instantiated for every USB device connected to the bus. The provider for an `IOUSBDevice` object is an `IOUSBController` object (which is an abstraction of a USB controller).
- **IOUSBInterface:** The `IOUSBInterface` class is an abstraction of one of the interfaces of a USB device. There is one of these classes instantiated for every interface in a device. When it is created, the `IOUSBInterface` creates `IOUSBPipe` objects for each endpoint described in the interface descriptor of the interface. The provider for an `IOUSBInterface` object is an `IOUSBDevice` object.
- **IOUSBPipe:** The `IOUSBPipe` class contains the methods that are used for communicating with a USB device or a USB interface. There is one `IOUSBPipe` object created for the default control endpoint and an additional one for each endpoint described in the interface descriptor. The provider for an `IOUSBPipe` object is an `IOUSBInterface` object.

Kernel-resident USB drivers are clients of the family that provides the transport services and are members of the family from which they get their class inheritance. For example, a driver for a USB keyboard is a client of the `IOUSBInterface` object (that is, the `IOUSBInterface` object is the provider for the driver) but the keyboard driver is a member of the `IOHIDFamily`. The USB family provides the mechanism for getting at the key presses in the keyboard. The keyboard driver supports methods from the HID family for sending those key presses to the event system.

As mentioned above, you can write a driver to match against a USB device or a USB interface. The `IOUSBDevice` or `IOUSBInterface` classes are the providers for the drivers. The drivers themselves can be members of a separate family (such as the `IOHIDFamily` or `IOAudioFamily`) or can be members of the `IOService` family.

Power management:

All in-kernel USB device drivers should implement at least basic power management to increase power saving in the system. In OS X v10.5, the USB family introduced the `IOWUSBHubPolicyMaker` object, which is an abstraction of a USB hub that includes power-management capabilities. When you develop a USB device driver to run in OS X v10.5 and later and you call `joinPMtree`, your driver is attached into the power plane as a power child of an `IOWUSBHubPolicyMaker` object. (In earlier versions of OS X, the power parent of a USB device driver was an `IOWUSBController` object representing the controller to which the device was attached.)

Your USB device driver can communicate with its `IOWUSBHubPolicyMaker` object to determine the power state of its hub. This can be useful if, for example, you need to handle an imminent shutdown differently from a restart. The `IOWUSBHubPolicyMaker` object supports the following five power states for a hub:

- **On.** The hub is fully functional and at least one of its ports is active (that is, not suspended).
- **Sleep.** If the hub supports sleep, its ports are inactive and it is not supplying power to any attached devices; if not, the sleep state is identical to the off state.
- **Doze.** This is an idle power-saving state a hub can enter when all its ports are suspended or disconnected and all attached devices are in either the off or doze state. Not all hubs support the doze state.
- **Off.** The hub enters this state when the system is about to shut down.
- **Restart.** This state is identical to the off state, except that a hub enters it when the system is about to restart.

USB devices seldom support more than the first four of these power states, and many support only on, sleep, and off. If you're developing a driver for a USB device that supports doze, you should call `SuspendDevice` when you switch the device's power state to doze, so the hub can suspend the port to which the device is attached. If the drivers for all the devices attached to a hub do this, the hub can enter the doze state, which can result in significant power saving for the system. However, if your driver calls `SuspendDevice` without also changing the device's power state to doze, you prevent the hub from entering the doze state and saving power.

If the hub to which your device is attached does not support sleep, the device must go to its off state when the system goes to sleep. Of course, if your device does not support sleep, it must also go to its off state when the system goes to sleep, even if its hub does support sleep.

During a power-state change, note that a USB hub's power state is not updated until the hub receives a `powerChangeDone` call, which does not happen until after all downstream hubs and devices have received their `powerChangeDone` calls. This means that if you use the `getPowerState` call introduced in OS X v10.5 to get an upstream hub's current power state, you might receive stale power-state information. However, if your device is changing to its on state, you can assume that the upstream hubs are actually on, even if their power-state values haven't yet been updated.

Table A-14 Clients and providers of the USB family

	Client of the nub	Provider for the nub
Action	Drives a device that plugs into a USB port.	Drives a USB bus controller.
Example	A keyboard driver is an interface driver; its provider is the driver for the USB <i>device</i> . The keyboard driver inherits from the IOHIDKeyboard class—it's a member of the HID family.	
Classes	IOUSBDevice drivers can only communicate with the USB device through the default control pipe. IOUSBInterface drivers have IOUSBPipe objects created for all the endpoints that are described in the interface descriptor for the current configuration. The driver uses these objects to communicate with the device.	USB family-member drivers should inherit from the IOUSBController class.
Notes	Common client families include the HID family (IOHIDPointing and IOHIDKeyboard classes) and the Transport Drivers for Storage family devices. Most kernel-mode clients of the USB family are interface drivers, and only occasionally device drivers.	OS X includes generic drivers that support all Open Host Controller Interface (OHCI) bus controllers. In general, third-party developers do not need to write drivers for the USB family.

Driver matching: The USB family uses the USB Common Class Specification, revision 1.0 to match devices and interfaces to drivers (for a link to this specification, see the section above titled “References and specifications”). The driver should use the keys defined in this specification to create a matching dictionary that defines a particular driver personality. There are two tables of keys in the specification. The first table contains keys for finding devices and the second table contains keys for finding interfaces. Both tables present the keys in order of specificity: the first key in each table defines the most specific search and the last key defines the broadest search. Each key consists of the combination of elements listed in the left column of the table.

For a successful match, you must add the elements of exactly one key to your personality. If your personality contains a combination of elements not defined by any key, the matching process will not find your driver. For example, if you're attempting to match a device and you add values representing that device's vendor, product, and protocol to your matching dictionary, the matching process is unsuccessful even if a device with those same values in its device descriptor is currently represented by an IOUSBDevice nub in the I/O Registry. This is because there is no key that combines the elements of vendor number, product number, and protocol.

Devices Without I/O Kit Families

Some categories of devices do not have family support from I/O Kit. In general, there are three reasons why a particular device may not be supported by an I/O Kit family.

- Support for certain devices is provided by other frameworks. The I/O Kit is not the most appropriate place for the abstractions that represent these devices. Examples of such devices include printers, scanners, digital cameras, and other imaging devices. If you are developing a driver for this category of device, you should use the appropriate imaging SDK.
- For some devices, it is not possible to provide a set of useful, common abstractions. Such devices might include USB security dongles, data acquisition cards, and other vendor-specific devices. These devices do not share a sufficiently large number of characteristics to make creation of I/O Kit families worthwhile. For example, although security dongles all connect via USB, there is no easily defined set of abstractions common to all such devices. An I/O Kit family would not provide substantial assistance to developers. It should not be assumed, however, that a family is required to write a new driver. In many cases, the `IOService` class provides everything a driver requires to write a “family-less” driver.
- For many devices, it is possible to define a set of useful abstractions; however, Apple has not chosen to create a family for one or more reasons. These devices may be part of a technology that is not a common Macintosh market. Or, Apple’s engineers may not have sufficient in-house expertise with certain devices to create the best family definition. In these cases, an opportunity exists for third-party developers to extend the I/O Kit model by developing families of their own. In addition, families developed under the Apple Public Source License can be sent back to Apple for possible inclusion in future releases of OS X.

Imaging Devices

There is no I/O Kit family for imaging devices (cameras, printers, and scanners). Instead, support for particular imaging device characteristics is handled by user-space code (see [Controlling Devices From Outside the Kernel](#) (page 35) for further discussion). Developers who need to support imaging devices should employ the appropriate imaging SDK.

Digital Video

To add digital video capabilities to your software, use the QuickTime APIs.

Sequential Access Devices (Tape Drives)

There is, at present, no I/O Kit family specifically designed for sequential access devices, such as tape drives. However, third-party developers can use the SAM device interface to create plug-in components for such devices.

Telephony Devices

There is, at present, no I/O Kit family for telephony devices. Apple is evaluating plans for a Telephony family for the future.

Vendor-Specific Devices

For some devices, it is not possible to provide a set of useful, common abstractions. Because families define the set of abstractions shared by all devices within the family, it is not feasible to create a family for these devices.

In most cases, however, a family is not necessary in order to write a driver for these devices. Developers should start by inheriting functionality from `IOService`, then use the `GetProperty` and `SetProperty` calls to communicate with their driver. In many cases, this should suffice. In some cases, however, such as data acquisition cards requiring high bandwidth, the developer should create their own user client (for a device-interface plug-in). Such objects can provide shared memory and procedure-call interfaces to a user-space library (see `IOUserClient.h`). You can find several good examples in `IOKitExamples` on the Darwin Open Source site.

Document Revision History

This table describes the changes to *IOKit Fundamentals*.

Date	Notes
2014-04-09	Removed description of completion chaining; this technique is not recommended.
2007-05-17	Updated and clarified the "Managing Power" chapter.
2006-11-07	Added information about using IODMACommand.
2006-10-03	Made minor corrections.
2006-05-23	Added caveat regarding in-function static constructors.
2006-02-07	Added information on the use of namespaces and a caveat about nested classes.
2005-12-06	Made minor corrections.
2005-11-09	Fixed minor typos.
2005-04-08	Fixed links, typos. Added note that Objective-C does not supply I/O Kit interfaces.
2004-10-05	Removed information about deprecated SCSI family; added information about new SCSI Parallel family.
2004-08-31	Added information about handling interrupts when implementing pseudo-DMA.
2004-05-27	Changed outdated links.

Date	Notes
2004-04-22	Updated documentation references, added information on Xcode.
2004-02-13	Fixed minor errors.
2003-10-10	Added information about the use of IOBufferMemoryDescriptor objects to represent kernel-allocated buffers in user-space tasks.
2003-09-18	Updated for OS X v10.3. Added information about changes in memory subsystem to support 64-bit architectures.



Apple Inc.
Copyright © 2001, 2014 Apple Inc.
All rights reserved.

No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by any means, mechanical, electronic, photocopying, recording, or otherwise, without prior written permission of Apple Inc., with the following exceptions: Any person is hereby authorized to store documentation on a single computer or device for personal use only and to print copies of documentation for personal use provided that the documentation contains Apple's copyright notice.

No licenses, express or implied, are granted with respect to any of the technology described in this document. Apple retains all intellectual property rights associated with the technology described in this document. This document is intended to assist application developers to develop applications only for Apple-branded products.

Apple Inc.
1 Infinite Loop
Cupertino, CA 95014
408-996-1010

Apple, the Apple logo, Carbon, Cocoa, Finder, FireWire, Logic, Mac, Mac OS, Macintosh, Objective-C, OS X, Pages, Power Mac, Quartz, QuickTime, and Xcode are trademarks of Apple Inc., registered in the U.S. and other countries.

.Mac is a service mark of Apple Inc., registered in the U.S. and other countries.

CDB is a trademark of Third Eye Software, Inc.

Intel and Intel Core are registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.

OpenGL is a registered trademark of Silicon Graphics, Inc.

PowerPC and the PowerPC logo are trademarks of International Business Machines Corporation, used under license therefrom.

UNIX is a registered trademark of The Open Group.

VMS is a trademark of Digital Equipment Corporation.

APPLE MAKES NO WARRANTY OR REPRESENTATION, EITHER EXPRESS OR IMPLIED, WITH RESPECT TO THIS DOCUMENT, ITS QUALITY, ACCURACY, MERCHANTABILITY, OR FITNESS FOR A PARTICULAR PURPOSE. AS A RESULT, THIS DOCUMENT IS PROVIDED "AS IS," AND YOU, THE READER, ARE ASSUMING THE ENTIRE RISK AS TO ITS QUALITY AND ACCURACY.

IN NO EVENT WILL APPLE BE LIABLE FOR DIRECT, INDIRECT, SPECIAL, INCIDENTAL, OR CONSEQUENTIAL DAMAGES RESULTING FROM ANY DEFECT, ERROR OR INACCURACY IN THIS DOCUMENT, even if advised of the possibility of such damages.

Some jurisdictions do not allow the exclusion of implied warranties or liability, so the above exclusion may not apply to you.

Index

Numerals

16-bit PC cards [156](#)

64-bit architectures, issues with [100–103](#)

A

accessor functions, IOService class [67](#)

acknowledgePowerChange method [126](#)

Action function [78, 82, 89](#)

active matching [31, 49, 65](#)

activityTickle method [125](#)

ADB family [142](#)

addEventSource function [84, 91](#)

AGP family (PCI and AGP family) [158](#)

alloc function [56](#)

allocClassWithName function [56](#)

Apple Developer Connection [12](#)

applications

 accessing hardware from [35](#)

 controlling imaging devices from [16](#)

 for I/O Kit development [19–20](#)

arbitration [25](#)

architecture of I/O Kit [23–38](#)

ATA and ATAPI family [143](#)

attach function [49–50, 64, 65](#)

Audio family [145](#)

Audio HAL [145](#)

Audio Hardware Abstraction Layer. *See* Audio HAL

Audio plane [40](#)

authentication [36](#)

B

base classes

 in I/O Kit [61, 68](#)

 in libkern [53, 61](#)

big-endian format [28](#)

block devices [38](#)

BSD device interface [166, 171](#)

BSD network stack [153](#)

BSD sockets [38](#)

BSD

 and custom device interfaces [37](#)

 and kernel-resident code [29](#)

 as design principle of I/O Kit [15](#)

 device files in [38](#)

bus controller drivers [48, 105, 106](#)

byte swapping [28](#)

C

C++ [16](#)

 and I/O Kit implementation [13](#)

 subset of used by libkern [53](#)

callback functions. *See* Action functions

Carbon Draw Sprockets [149](#)

Carbon environment, using BSD sockets in [38](#)

CardBus cards [156](#)

CFPlugIn [36, 149](#)

character devices [38](#)

checkForWork function [82](#)

class hierarchy in I/O Kit [31, 52–53](#)

class information functions [57](#)

class matching [30, 48, 65](#)

client configuration [24](#)

close function [66](#)

Cocoa environment, using BSD sockets in [38](#)

command gates [90](#)

- creating 91
- disposing of 93
- issuing I/O requests through 92
- registering 91
- command-based events 83
- command-line tools 19–20
- commandGate method 91
- commands, as event sources 34
- completion routines 79
 - See also Action functions
- Component Object Model (COM) 36
- constructors 54, 58
- context switching 78
- controller drivers 48, 105, 106
- Core Audio framework 41, 145
- Core Foundation Plug-in Services. See CFPlugIn
- Core MIDI framework 145

D

- DART (device address resolution table) 100
- Darwin Open Source 21
- data containers 33
- data management 95–107
- deadlocks
 - and hardware constraints 106
 - and multiple work loops 78
 - avoiding 78, 91
 - paging 91
- dedicated work loops 80
- delete operator 54, 55
- design principles 15–16
- destructors
 - defining 58, 59
 - of OSObject class 54
- detach function 49–50, 64, 65
- detach method 131
- developer resources 11–12
- development applications 19
- device drivers 19, 20

- device interfaces 14, 22, 36–37
- device matching 37, 51
- device memory 34, 68
- device probing 49–50, 63, 65
- device removal 129, 131
- Device Tree plane 41
- devices
 - determining idleness of 123–125
 - imaging 16
 - network 22
 - power state of 111
 - removal of 129–131
 - serial 22
 - storage 22
 - unsupported 16
- dictionaries. See matching dictionaries
- didTerminate method 131
- digital cameras 16
- digital video 176
- direct interrupts 80, 83, 84
- direct memory access. See DMA
- display devices 148
- DLIL (Data Link Interface Layer) 155
- DMA
 - and controller drivers 105–106
 - engines 96, 105, 107
- down calls 91
- driver development 19–20
- driver layering 23–27
- driver loading 50, 64–66
- driver matching 48
- driver matching 43–49
 - and I/O Catalog 40
 - as service of IOService 63, 64, 66
 - overview of 30
- driver object life cycles 49
- driver personalities 30, 43–48
- driver shutdown 66
- drivers

- and families [24, 69, 71](#)
- and nubs [25](#)
- life cycle of [27, 34](#)
- loading [25](#)
- dynamic allocation of objects [56](#)
- dynamic type casting [57](#)

E

- Ethernet controller [26](#)
- event handling [77](#)
 - and work loops [77, 81](#)
 - deferring work [82](#)
 - device removal [129](#)
 - of timer events [89, 90](#)
- event sources [28, 34, 82](#)
 - adding to work loops [82](#)
 - and interrupt handling [83–89](#)
 - categories of [77](#)
 - classes for [78, 82](#)
 - disposing of [82](#)
 - timer events [89–90](#)
- exceptions, as disallowed feature of C++ [16, 53](#)

F

- families, I/O Kit [69–76](#)
 - and drivers [69–71](#)
 - as libraries [71–73](#)
 - creating [75](#)
 - defined [24](#)
 - devices with no families [176](#)
 - list of [34](#)
 - loading [25](#)
 - naming conventions for [74](#)
 - reference for [142–175](#)
 - structure of [73–75](#)
 - typical classes in [74](#)
 - unsupported [35](#)
- features of I/O Kit [14–15](#)

- file system, extensions to [22](#)
- filter interrupt event sources [87–88](#)
- filter scheme drivers [168, 169](#)
- FireWire config ROM [147](#)
- FireWire family [146](#)
- FireWire plane [41](#)
- frame buffers [148](#)
- frameworks [18](#)
- free function [49, 55, 59, 66](#)
- function call stacks [29](#)

G

- gating mechanism [28](#)
- gcc command-line tool [20](#)
- gdb command-line tool [20](#)
- getClassname function [58](#)
- getClassSize function [58](#)
- getInstanceCount function [58](#)
- getPhysicalSegments function [96](#)
- getProperty function [62](#)
- getSuperClass function [58](#)
- getWorkLoop function [79](#)
- graphics acceleration [14, 149](#)
- Graphics family [148](#)

H

- hardware constraints on bus controller drivers [106](#)
- hardware modeling [23](#)
- hardware, support for [16](#)
- header files [18](#)
- HID family [151](#)
- HID Manager [151](#)
- hot-swapping [129, 131](#)
- Human Interface Device family. *See* HID family

I

- i.LINK standard [146](#)
- I/O addresses [100](#)

- I/O Catalog [14, 30, 39](#)
- I/O commands [77](#)
- I/O Kit families. *See* families, I/O Kit
- I/O Kit framework [19, 39](#)
- I/O queues, clearing [131](#)
- I/O Registry [39–41](#)
 - and storage devices [171](#)
 - application for [19](#)
 - architecture [39](#)
 - device matching and [51](#)
 - driver matching and [49](#)
 - examining [41–42](#)
 - introduced [29–30](#)
 - planes defined in [40](#)
- I/O Registry Explorer application [19, 23, 41–42](#)
- I/O requests [24, 90](#)
 - buffers for [103](#)
 - memory in [97, 100](#)
 - relaying [103](#)
- I/O transfers [95–100](#)
- I/O vectors. *See* scatter/gather lists
- idleness of devices [123](#)
- IEEE 1394 standard [146](#)
- imaging devices [16, 176](#)
- in-function static constructors [17](#)
- indirect interrupts [77, 83](#)
- information property lists [30, 40](#)
- init function [49, 50, 54, 59, 64](#)
- initialization methods [59](#)
- interrupt controllers
 - and direct interrupts [78, 83](#)
 - and indirect interrupts [84](#)
- interrupt event sources
 - disposing of [86](#)
 - ordering of [84–85](#)
 - setting up [85](#)
- interrupt events [77, 83](#)
- interrupt handlers [85, 88](#)
- interrupt handling [68, 83, 89](#)
- interrupts [34](#)
- introspection, of objects [57](#)
- IOADBController class [142, 143](#)
- IOADBDevice [142, 143](#)
- IOAGPDevice class [158](#)
- ioalloccount command-line tool [20](#)
- IOApplePartitionScheme [170](#)
- IOATACCommand class [145](#)
- IOATADevice [145](#)
- IOATAPIProtocolTransport class [162](#)
- IOAudioDevice class [146](#)
- IOAudioEngine class [146](#)
- IOBigMemoryCursor class [107](#)
- IOBlockStorageDevice class [166, 167](#)
- IOBlockStorageDriver class [167](#)
- IOBlockStorageServices class [161](#)
- IOBufferMemoryDescriptor class [96](#)
- IOCardBusDevice class [156](#)
- IOCDMedia class [171](#)
- ioclasscount command-line tool [20](#)
- IOCommandGate class [78, 82, 90](#)
- IOCompactDiscServices class [161](#)
- IODisplay class [143](#)
- IODMACommand class [101](#)
- IODVDServices class [161](#)
- IOEventSource class [82](#)
- IOFDiskPartitionScheme [170](#)
- IOFilterInterruptEventSource class [87](#)
- IOFireWireController class [148](#)
- IOFireWireSBP2LUN class [159](#)
- IOFireWireSBP2Target class [159](#)
- IOFireWireSerialBusProtocolTransport class [162](#)
- IOFireWireUnit class [148](#)
- IOFramebuffer class [149](#)
- IOHIDKeyboard class [143, 175](#)
- IOHIDPointing class [143, 175](#)
- IOInterruptEventSource class [78, 82, 84, 87](#)
- IOMKitPersonalities dictionary [43](#)
- IOLittleMemoryCursor class [107](#)

- IOMatchCategory 46
 - IOMbufMemoryCursor class 155
 - IOMedia class 168
 - IOMedia filter schemes 169
 - IOMedia properties 170
 - IOMemoryCursor class 96, 107
 - IOMemoryDescriptor class 96, 99, 103
 - IOMemoryMap 104
 - IOModemSerialStreamSync class 166
 - IOMultiMemoryDescriptor class 104
 - IONaturalMemoryCursor class 107
 - IONDRVFramebuffer class 149
 - IONetworkController class 155
 - IONetworkData class 155
 - IONetworkMedium class 155
 - IOOutputQueue class 155
 - IOPacketQueue class 155
 - IOPartitionScheme class 170
 - IOPCCard16Device class 156
 - IOPCCard16Enabler class 156
 - IOPCCardBridge class 156
 - IOPCI2PCIBridge class 156
 - IOPCIBridge class 156, 158
 - IOPCIDevice class 80, 156, 158
 - IOPMPowerState structure 119
 - IOPower plane. *See* Power plane
 - IOReducedBlockServices class 161
 - ioreg command-line tool 20, 23, 30, 41
 - IORegistryEntry class 62–63
 - in class hierarchy 33, 52, 61
 - member functions in 62
 - IOResources 45
 - IORS232SerialStreamSync class 166
 - IOSCSICommand class 161
 - IOSCSIDevice class 161
 - IOSCSIParallelInterfaceController class 161
 - IOSCSIParallelInterfaceProtocolTransport class 162
 - IOSCSIPeripheralDeviceType classes 161
 - IOSCSIPeripheralDeviceNub class 161
 - IOSCSIPrimaryCommandsDevice class 162
 - IOSCSIProtocolServices class 162
 - IOSerialDriverSync class 166
 - IOSerialStreamSync class 166
 - IOService class 33, 49, 61, 63–68
 - accessor functions of 67
 - in class hierarchy 52
 - IOServiceNotificationHandler type 67
 - IOServiceOpen function 68
 - iostat command-line tool 20
 - IOStorage class 167, 169
 - IOSubMemoryDescriptor class 104
 - IOTimerEventSource class 78, 82, 89
 - IOUSBController class 173, 175
 - IOUSBDevice class 173
 - IOUSBInterface class 173
 - IOUSBMassStorageClass class 162
 - IOUSBPipe class 173
 - IOWorkLoop class 77
 - IOZoomVideoDevice class 156
 - isEqualTo function 58
 - iteration functions, IORegistryEntry class 63
- ## K
-
- KDP (Kernel Debugger Protocol) 155
 - kernel development kit (KDK) 15
 - kernel environment
 - alternatives to programming in 22
 - and memory protection 21
 - caveats for programming in 21, 28
 - kernel extensions 22, 33, 43, 71
 - kernel modules
 - and I/O Kit classes 31
 - loading and unloading 53, 56
 - OSObject class and 33
 - setting version in 72
 - KEXT manager 72
 - kext load command-line tool 20
 - KEXTs. *See* kernel extensions

kextstat command-line tool [20](#)
kextunload command-line tool [20](#)
kIOMessageServiceIsTerminated message [66](#)
KMODs. *See* kernel modules

L

language used in I/O Kit [16](#)
layered architecture [23–27](#)
libkern library
 base classes in [53–61](#)
 in class hierarchy [31, 52](#)
 overview in I/O Kit [18](#)
 services provided by [28](#)
libraries [18, 71–73](#)
library versioning [72](#)
life cycle, driver object [64–66](#)
 device probing [49](#)
 functions of [64](#)
 managed by IOService class [63](#)
 of device drivers [27, 34](#)
 overriding functions [64](#)
limitations of I/O Kit [16](#)
little-endian format [28](#)

M

Mac OS 9 [15](#)
Mach [29](#)
Mach IPC [37](#)
Mach scheduler [83, 89](#)
Mach shared memory [37](#)
macros of OSMetaClass class
 object construction [57](#)
 object introspection [57](#)
 runtime type declaration [57, 58](#)
 type casting [57](#)
matching dictionaries [43, 62](#)
matchPropertyTable function [66](#)
mbuf structure [97](#)

memory cursors [96, 97, 104–105, 107](#)
memory descriptors [96–97, 103–104](#)
memory protection [29](#)
memory
 virtual. *See* virtual memory
 and I/O transfer interfaces [105](#)
 in I/O transfers [97–100](#)
 paging and [21](#)
 protection [15, 21](#)
 wiring down [28](#)
message function [66, 67](#)
messageClient function [67](#)
messaging [34, 67](#)
metaCast functions [58](#)
MIDI [145](#)
mkextcache command-line tool [20](#)
MMCDeviceInterface [163](#)
modHasInstance function [58](#)
multimedia support [14](#)
multiple inheritance, as disallowed feature of C++
 [16, 53](#)
multiprocessing [14](#)

N

namespaces [17](#)
naming conventions for I/O Kit families [74](#)
NDRV graphics drivers compatibility [150](#)
negotiation [36](#)
nested class [54](#)
network devices [22](#)
network drivers [97, 103](#)
Network family [152](#)
Network Kernel Extension (NKE) [155](#)
networking services [38](#)
new operator [54](#)
newUserClient function [68](#)
notification handler [67](#)
notifications [34, 63, 67](#)
nubs

- driver matching and 48
- overview of 24–25
- registering 64

O

- object allocation 56
- object construction 54, 56
- object disposal 55
- object introspection 57
- object retention 55
- Open Firmware 158
- open function 66
- Open Host Controller Interface (OHCI) 148, 175
- Open Transport 22
- OpenGL 149
- OSCheckTypeInst macro 58
- OSDeclareAbstractStructors macro 57
- OSDeclareDefaultStructors macro 57, 58
- OSDefineMetaClassAndAbstractStructors macro 57
- OSDefineMetaClassAndAbstractStructorsWithInit macro 57
- OSDefineMetaClassAndStructors macro 57, 59
- OSDefineMetaClassAndStructorsWithInit macro 57
- OSDictionary 34
- OSDynamicCast macro 58
- OSIterator objects 68
- OSMetaClass class 33
 - and runtime typing 55
 - in class hierarchy 52
 - macros in 57, 58
 - object construction and 54
 - type-casting macros of 57
- OSObject class 33, 52, 54–55
- OSTypeID macro 58
- OSTypeIDInst macro 58
- outputPacket function 103

P

- Package Maker application 19
- page-out threads 83
- paging
 - and deadlocks 91
 - and kernel code 21
- partition schemes 170
- passive matching 31, 48, 65
- PC Card family 156
- PCI and AGP family 26, 157
- PCI Localbus 2.1 specification 157
- personalities. *See* driver personalities
- physical memory
 - and I/O transfers 96, 98
 - kernel code restrictions 28
 - paging virtual memory into 28
- planes 40–41, 62
- Platform Expert 39
- plug-and-play feature 14
- plug-in interfaces 36
- policy makers
 - determining idleness 123–125
- positional functions, IORegistryEntry class 63
- POSIX
 - device files 38
 - device nodes 36
- power events 77
- power management 108
 - See also* notifications of power events
 - support for 14
- Power plane 41, 109
- power states
 - changing 120–123
 - defined 111
- PPP (Point-to-Point Protocol) 166
- preemptive multitasking support 14
- prepare function 96, 99
- prerequisites for driver development 13
- primary interrupts. *See* direct interrupts

- principal class of a personality 49
- printers 16
- probe function 49, 50, 64, 65
- probe scores 48, 49
- probing, devices 49, 50, 63, 65
- Programmed Input/Output (PIO) interface 105
- programming language 16
- property tables. *See* driver personalities
- protected memory 21
- provider matching 66
- provider-client relationships, tracking 29
- providers 48, 50

Q

- Quartz Compositor 22
- Quartz graphics system 149
- QuickTime, for video support 176

R

- read function 103
- reentrancy 28
- reference counting 55
- registerPowerDriver method 119
- registerService function 64
- release function 55
- resources for developers 11, 12
- responding to device removal 129–131
- retain function 55
- runAction function 92
- runCommand function 92
- runtime environment of I/O Kit 27–29
- runtime type information (RTTI), as disallowed
 - feature of C++ 16, 32, 53, 55
- runtime typing facility 53, 55, 58

S

- SBP-2 transport protocol 158
- SBP2 family 146, 158

- scanners 16
- scatter/gather lists 28, 96
- SCSI Architecture Model family 144, 161, 176
- SCSI disk drivers 26
- SCSI Parallel family 26, 160
- SCSI Parallel Interface-5 (SPI-5) specification 160
- SCSITaskDeviceInterface 163
- SCSITaskInterface 163
- SCSITaskLib 163
- SCSITaskUserClient 163
- secondary interrupts. *See* indirect interrupts
- sequential access devices 176
- serial devices 22
- Serial family 165
- Service plane 40
- setPowerState method 124
- shared interrupts 87
- shared work loops 78, 79
- sockets, BSD 38
- software development kit (SDK) 16
- SPI-5. *See* SCSI Parallel Interface-5 specification
- start function 50, 65, 78
- static constructors, in-function 17
- stop function 66
- storage devices 22
- storage drivers 166
- Storage family 26, 166
- structural events 78
- super macro 59
- symmetric multiprocessing support 14

T

- telephony devices 177
- templates, as disallowed feature of C++ 16, 53
- terminate method 130
- termios command 165
- threads, and work loops 77
- timeouts 89
- timer event sources 82, 89, 90

timer events [77, 89–90](#)
timers [34](#)
tools for development [19, 20](#)
Transport Driver layer [161](#)
type casting [57](#)

U

Unified Buffer Cache (UBC) [106](#)
Universal Page List (UPL) [106](#)
up calls [91](#)
UPS (uninterruptible power supply) devices [151](#)
USB Common Class Specification [175](#)
USB family [172](#)
USB plane [41](#)
user client [68, 98](#)
user client [36](#)
user thread interrupts [83](#)

V

virtual memory
 and I/O transfers [96, 98, 104](#)
 cache [106](#)
 in I/O Kit [15](#)
 kernel code restrictions [28](#)
 paging into physical memory [28](#)
virtual-memory (VM) pager [99, 106](#)

W

willTerminate method [131](#)
Window Manager interrupts [83](#)
work loops [77–81](#)
 adding and removing event sources [82](#)
 architecture of [78](#)
 dedicated [80](#)
 obtaining, examples of [80–81](#)
 querying of event sources [79, 84](#)
 shared [79](#)

X

Xcode application [19](#)

Z

Zoom Video cards [156](#)